


Memory Unit

- **Functional Units:**

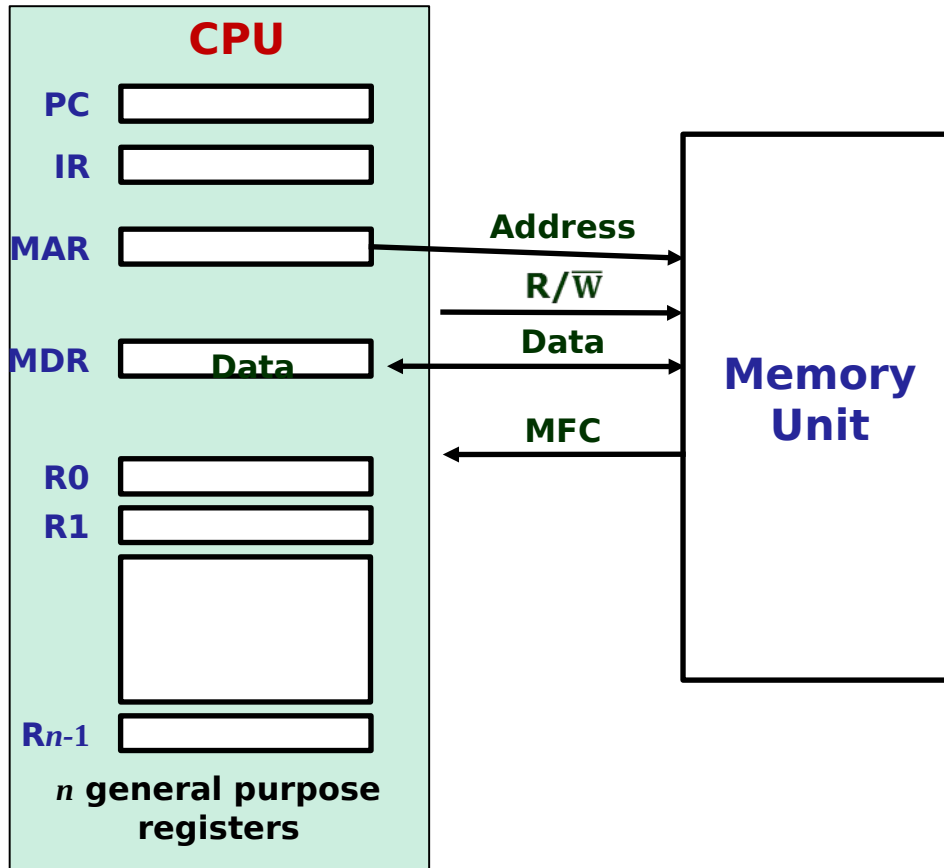
- Input Unit
- Memory Unit
- Arithmetic and Logic Unit (ALU)
- Output Unit
- Control Unit



Computer Memory

- Number & character operands, as well as instructions are stored in the memory of the computer
- Stored program concept
- CPU executes the instructions for which the instructions and operands have to come from memory unit
- Operations which involve memory:
 - Instruction fetch
 - Memory read
 - Memory operand fetch and store
 - Memory read
 - Memory write
- Instructions involving memory access:
 - LOAD and STORE instructions

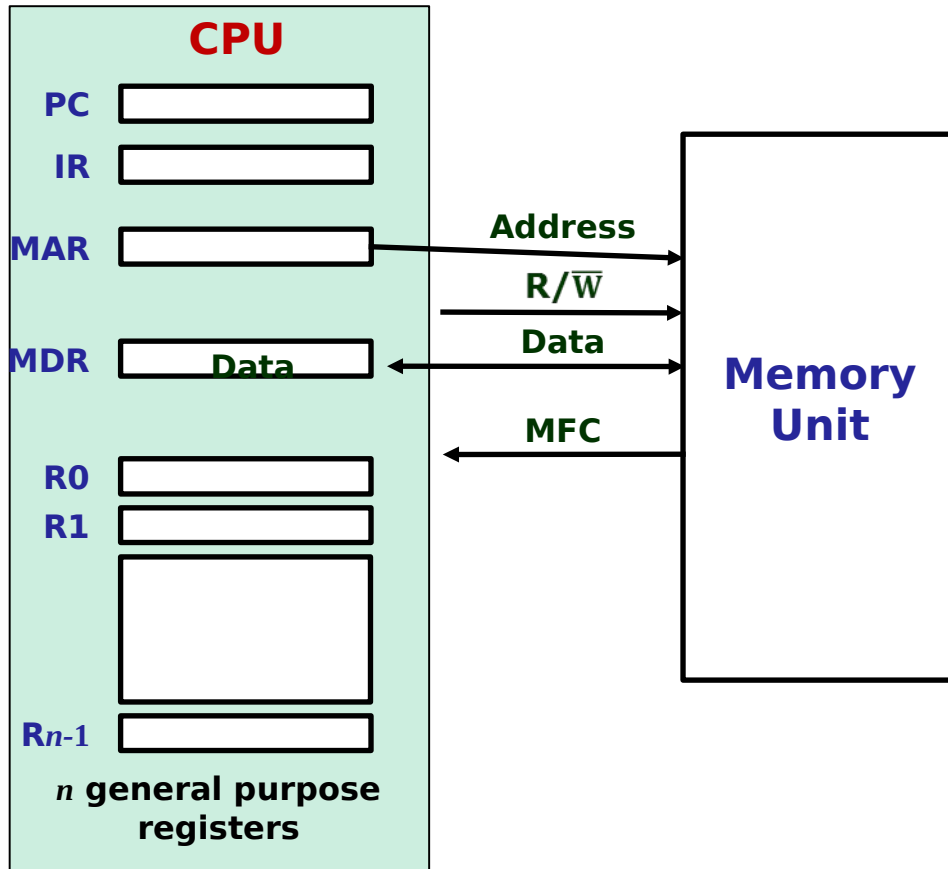
Memory Read and Write Operation



- **Read**

- Processor loads the address of memory location into MAR
- Set the R/\bar{W} line to 1
- Memory responds by placing the data from address location onto data line
- Confirm the action by asserting MFC (memory function complete) signal
- Upon receiving MFC signal, processor loads the data on data line into MDR

Memory Read and Write Operation



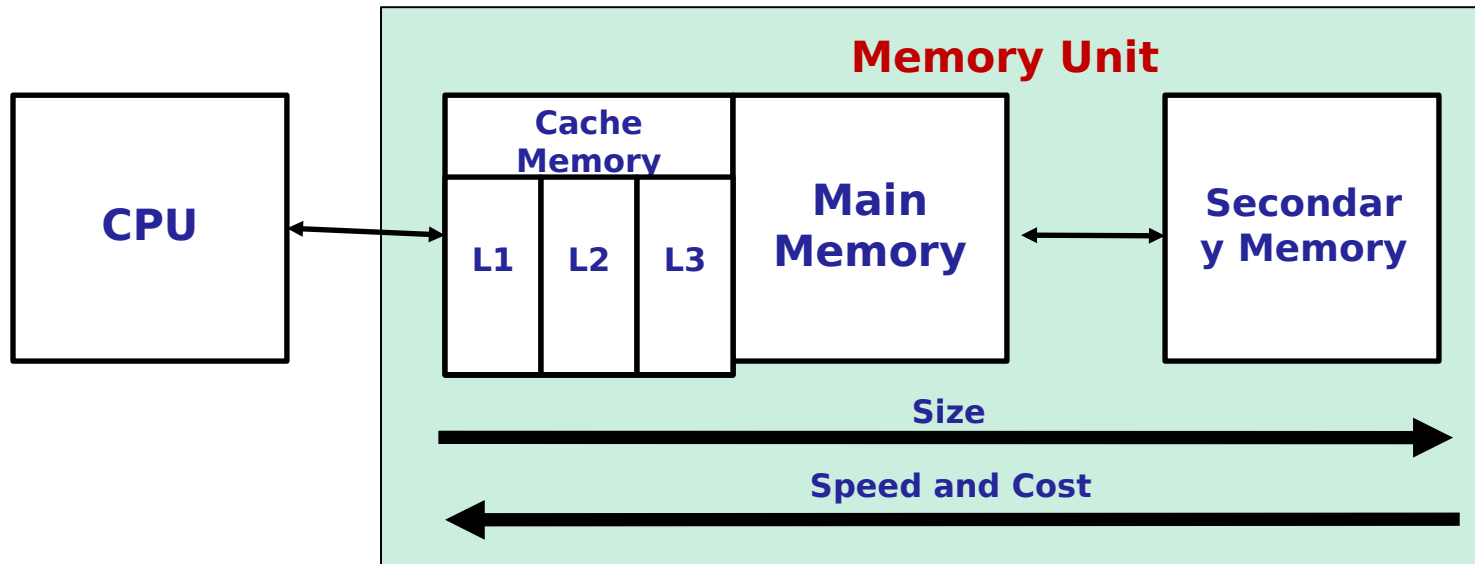
- **Write**

- Processor loads the address of memory location into MAR
- Processor loads data into MDR
- Set the R/\bar{W} line to 0 to indicate write operation
- Processor places the data in MDR onto data line
- Data on data line is written into memory location
- Memory confirms the action by asserting MFC signal

Memory Latency and Memory Organization

- **Latency:** Time to access the first of the sequence of memory words
- What is involved in determining the latency of the memory operation?
 - Processor issues the logical address to memory unit
 - The logical address need to be converted into physical address
- Memory unit is called random access memory (RAM)
 - Any location can be accessed for read/write operation independent of the location's address
- Memory unit is organised in hierarchical manner

Memory Hierarchy



- Processor processes instructions and data faster than it can be fetched from memory unit
- Memory access time is the bottleneck
- One way to reduce memory access time is to use faster memory
 - A small and faster memory bridge the gap between processor and main memory

Memory Performance Parameters

- **Memory Access Time:**

- Time interval between initiation of one operation and completion of that operation
- **Example:** Time between assertion of Read signal and MFC signal

- **Memory Cycle Time:**

- Minimum time delay between the initiation of two successive memory operations
- Time delay between start of a read/write operation to start of next memory operation

- Memory cycle time is usually slightly larger than access time

Internal Organization of Memory

- The memory is organised such that a group of n -bits can be stored or retrieved in a single basic operation
- Each group of n -bits is referred as one **memory word**
- Accessing the memory to store or retrieve information require **address for each location**
- Possible number of address locations are decided by the **number of address lines** in the processor
- For k -address lines, there will be 2^k locations, each of n -bit memory word
- 2^k addresses constitute **address space of computer**
- Example: Let $k=10$ and $n=32$
 - Number of locations: 2^{10}
 - Size of the memory: $2^{10} \times 2^5 \text{ bits} = 2^{15} \text{ bits}$
 $= 2^{12} \text{ Bytes}$
 $= 4 \text{ KB}$

Memory Content Example


1024 memory locations: 1x0 bit address

16 bit data


Memory address		Memory content
Binary	decimal	
0000000000	0	1011010101011101
0000000001	1	1010101110001001
0000000010	2	0000110101000110
	⋮	⋮
1111111101	1021	1001110100010100
1111111110	1022	0000110100011110
1111111111	1023	1101111000100101

Semiconductor Memories

- Two basic ways of designing memory
 - Static RAM (SRAM)
 - Dynamic RAM (DRAM)
- **Static RAM:**
 - Built using metal-oxide semiconductor (MOS) transistors
 - MOS transistors acts as switch
 - +5 v (when Gate input is 1): Transistor conducts: ON state
 - 0 v (when Gate input is 0): Transistor does not conducts: OFF state

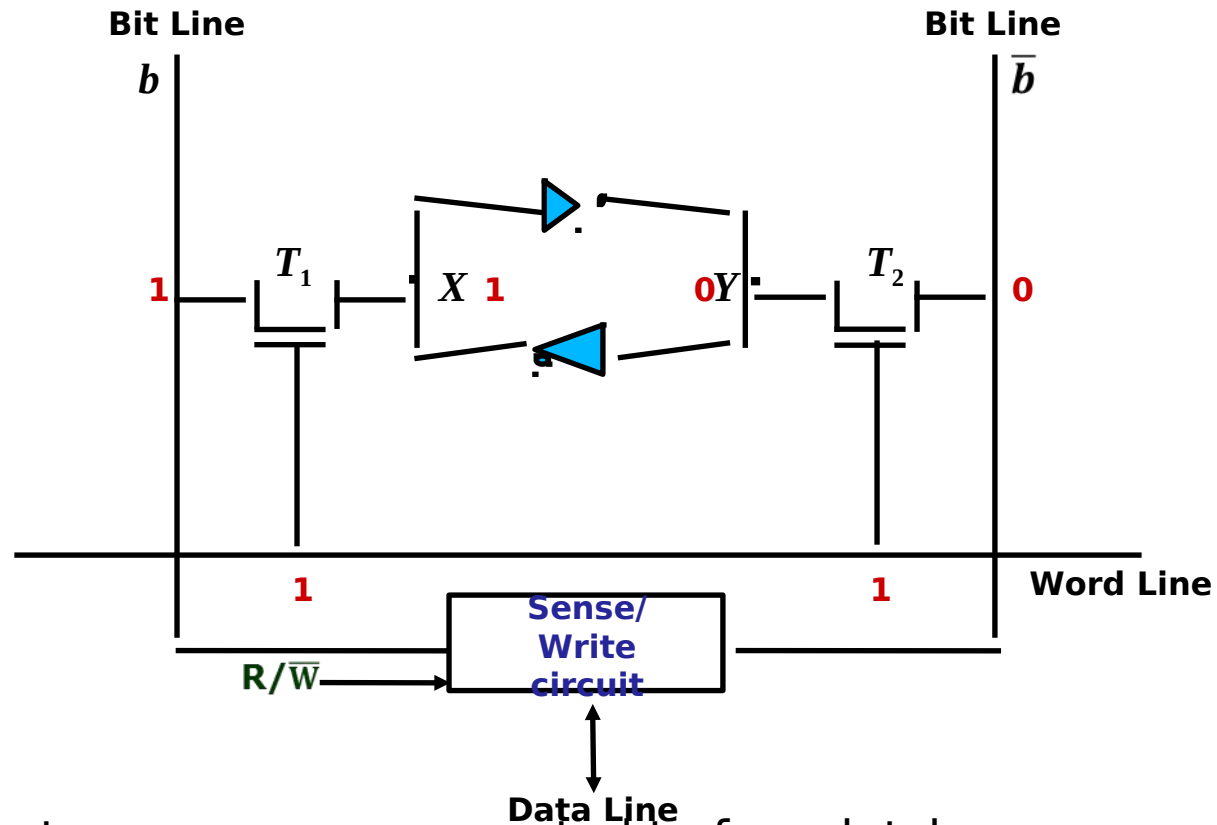


Static RAM Cell



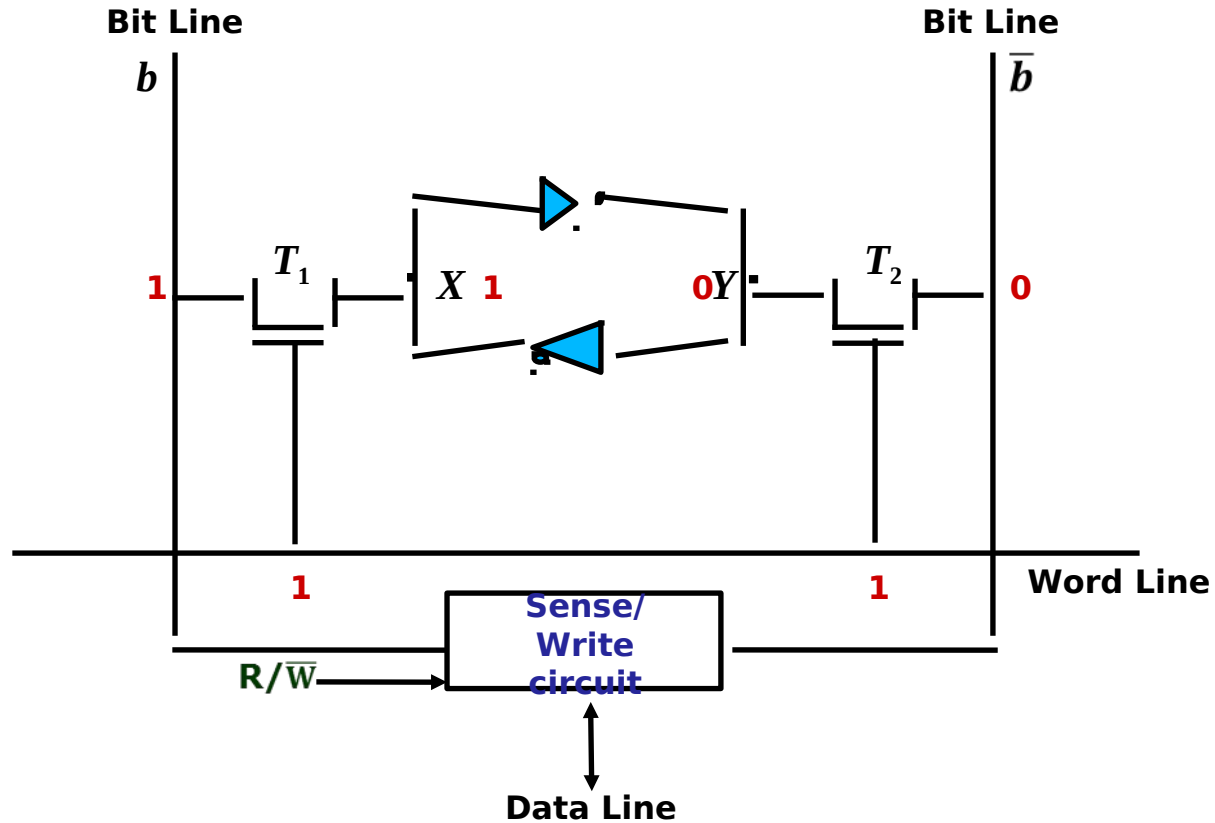
- Two inverters are cross connected to form latch
- Inverters are connected to 2 transistors which act as switches
- Switches are opened or closed under the control of word line
- This circuit retain the state (bit) as long as power is applied (Static Memory)

Static RAM Cell - Read



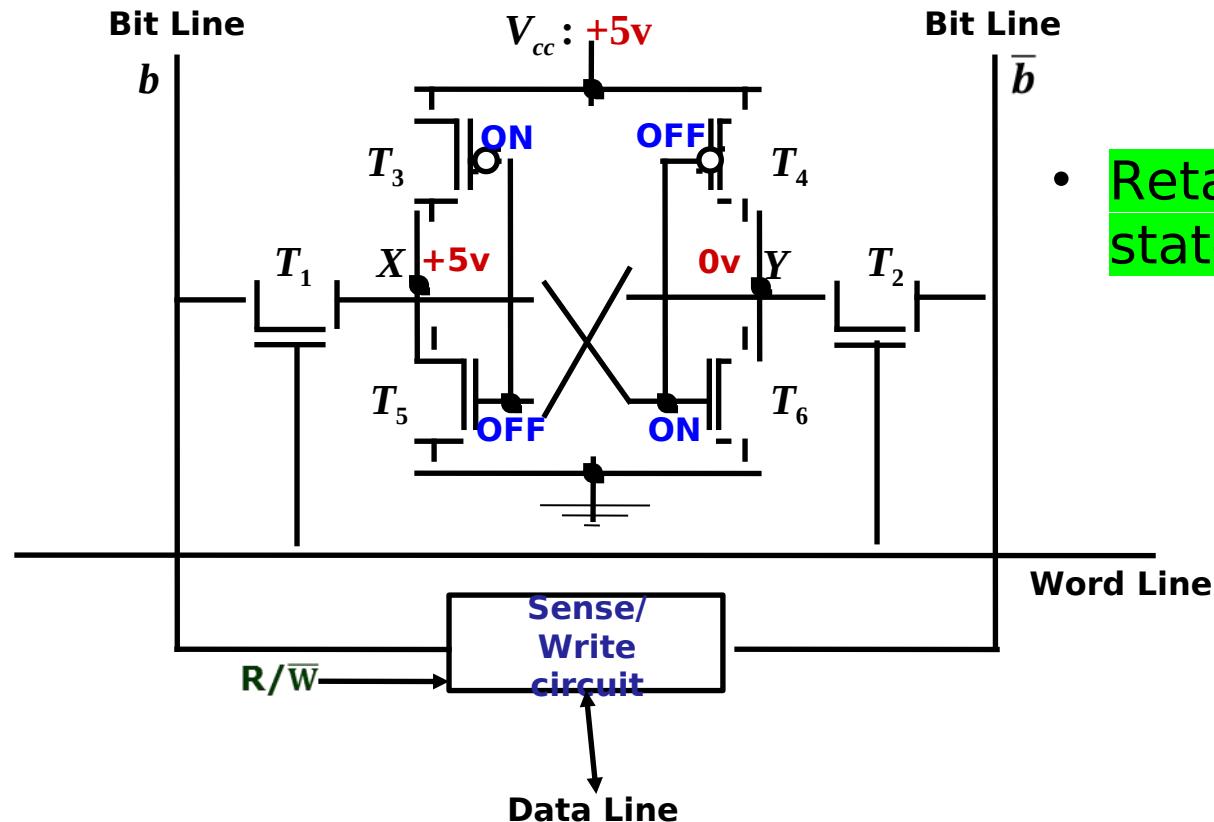
- Two inverters are cross connected to form latch
- Inverters are connected to 2 transistors which act as switches
- Switches are opened or closed under the control of word line
- This circuit retain the state (bit) as long as power is applied (Static Memory)

Static RAM Cell - Write



- Inverters are connected to 2 transistors which act as switches
- Two inverters are cross connected to form latch
- Switches are opened or closed under the control of word line
- This circuit retain the state (bit) as long as power is applied (Static Memory)

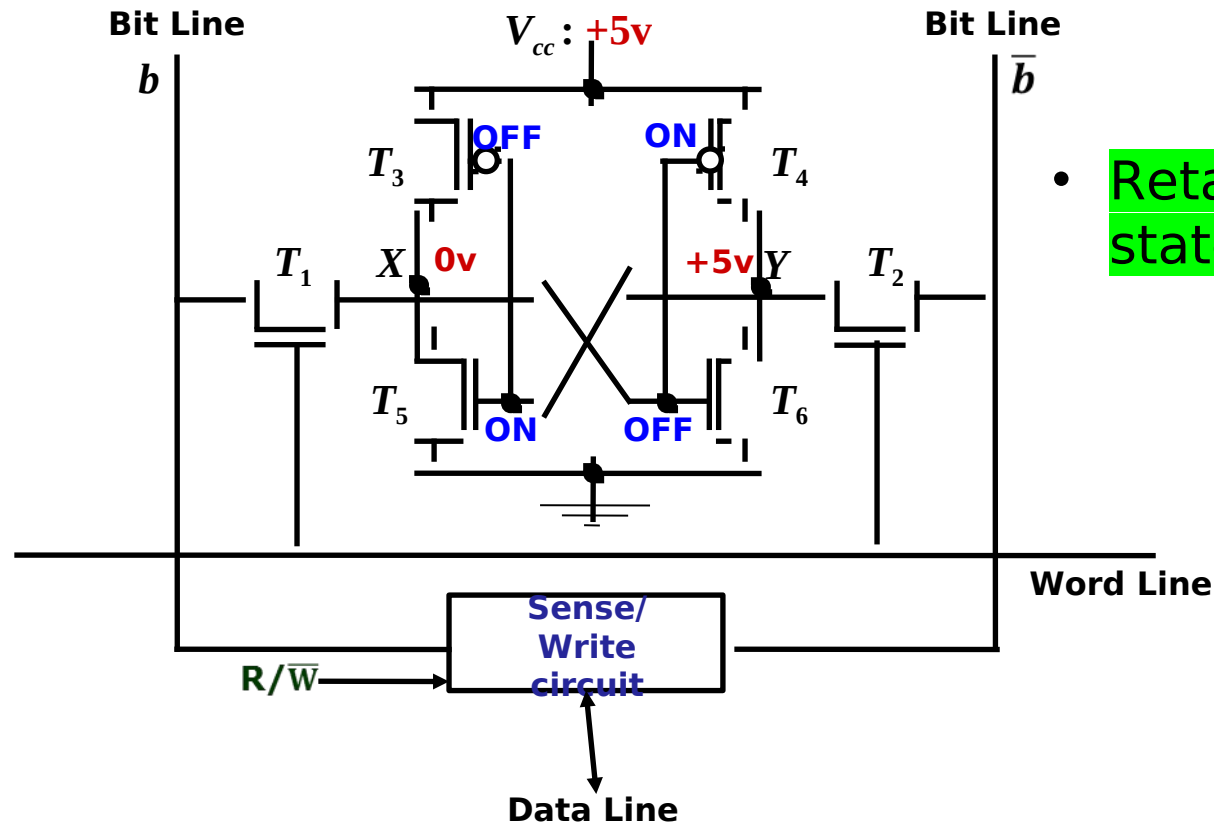
CMOS Static RAM Cell



- Retaining the state 1

- This circuit retain the state (bit) as long as power is applied (Static Memory)
- Continuous power is needed for a cell to retain the state

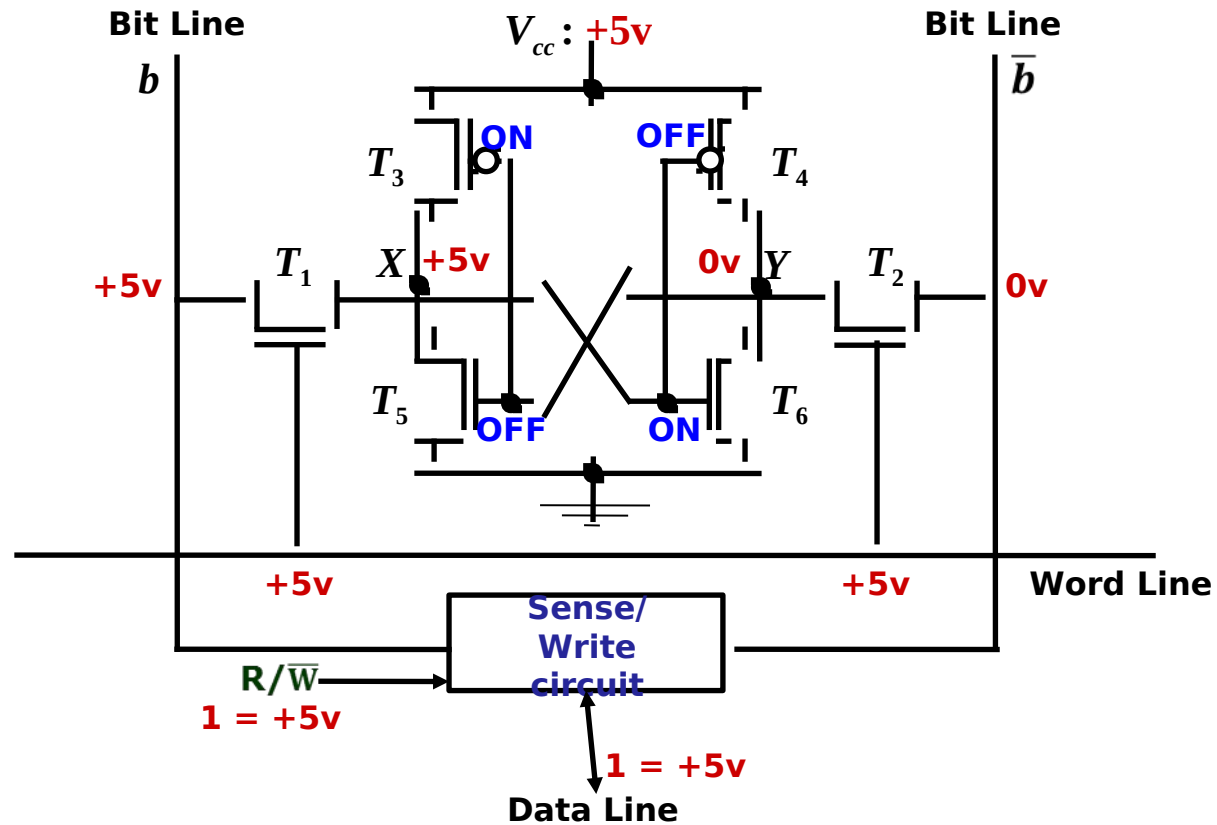
CMOS Static RAM Cell



- Retaining the state 0

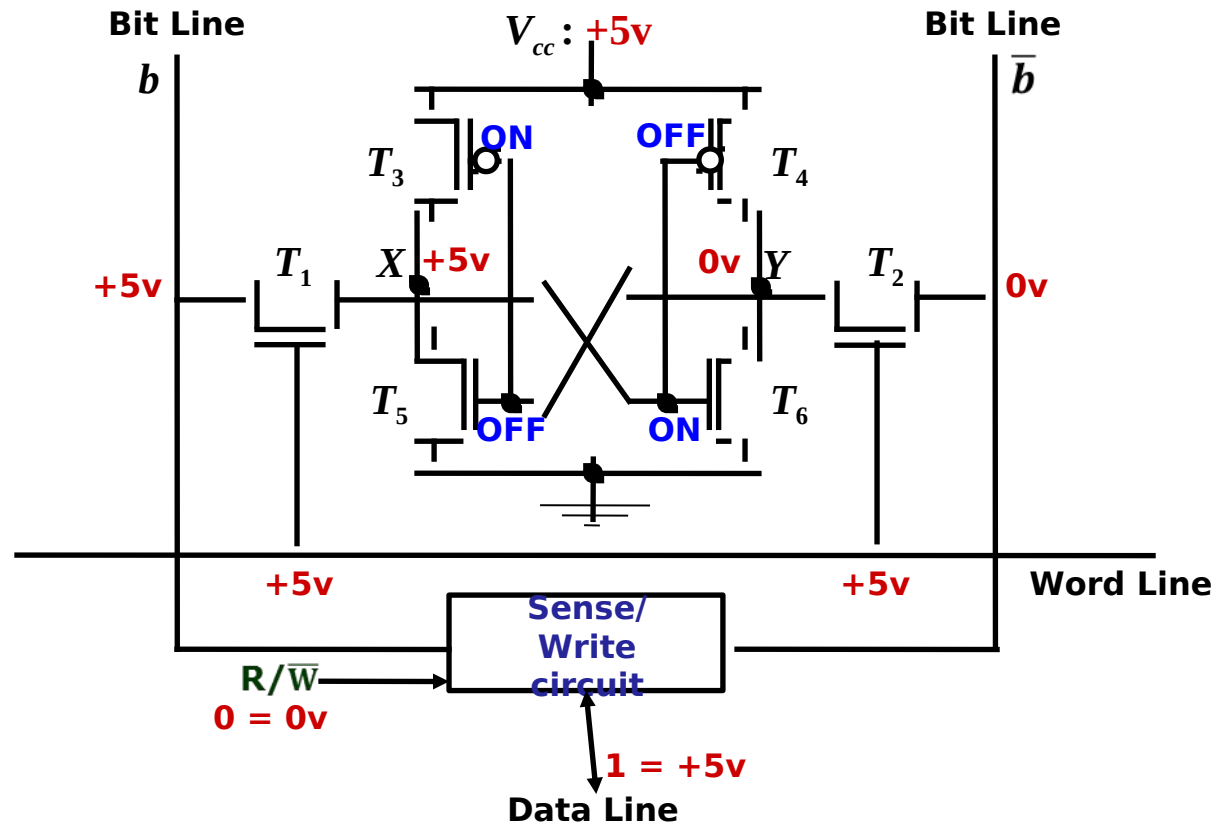
- This circuit retain the state (bit) as long as power is applied (Static Memory)
- Continuous power is needed for a cell to retain the state

CMOS Static RAM Cell - Read



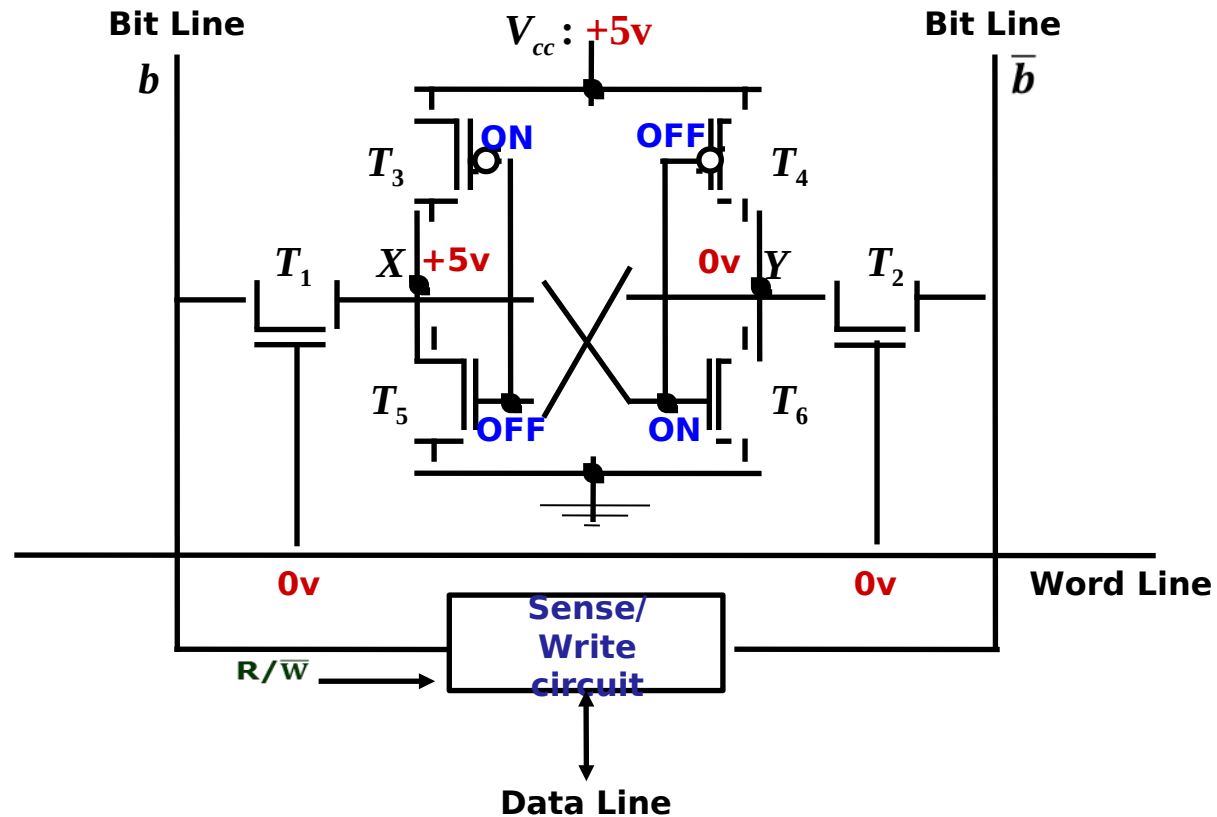
- This circuit retain the state (bit) as long as power is applied (Static Memory)
- Continuous power is needed for a cell to retain the state

CMOS Static RAM Cell - Write



- This circuit retain the state (bit) as long as power is applied (Static Memory)
- Continuous power is needed for a cell to retain the state


CMOS Static RAM Cell: Illustration



- Access time is less i.e. faster memory
- Low power consumption
- Uses 6 transistors: Costly
- Used in applications where speed is critical concern: Cache

Dynamic RAM (DRAM) Cell

- Less expensive and simpler cell
- Information is stored in the form of a charge on a capacitor (C)
 - Charge in capacitor is stored only for short time
 - However, a cell is required to store information for a much longer time
- To retain information for longer time, content of capacitor must be periodically refreshed



- Low speed as refresh needed
- Only 1 transistor is used
- Used to build main memory

Memory Content Example

1024 memory locations: 10 bit address


16 bit data

Memory address		Memory content
Binary	decimal	
0000000000	0	1011010101011101
0000000001	1	1010101110001001
0000000010	2	0000110101000110
	⋮	⋮
1111111101	1021	1001110100010100
1111111110	1022	0000110100011110
1111111111	1023	1101111000100101

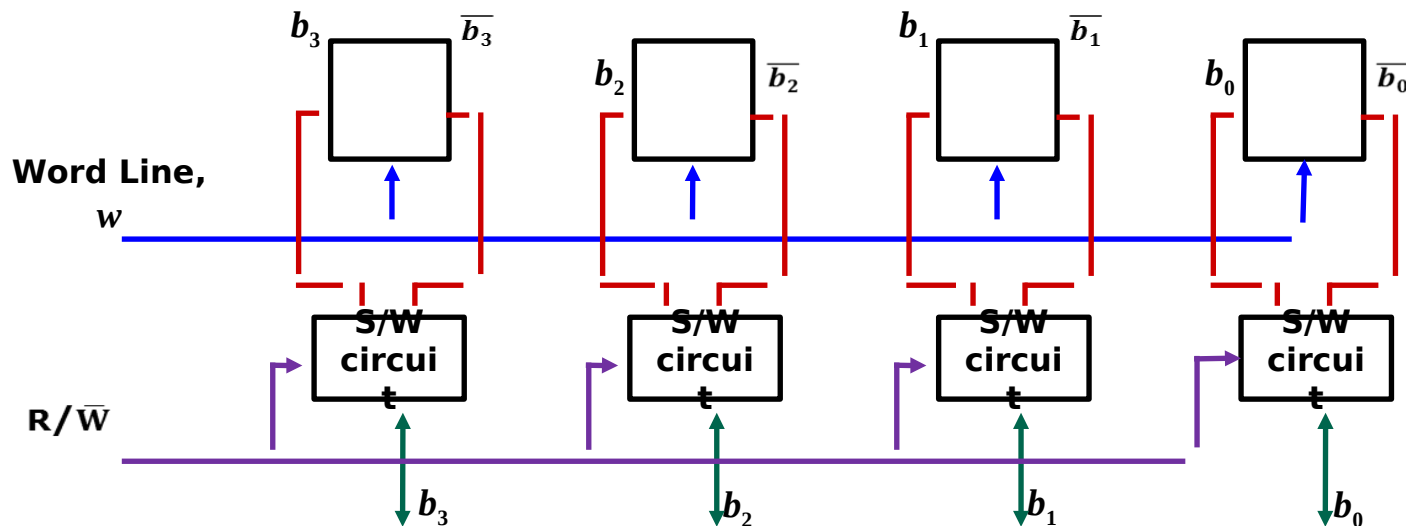
Fig. 7-3 Content of a 1024×16 Memory

Internal Organization of SRAM Chip

- SRAM cell (block diagram):



- Memory cells are usually organised in the form of an **array**
- At each memory location, n SRAM cells are placed next to each other to form **n -bit memory word**
- Example: **4-bit memory word**



Internal Organization of $2^k \times n$ SRAM Chip

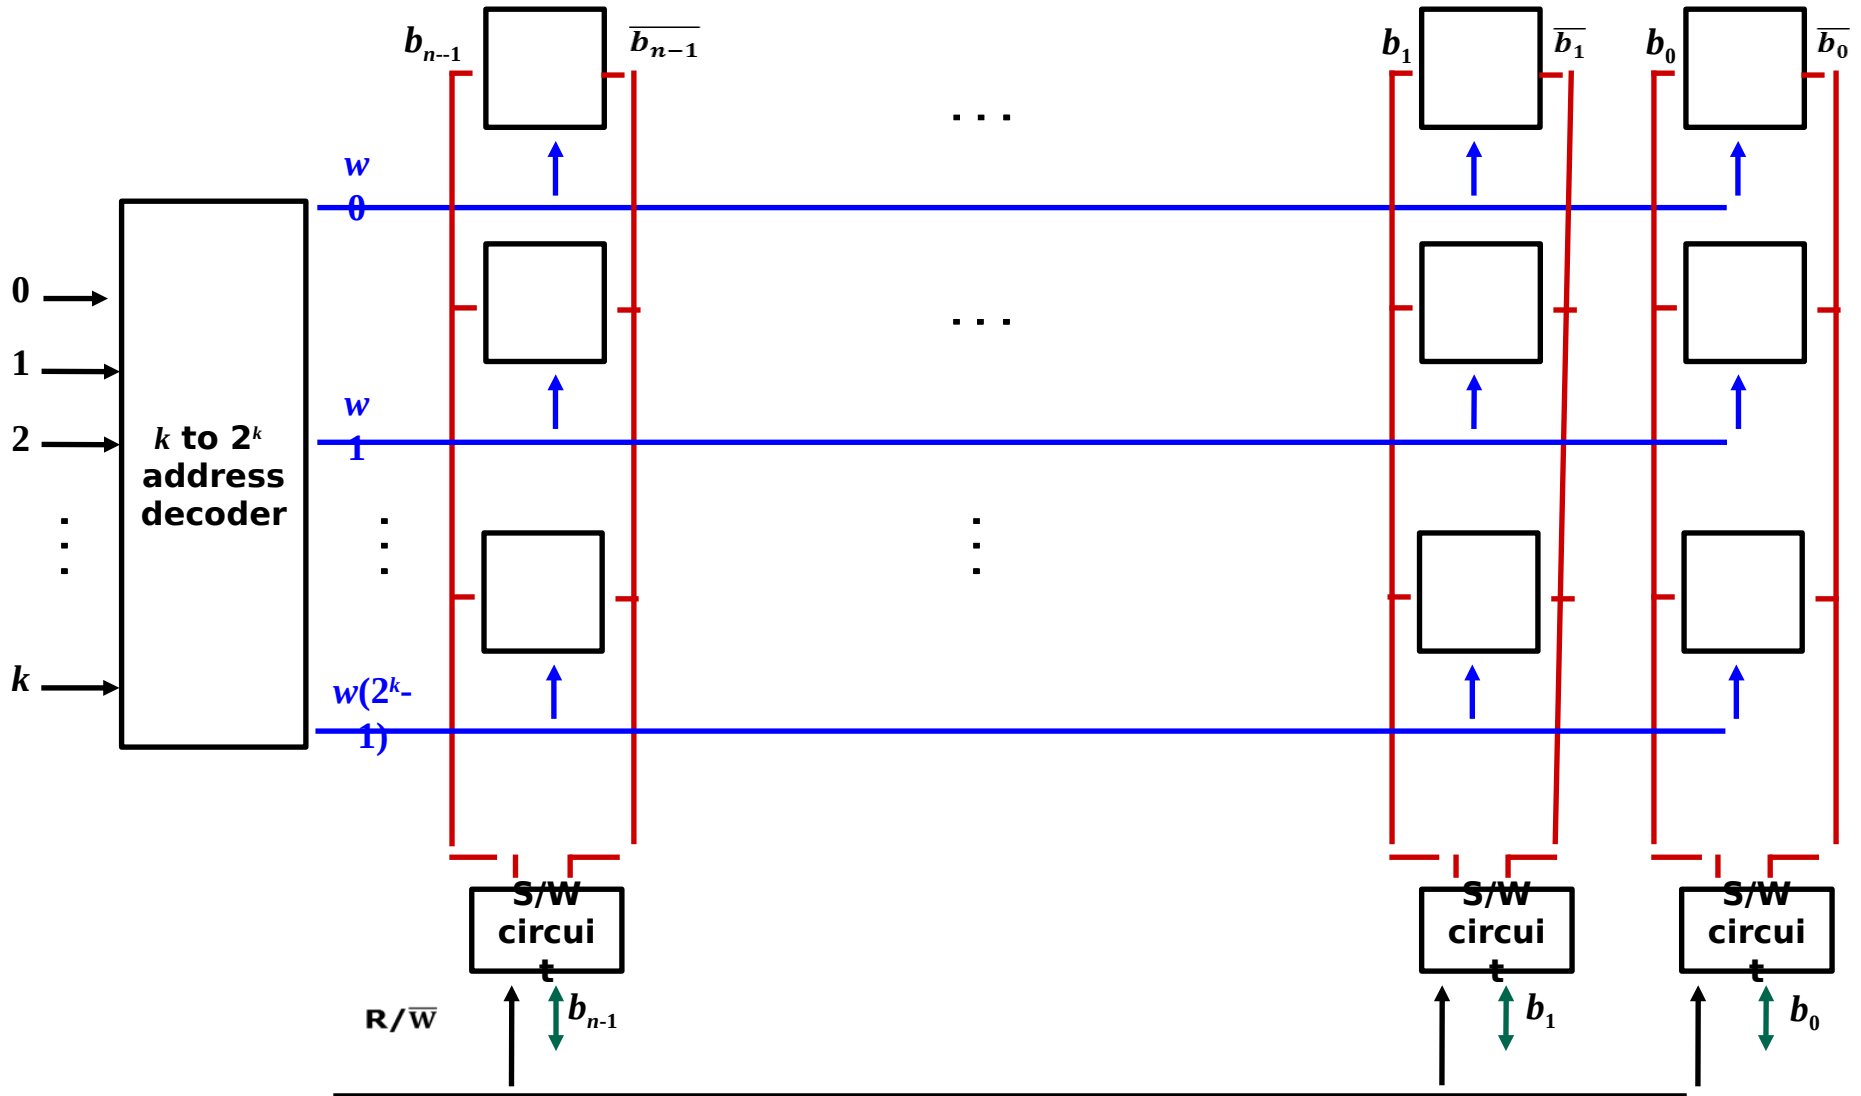


Illustration: 16 x 4 bit Memory Chip

- 1-dimensional address decoding:
 - Example: Design of 64b memory chip
 - Number of address line be : 4
 - Address decoder configuration : 4-to- 2^4

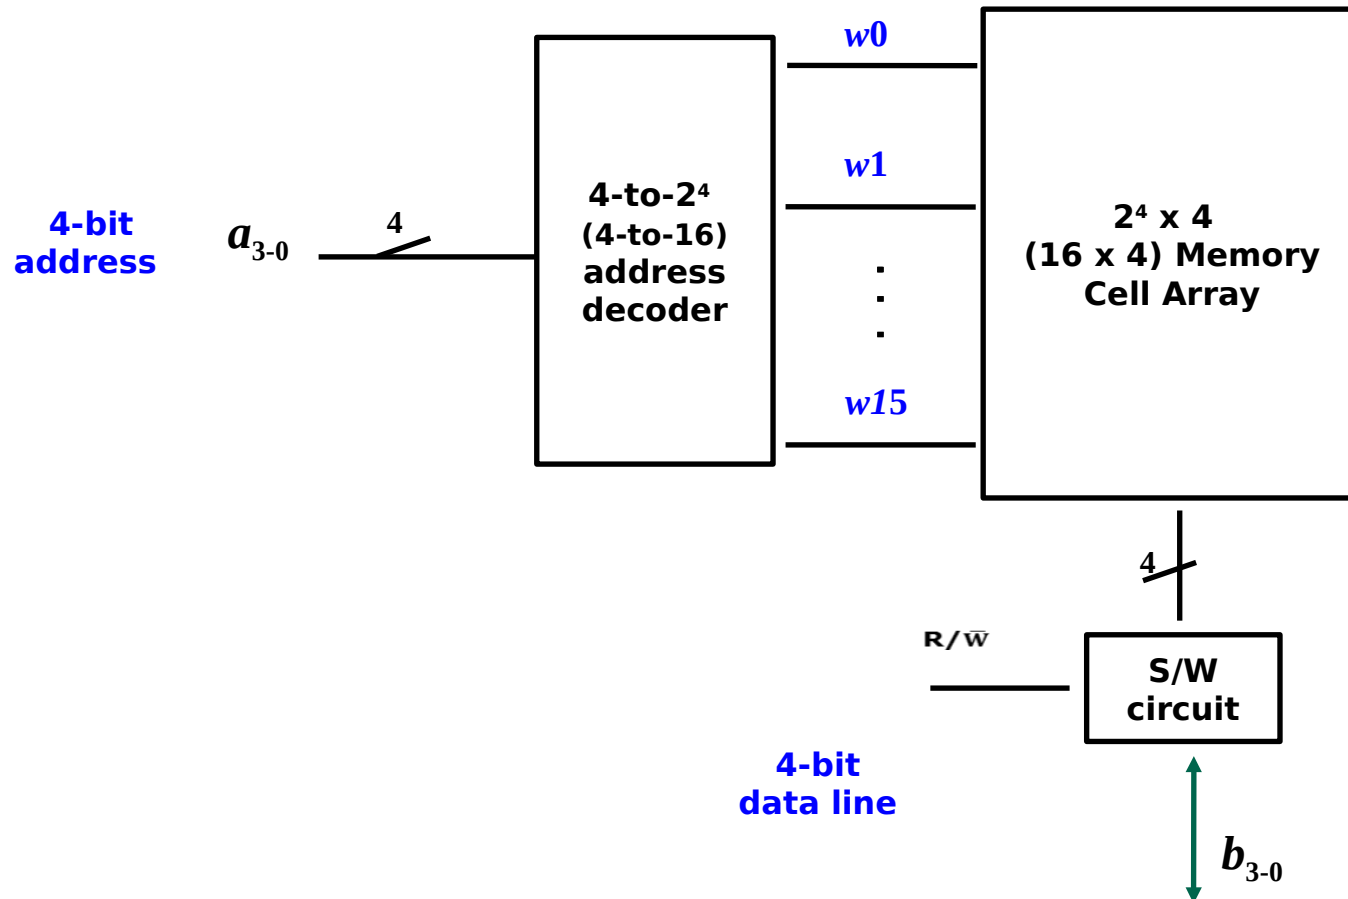


Illustration: 512 x 8 bit Memory Chip

- 1-dimensional address decoding:

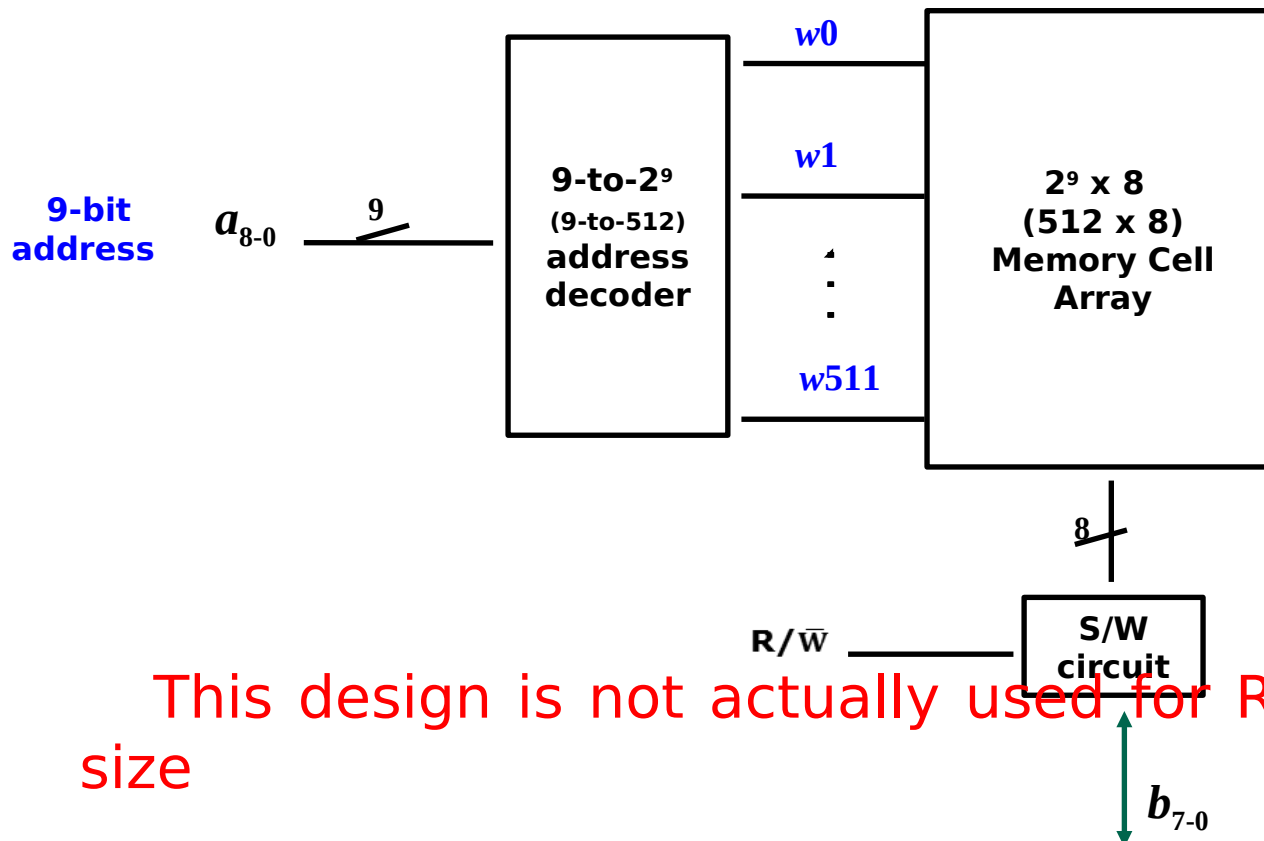
– *Example:*

Memory chip organization : 512 x 8

Memory Size : 512B

Number of address line be: 9

Address decoder configuration : 9-to- 2^9



Coincident Decoding

Regular decoding is costly:

A decoder with k inputs and 2^k outputs requires 2^k AND gates with k inputs per gate.

Total number of gates can be reduced by using two-dimensional decoding:

Basic idea: arrange memory cells in a (as close as possible to) square configuration.

Use two $k/2$ input decoders instead of one k input decoder

Two-Dimensional Decoding

Instead of using a single 10 x 1024 decoder
we use two 5x32 decoders.

One decoder picks the row, one the column




Fig. 7-7 Two-Dimensional Decoding Structure for a 1K-Word Memory

Two-Dimensional Decoding

Needs 64 5-input AND gates instead of 1024
10-input gates.

Address is divided to two equal parts


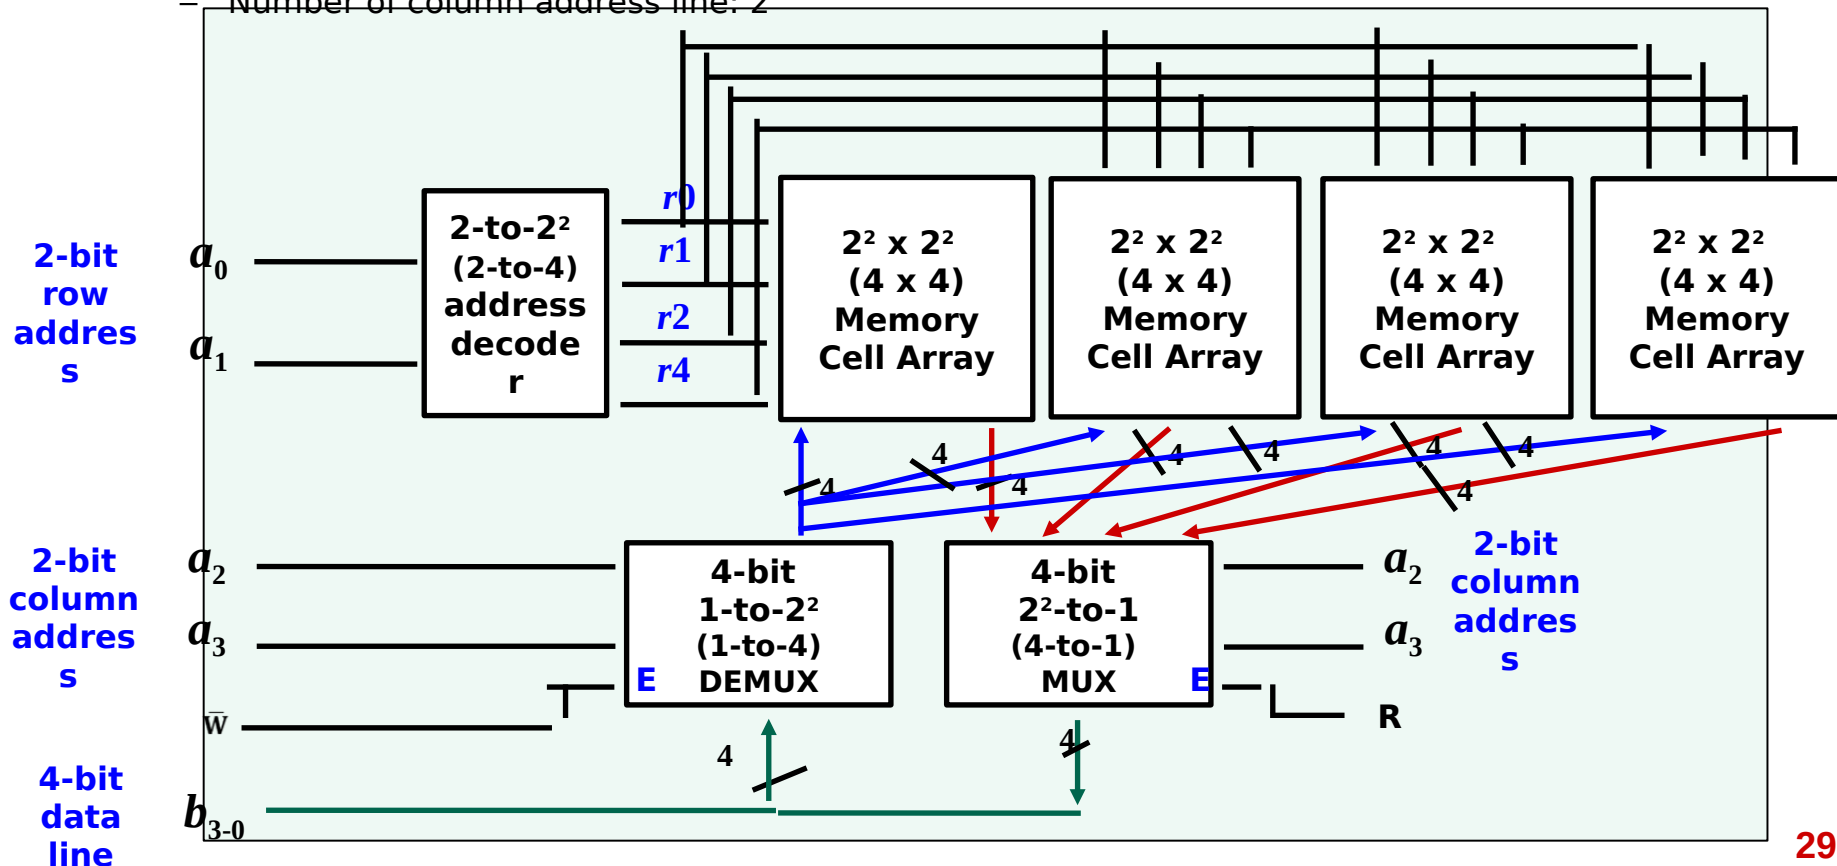


Fig. 7-7 Two-Dimensional Decoding Structure for a 1K-Word Memory

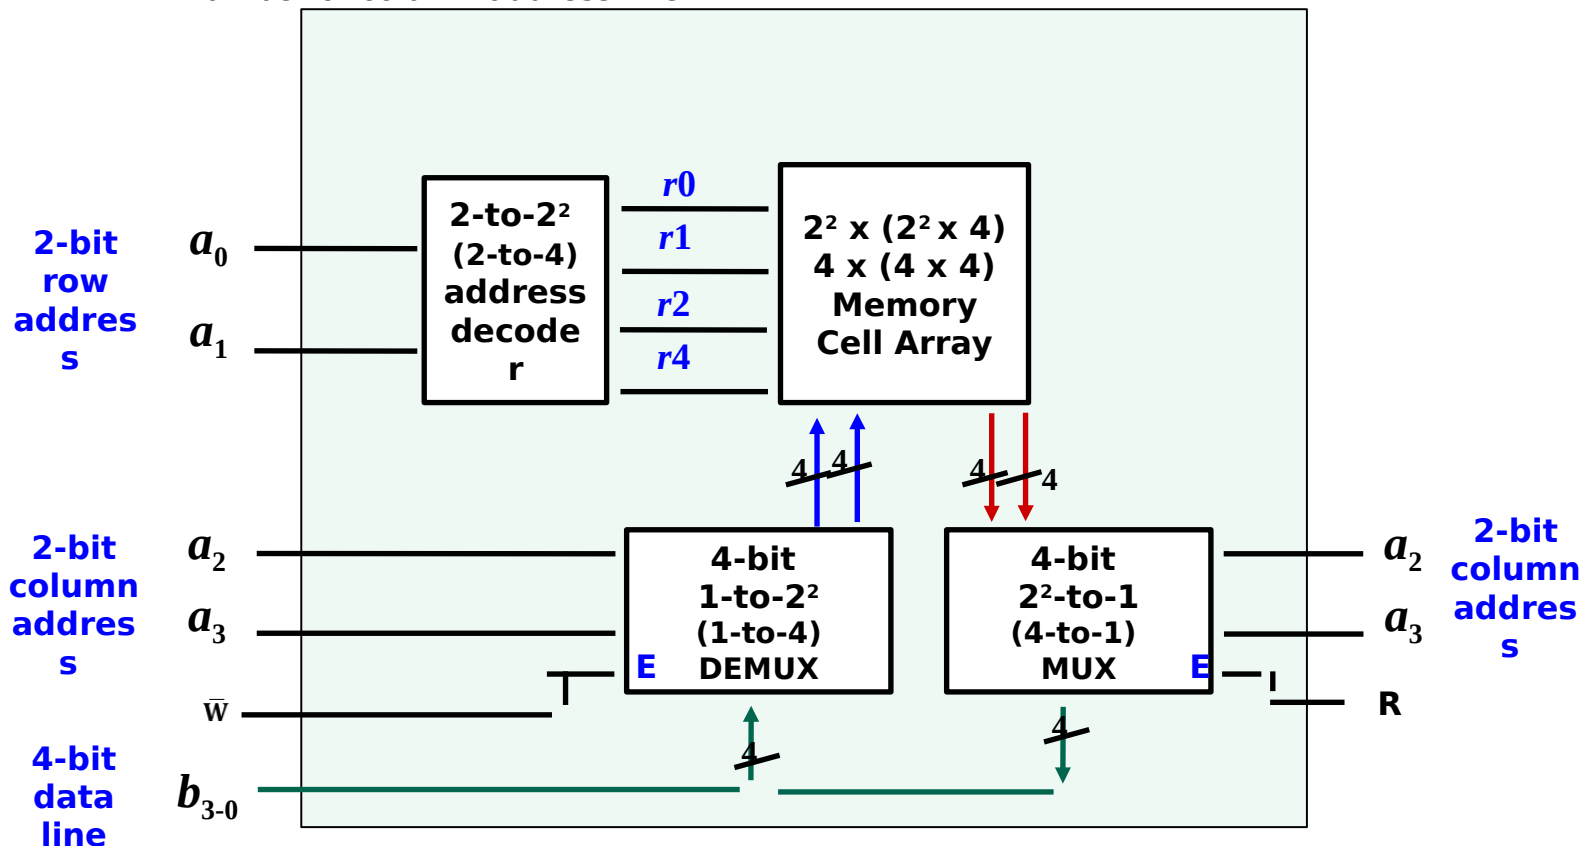
2-dimensional Address Decoding

- k -address line is divided into $k/2$ row address and $k/2$ column address
- Now, memory chip is considered as $2^{k/2} \times 2^{k/2}$ memory cell array
- **Example:** Design of 16 x 4 bit memory chip
 - Number of address lines: 4
 - Number of row address lines: 2
 - Number of column address line: 2



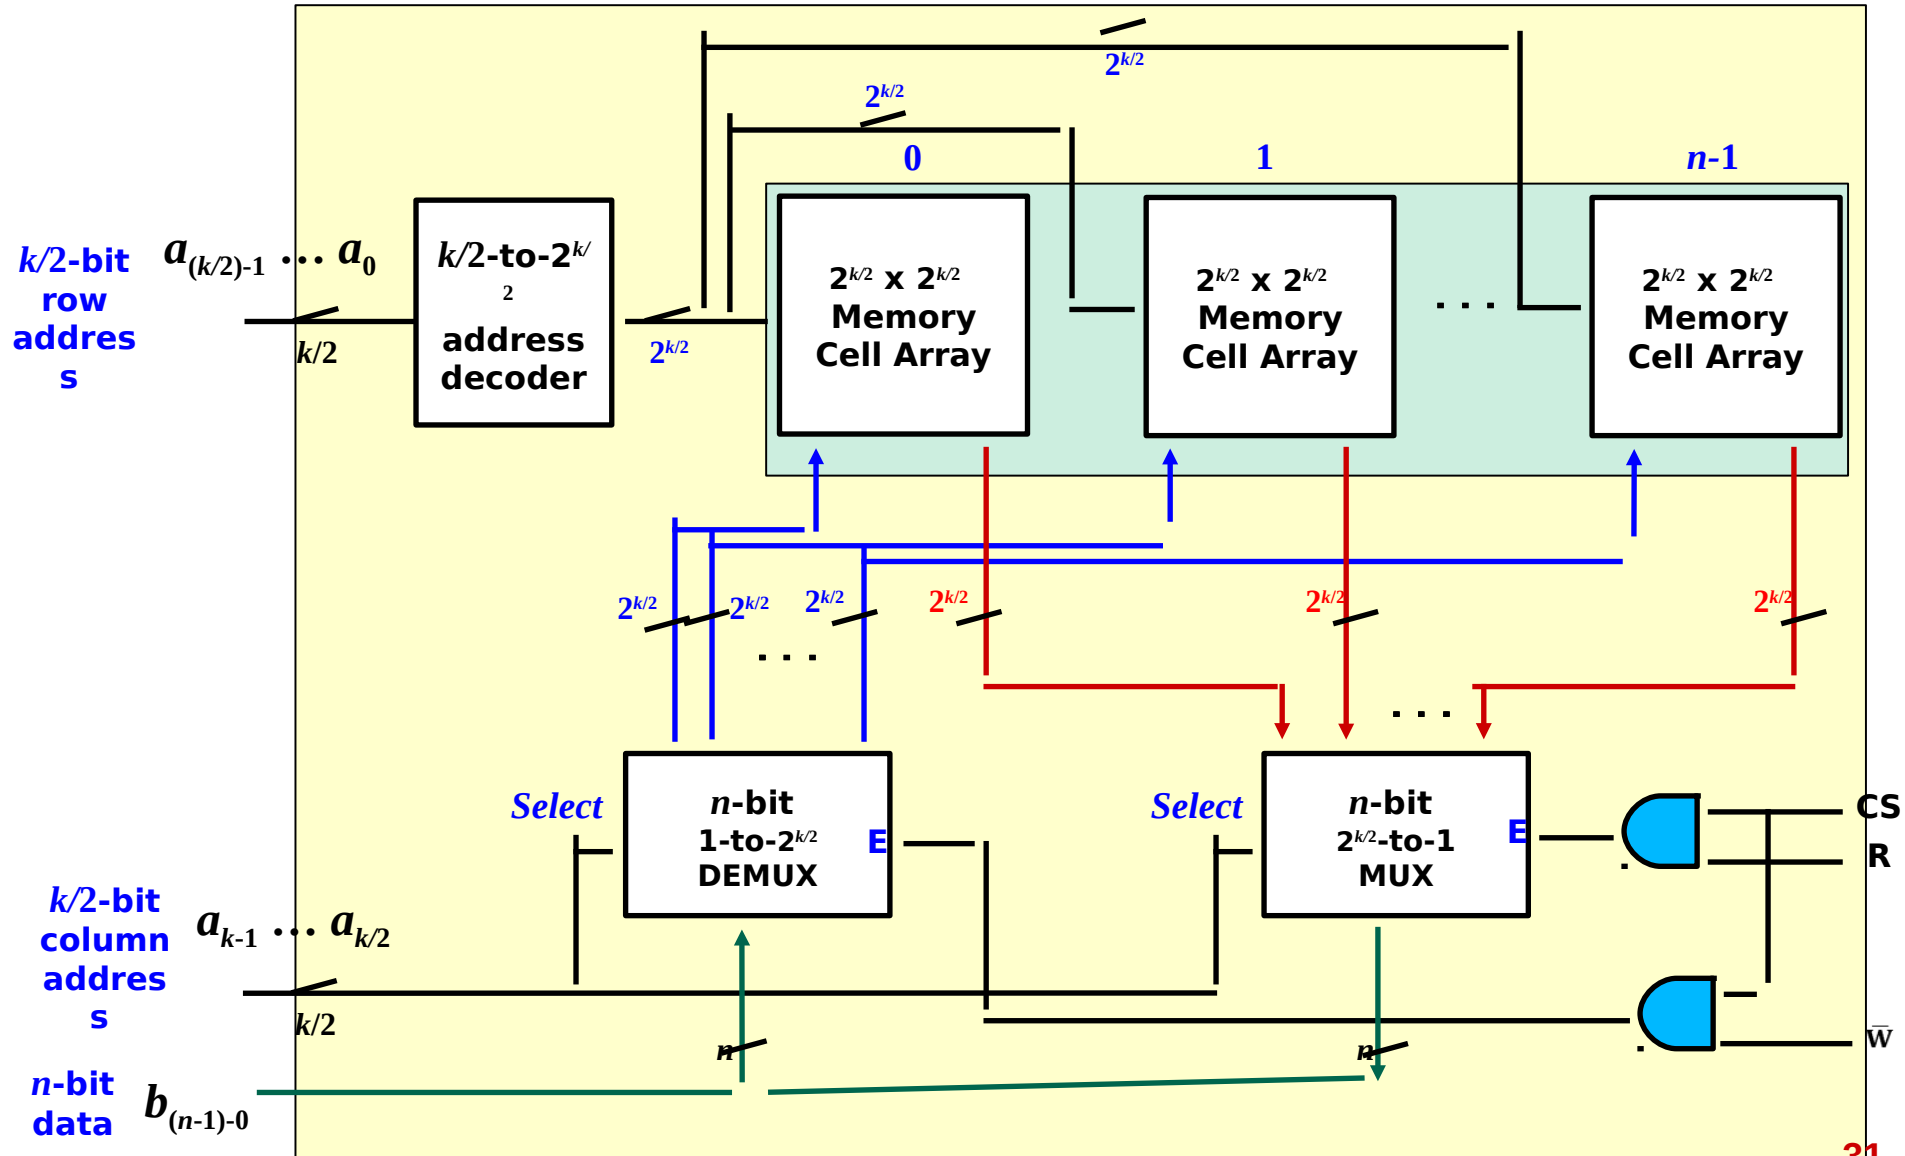
2-dimensional Address Decoding

- k -address line is divided into $k/2$ row address and $k/2$ column address
- Now, memory chip is considered as $2^{k/2} \times 2^{k/2}$ memory cell array
- Example: Design of 16 x 4 bit memory chip
 - Number of address lines: 4
 - Number of row address lines: 2
 - Number of column address line: 2



2-dimensional Address Decoding

- Organization of $2^k \times n$ SRAM Chip



2-dimensional Address Decoding

- Organization of $2^k \times n$ SRAM Chip

- The number of cell array depends on the width of the memory chip (word length) expected

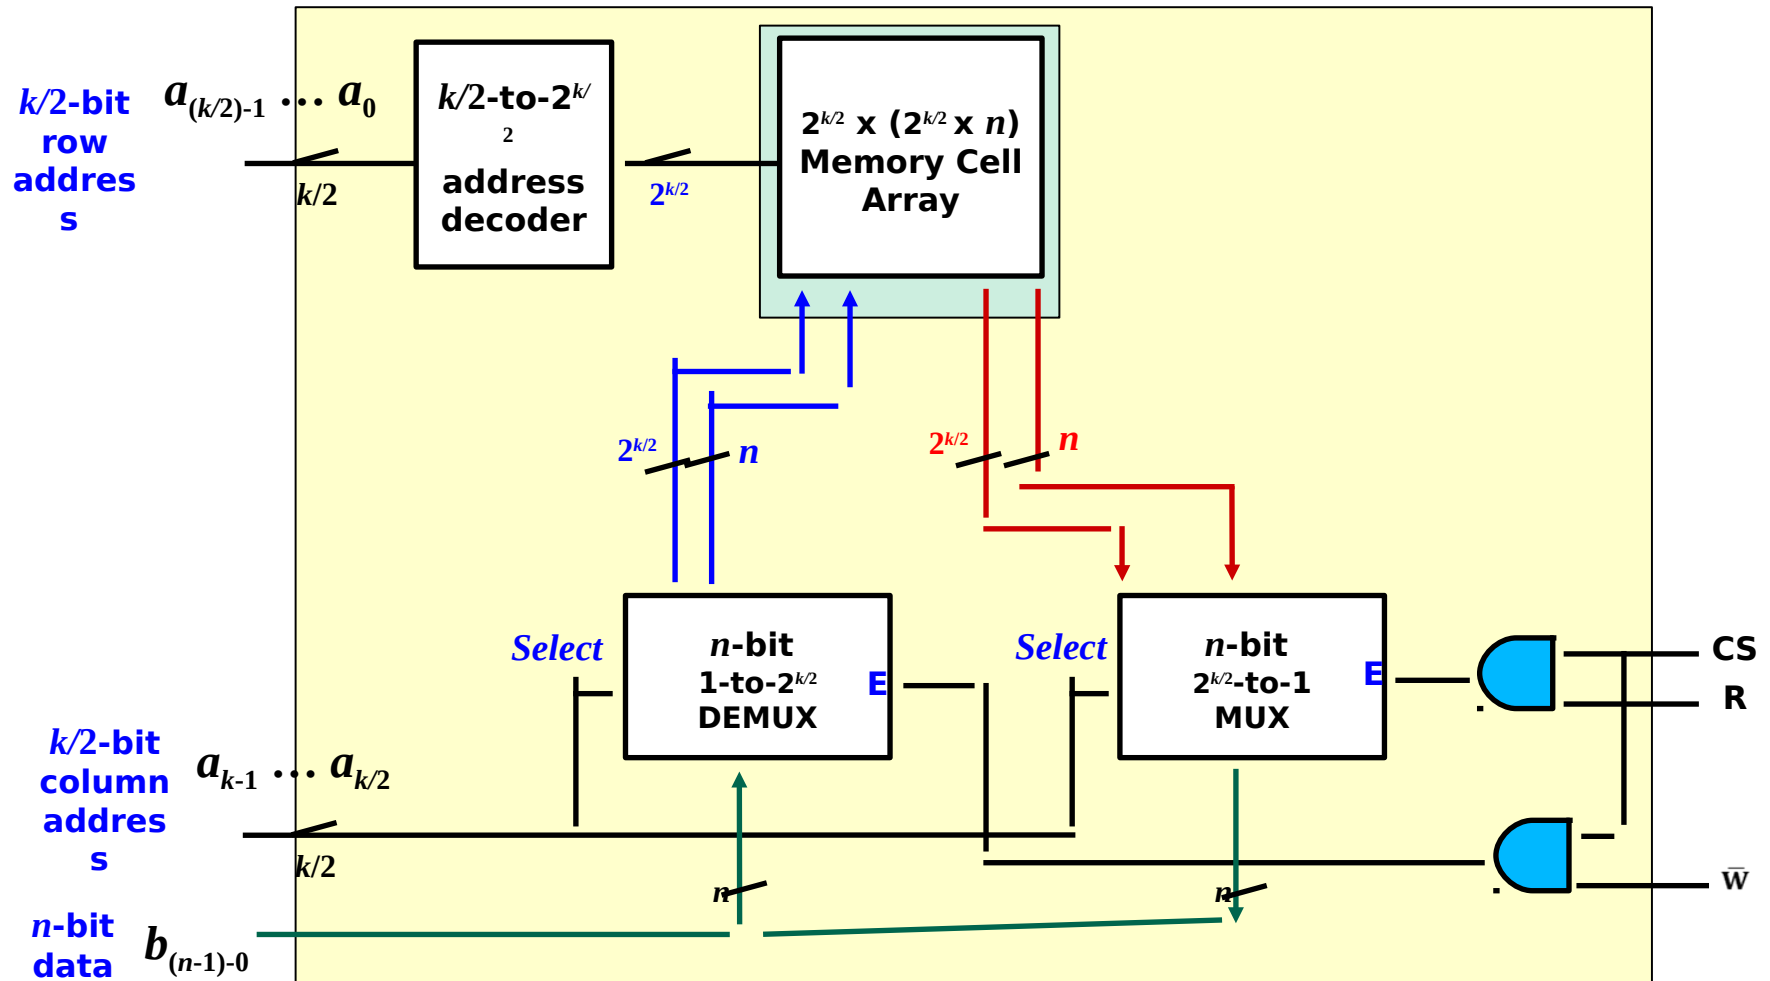


Illustration: Design of 1MB SRAM Chip

- Capacity of memory: 1 MB
- Word length: 8 bits
- Memory organization: 1M x 8

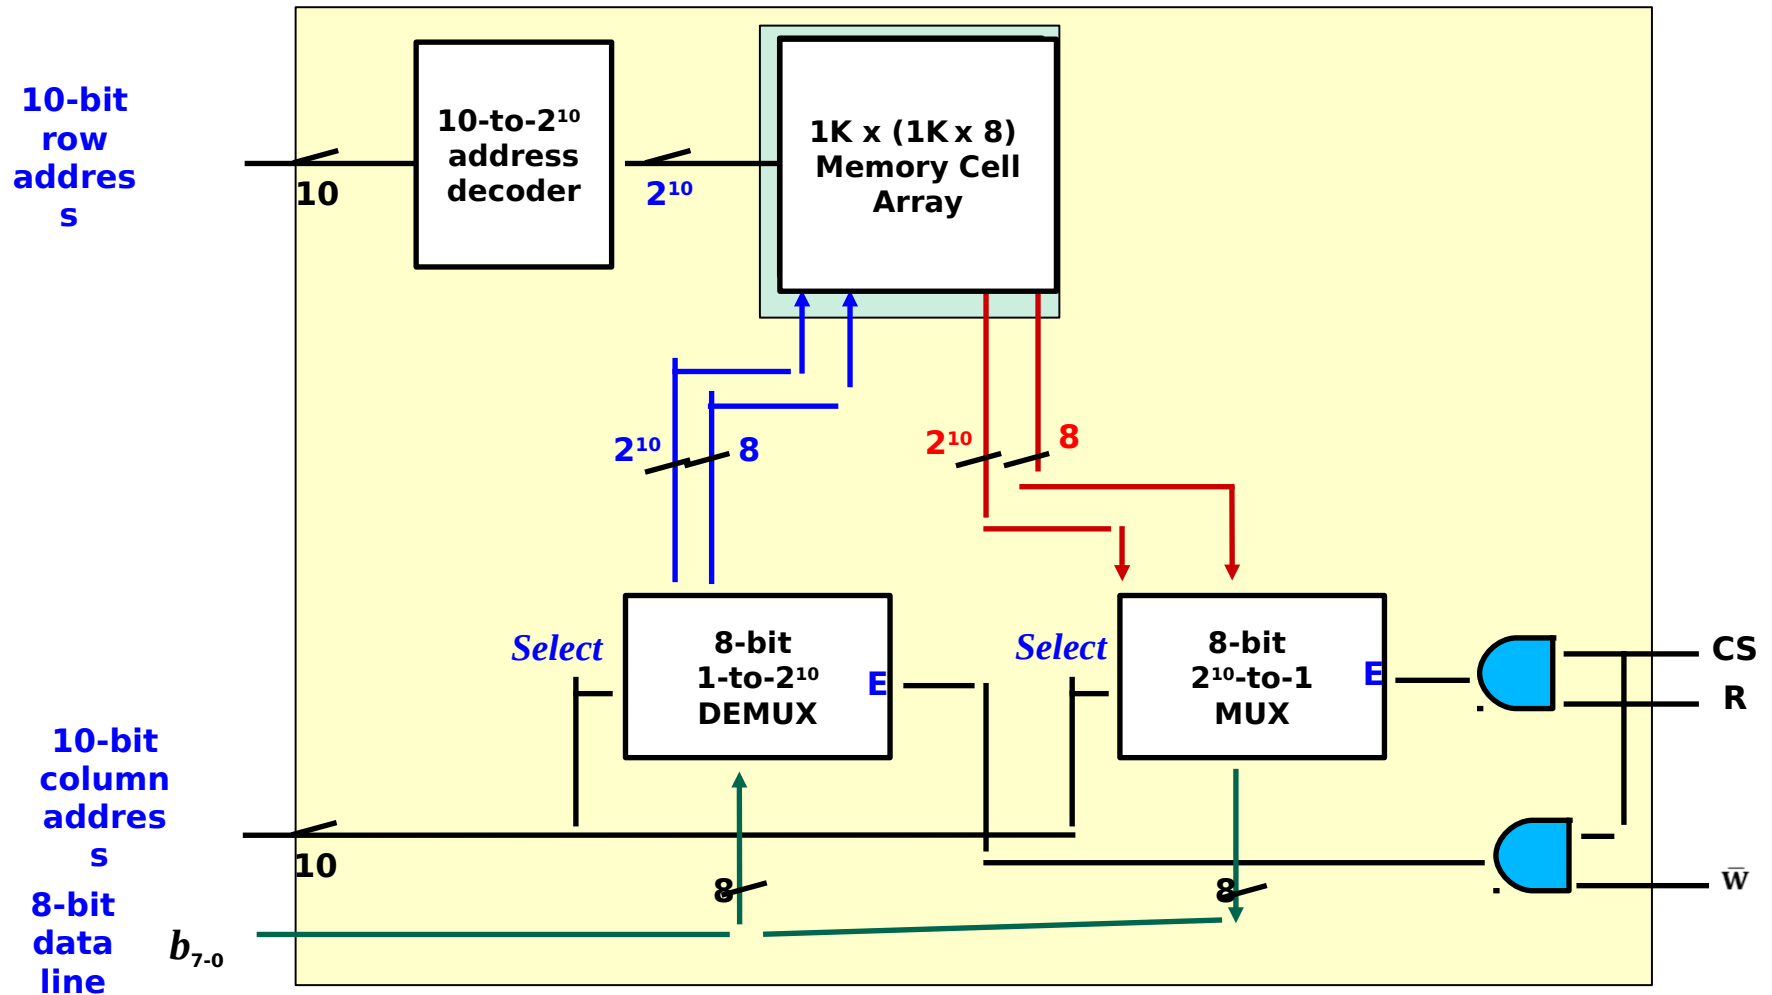
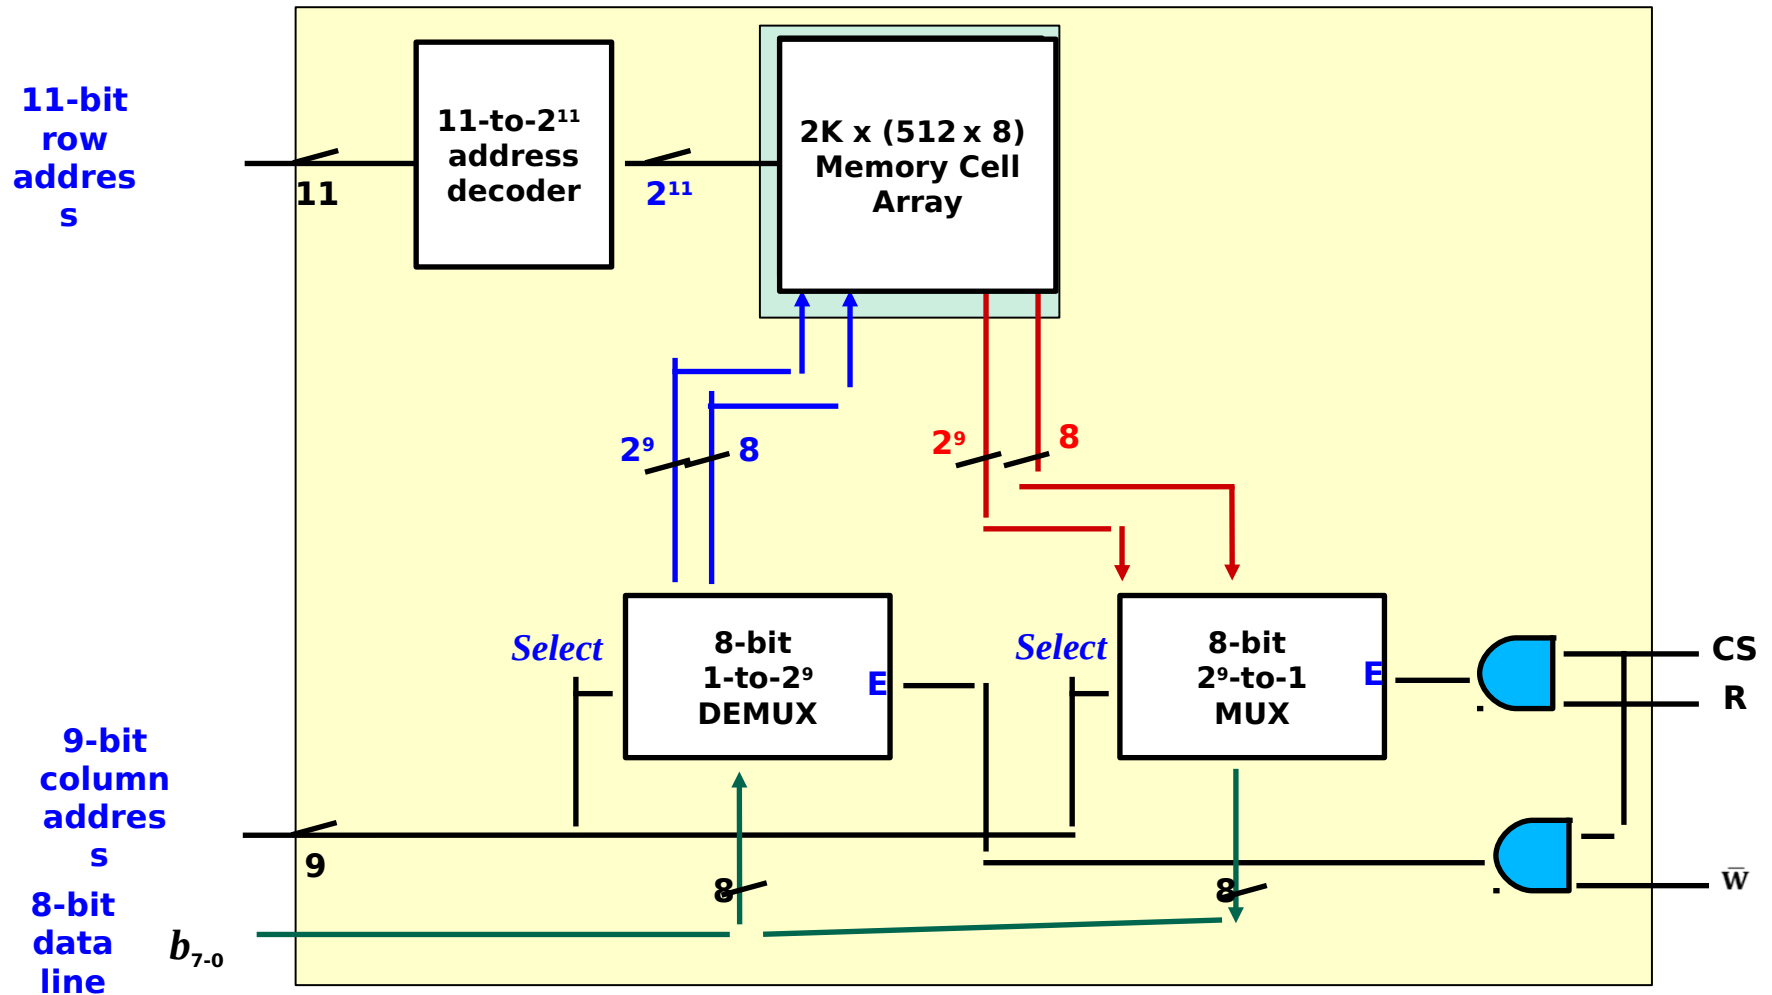


Illustration: Design of 1MB SRAM Chip

- Capacity of memory: 1 MB ($2^{20} \text{ B} = 2^{23} \text{ b}$)
- Word length: 8 bits
- Number of row address: 11
- Number of column address: $20 - 11 = 9$



Static RAM Module

- **Byte Addressable Memory**

- Illustration:

- **Task:** Design SRAM with capacity 8MB **(2^{23} B)**
- **Requirements:**
 - Memory organization depends on the **word size** and word size decides the width of the data bus
 - Let word size be 32 bit **(4 B)**
 - Now, memory organization: **2M x 32 bit**
 - Memory should be byte addressable
 - Let the organization of cell array to incorporate byte addressability be **1M x 8 bit**
- Organization include:
 - 2 rows of chips, each of size **1M x 8 bit**
 - Each row contain **4 chips**
- Number of address lines: **23** **(2^{23} B)**

Illustration: 8MB Static RAM Module

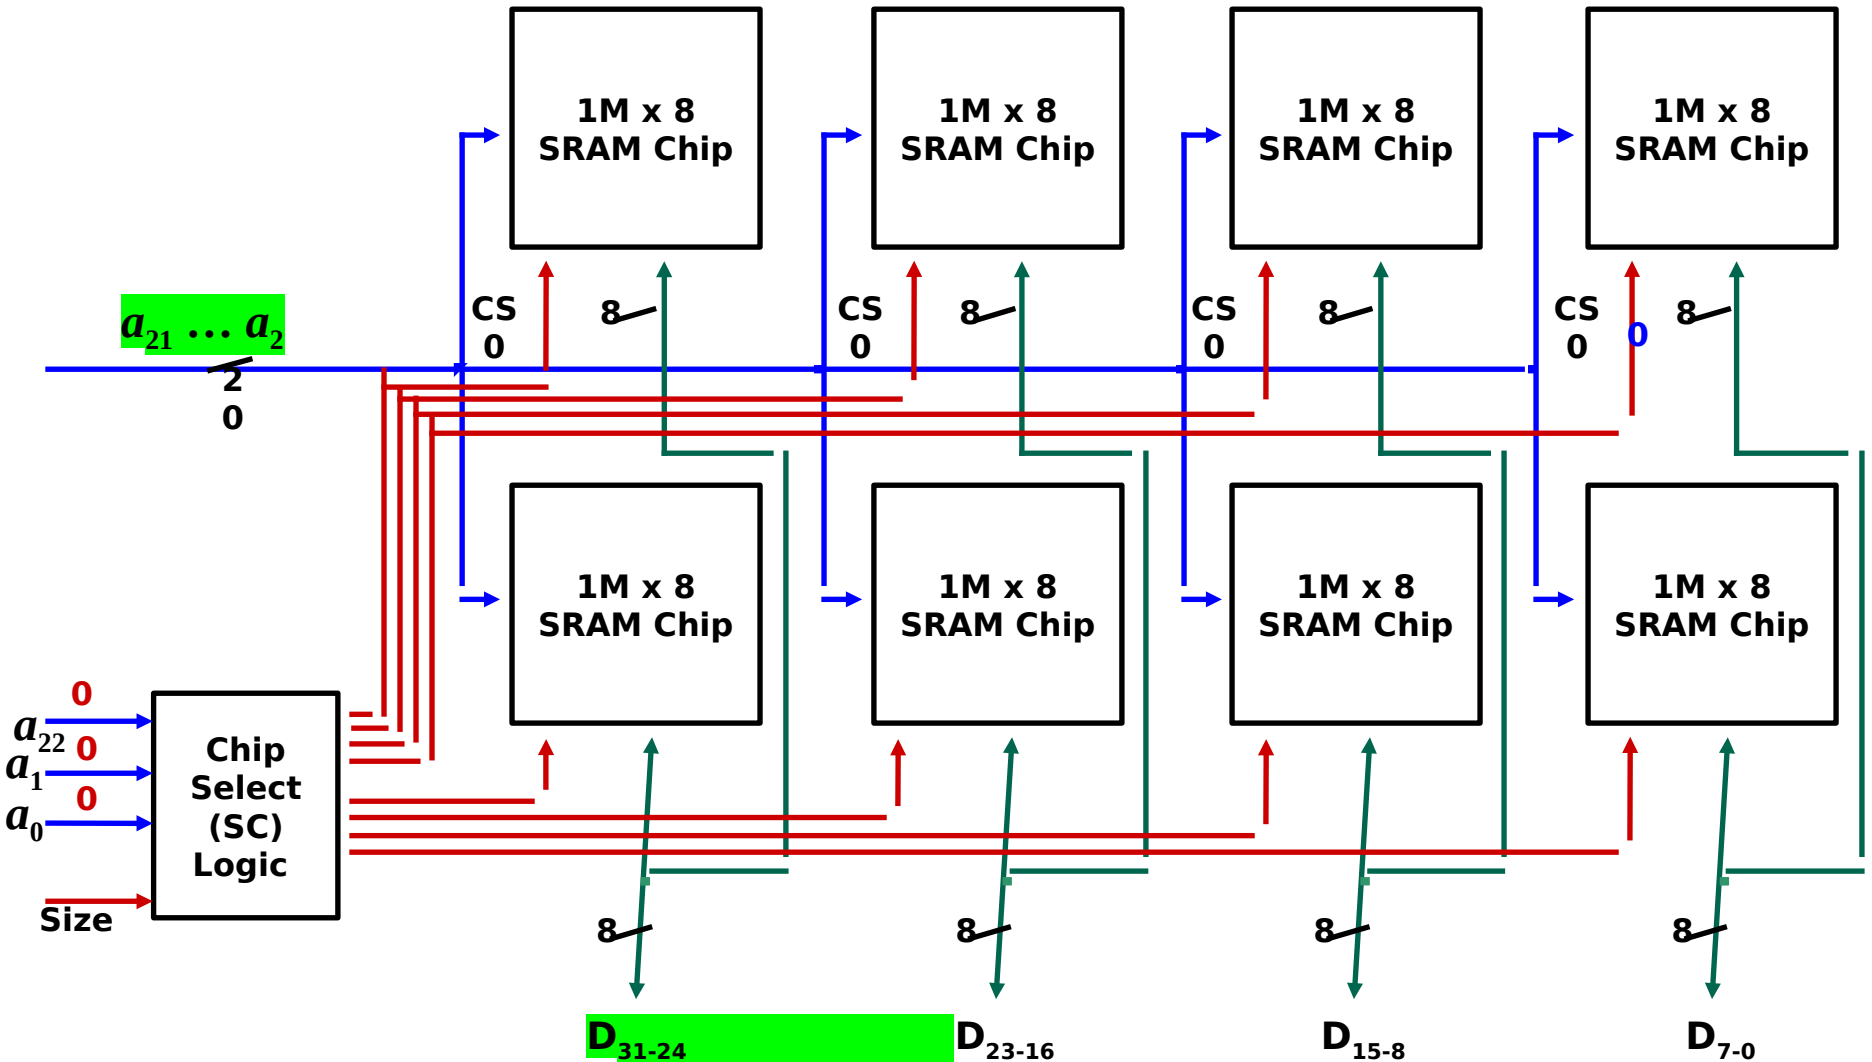


Illustration: 8MB Static RAM Module

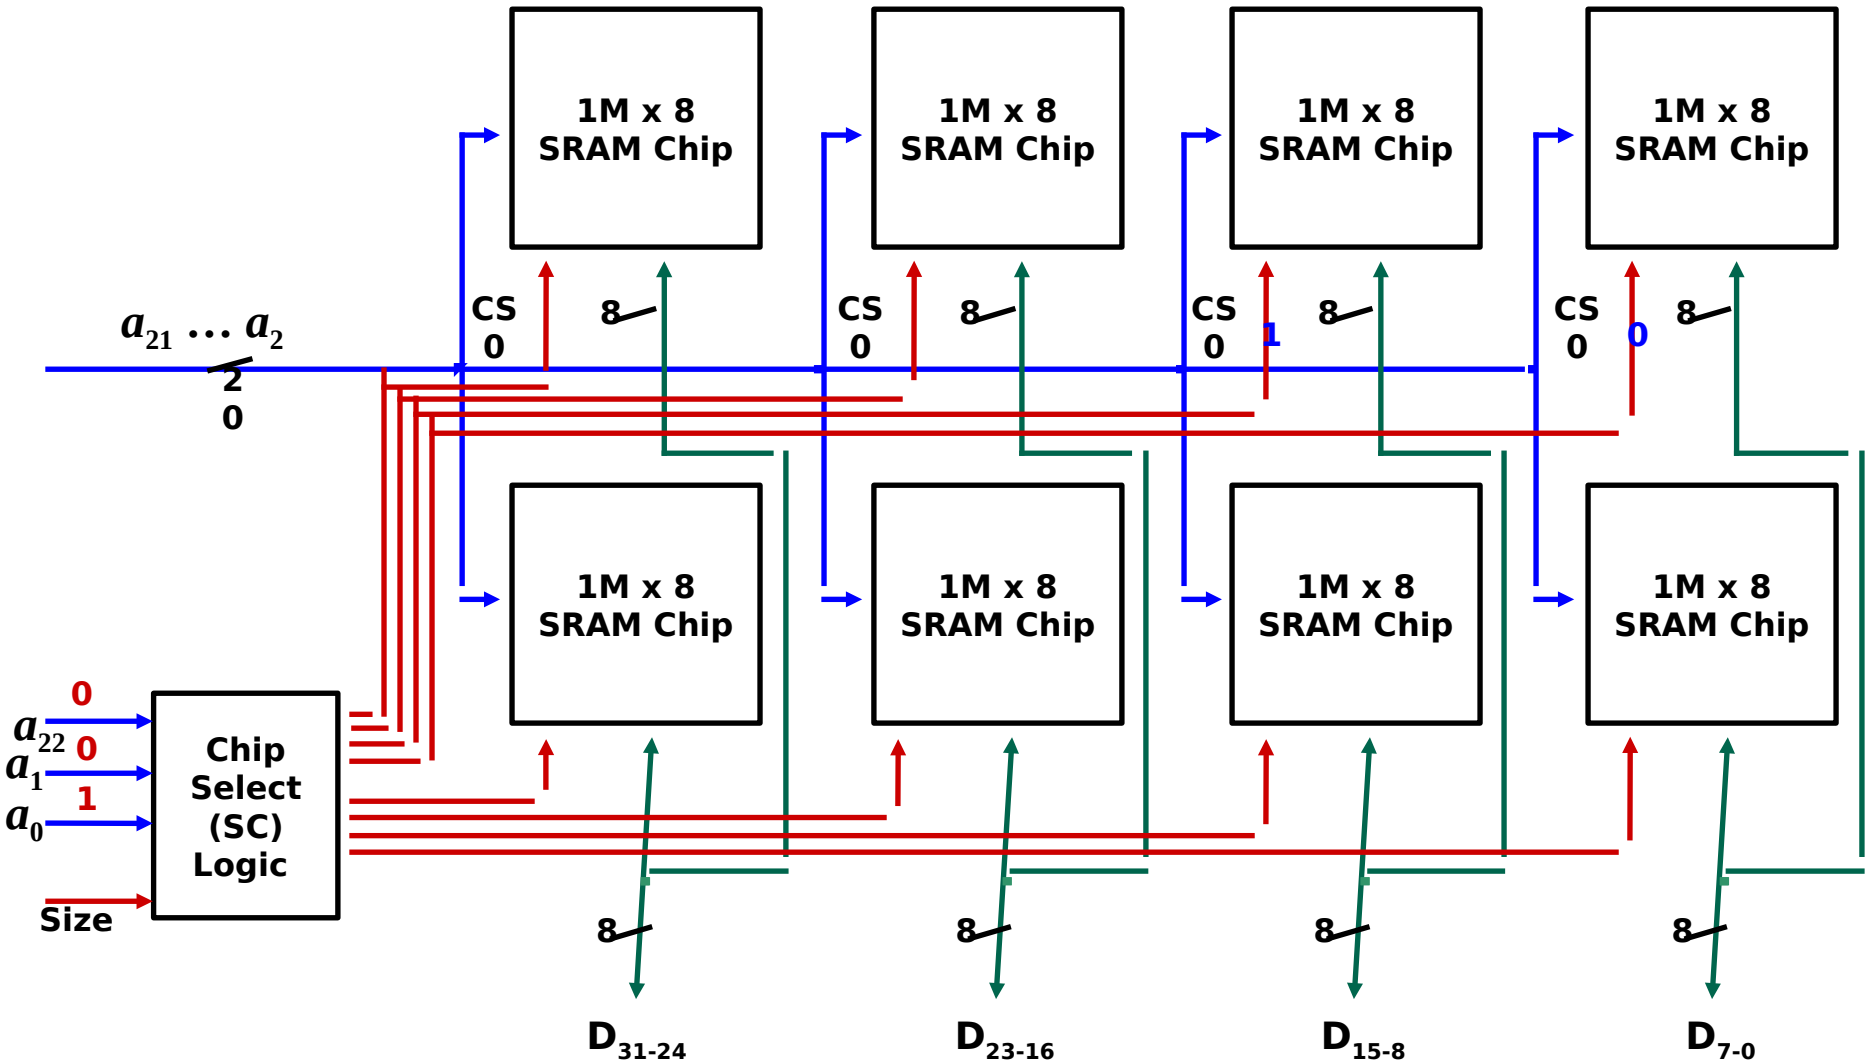


Illustration: 8MB Static RAM Module

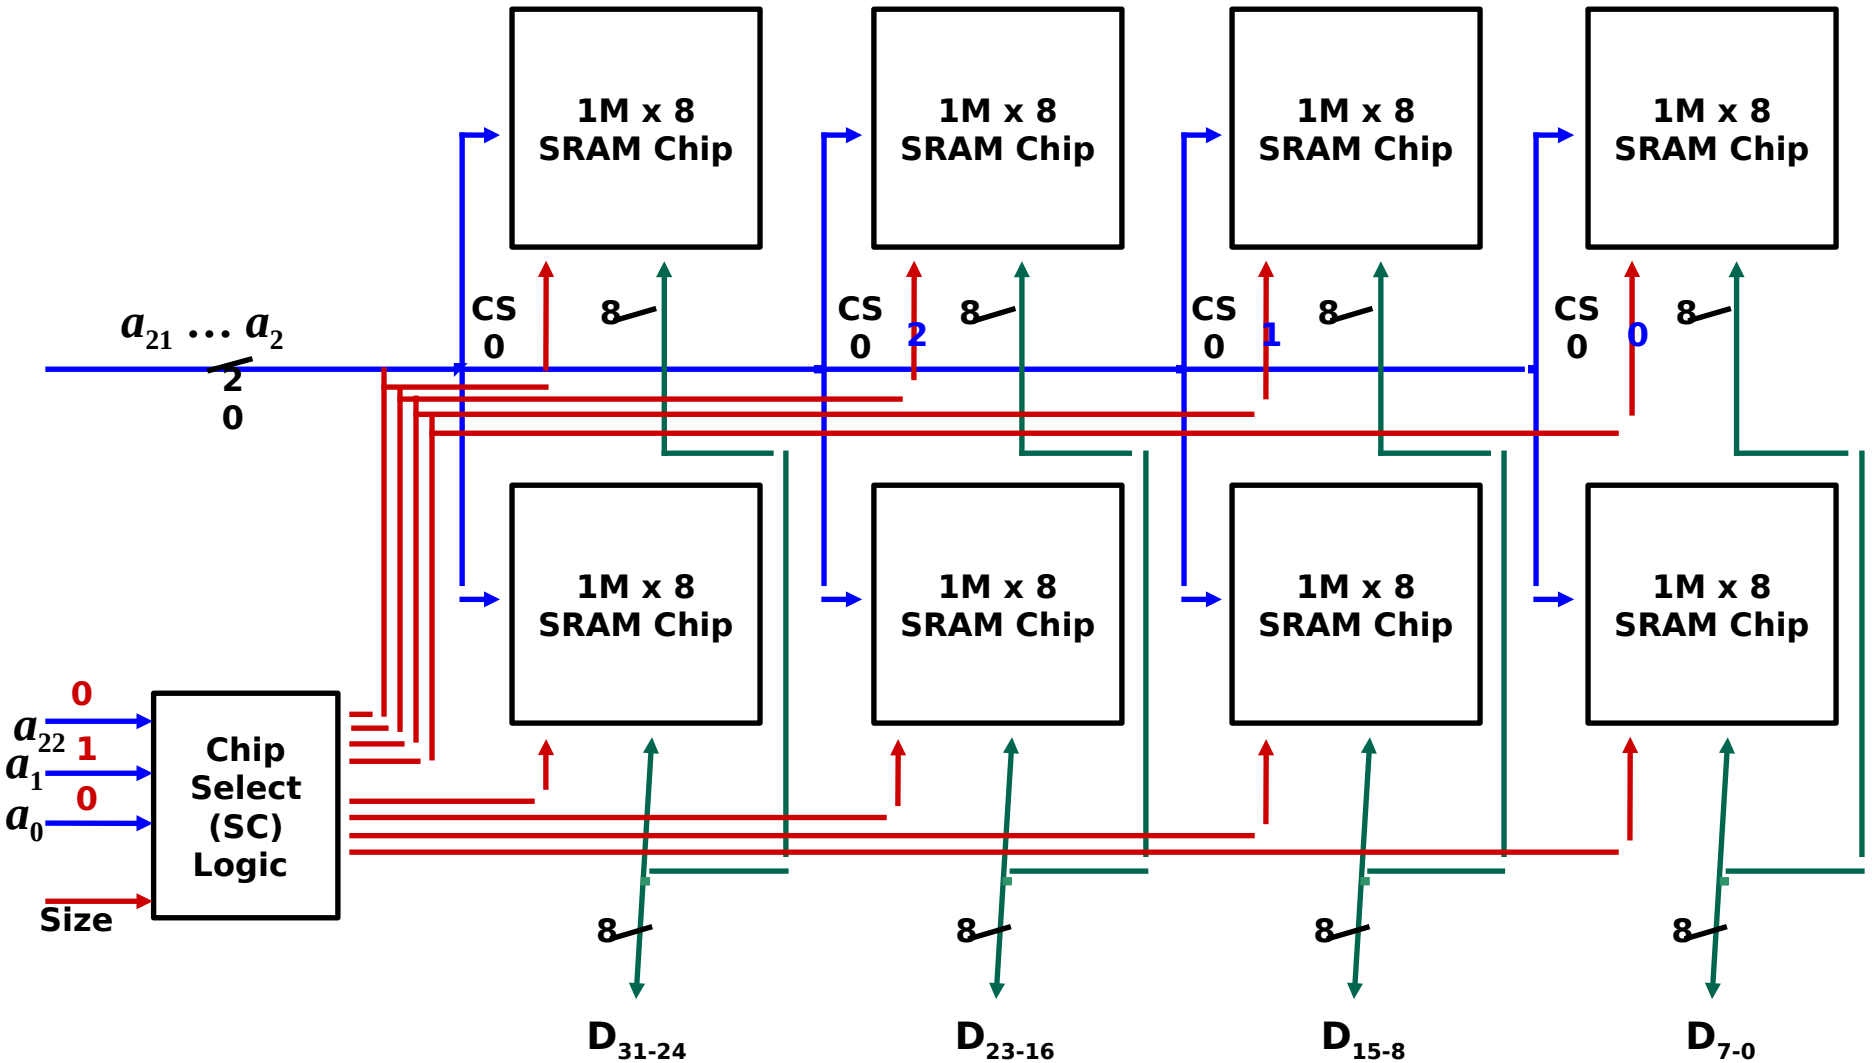


Illustration: 8MB Static RAM Module

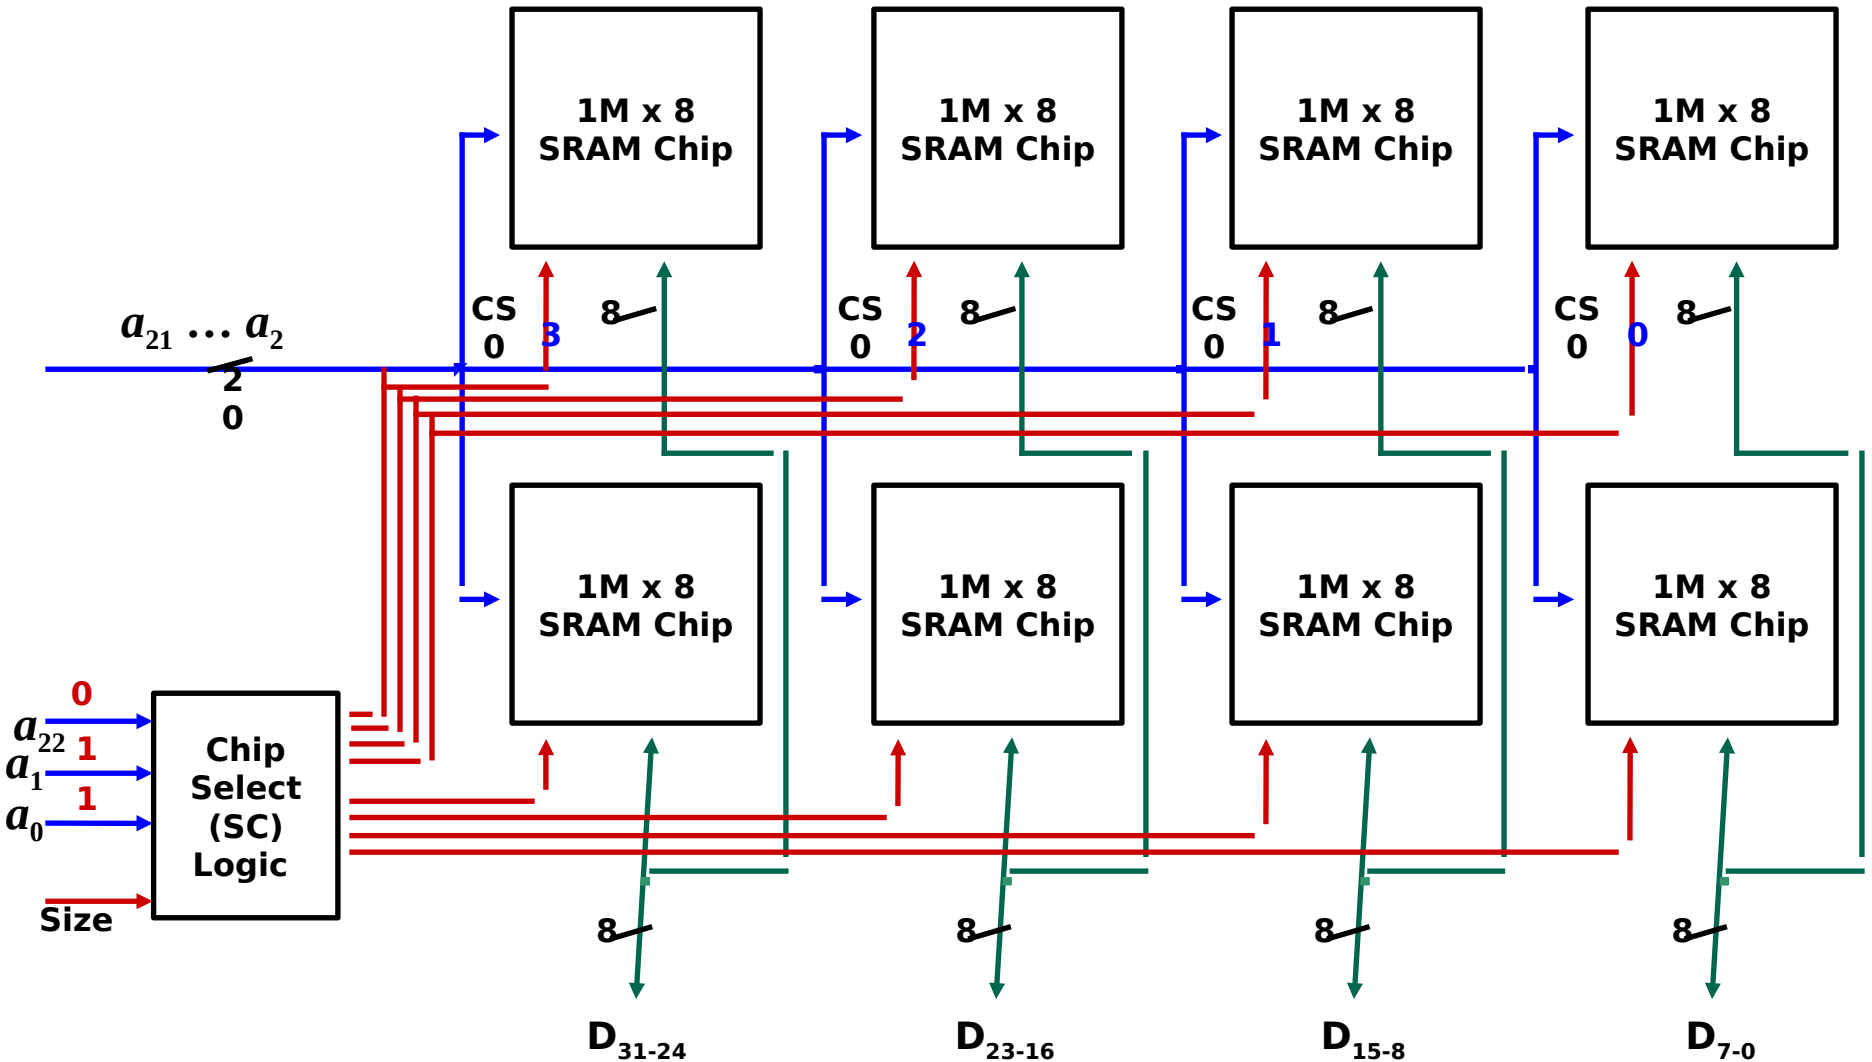


Illustration: 8MB Static RAM Module

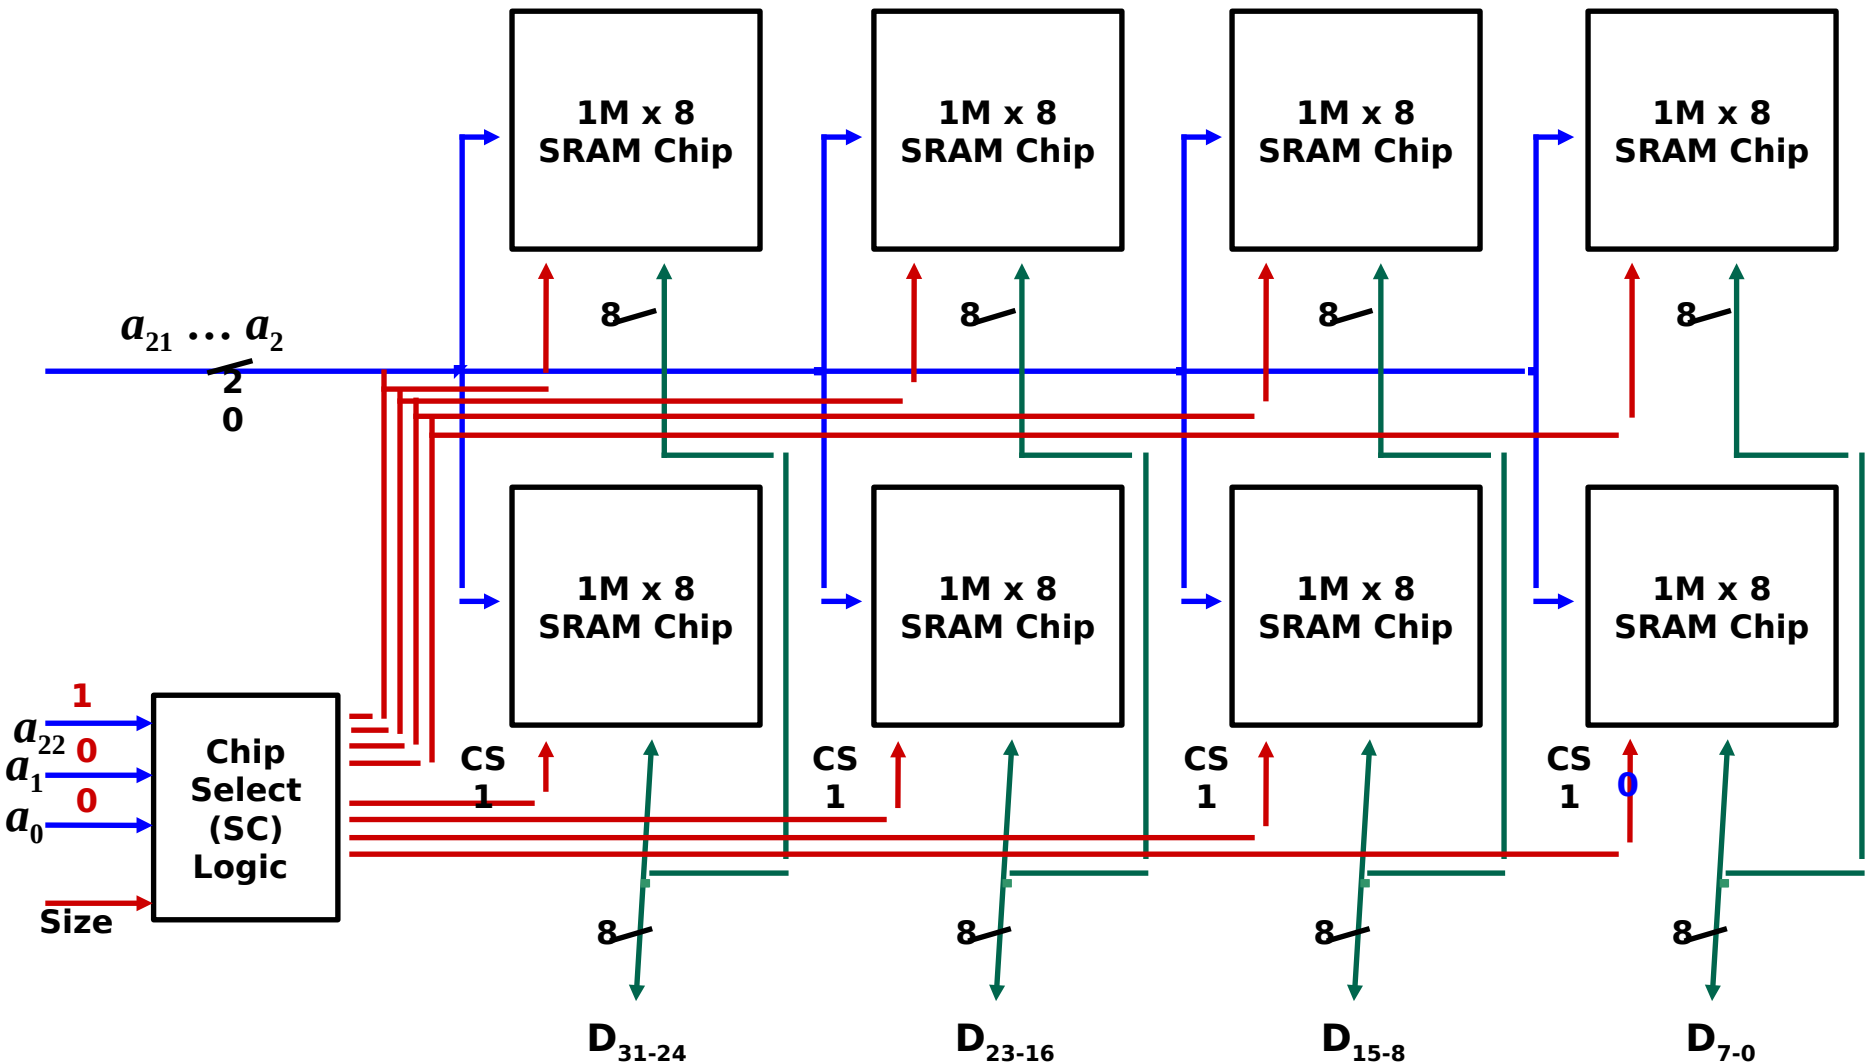


Illustration: 8MB Static RAM Module

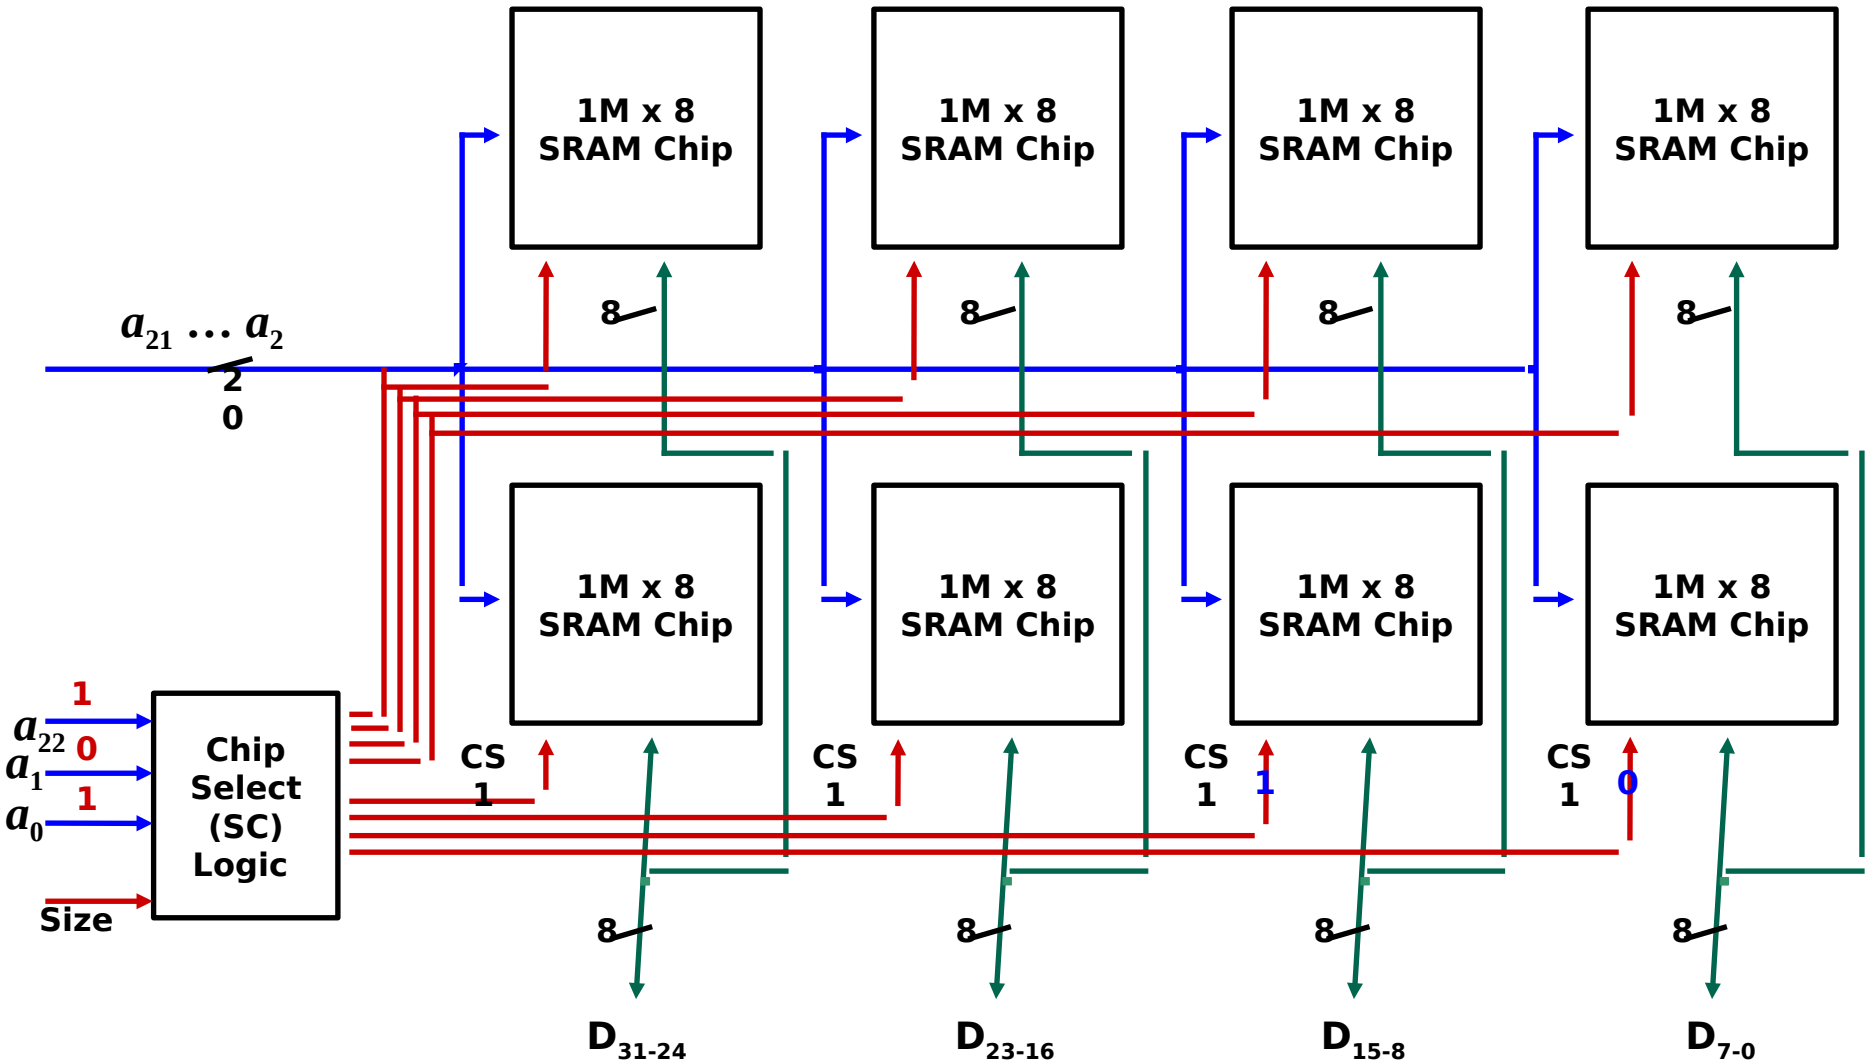


Illustration: 8MB Static RAM Module

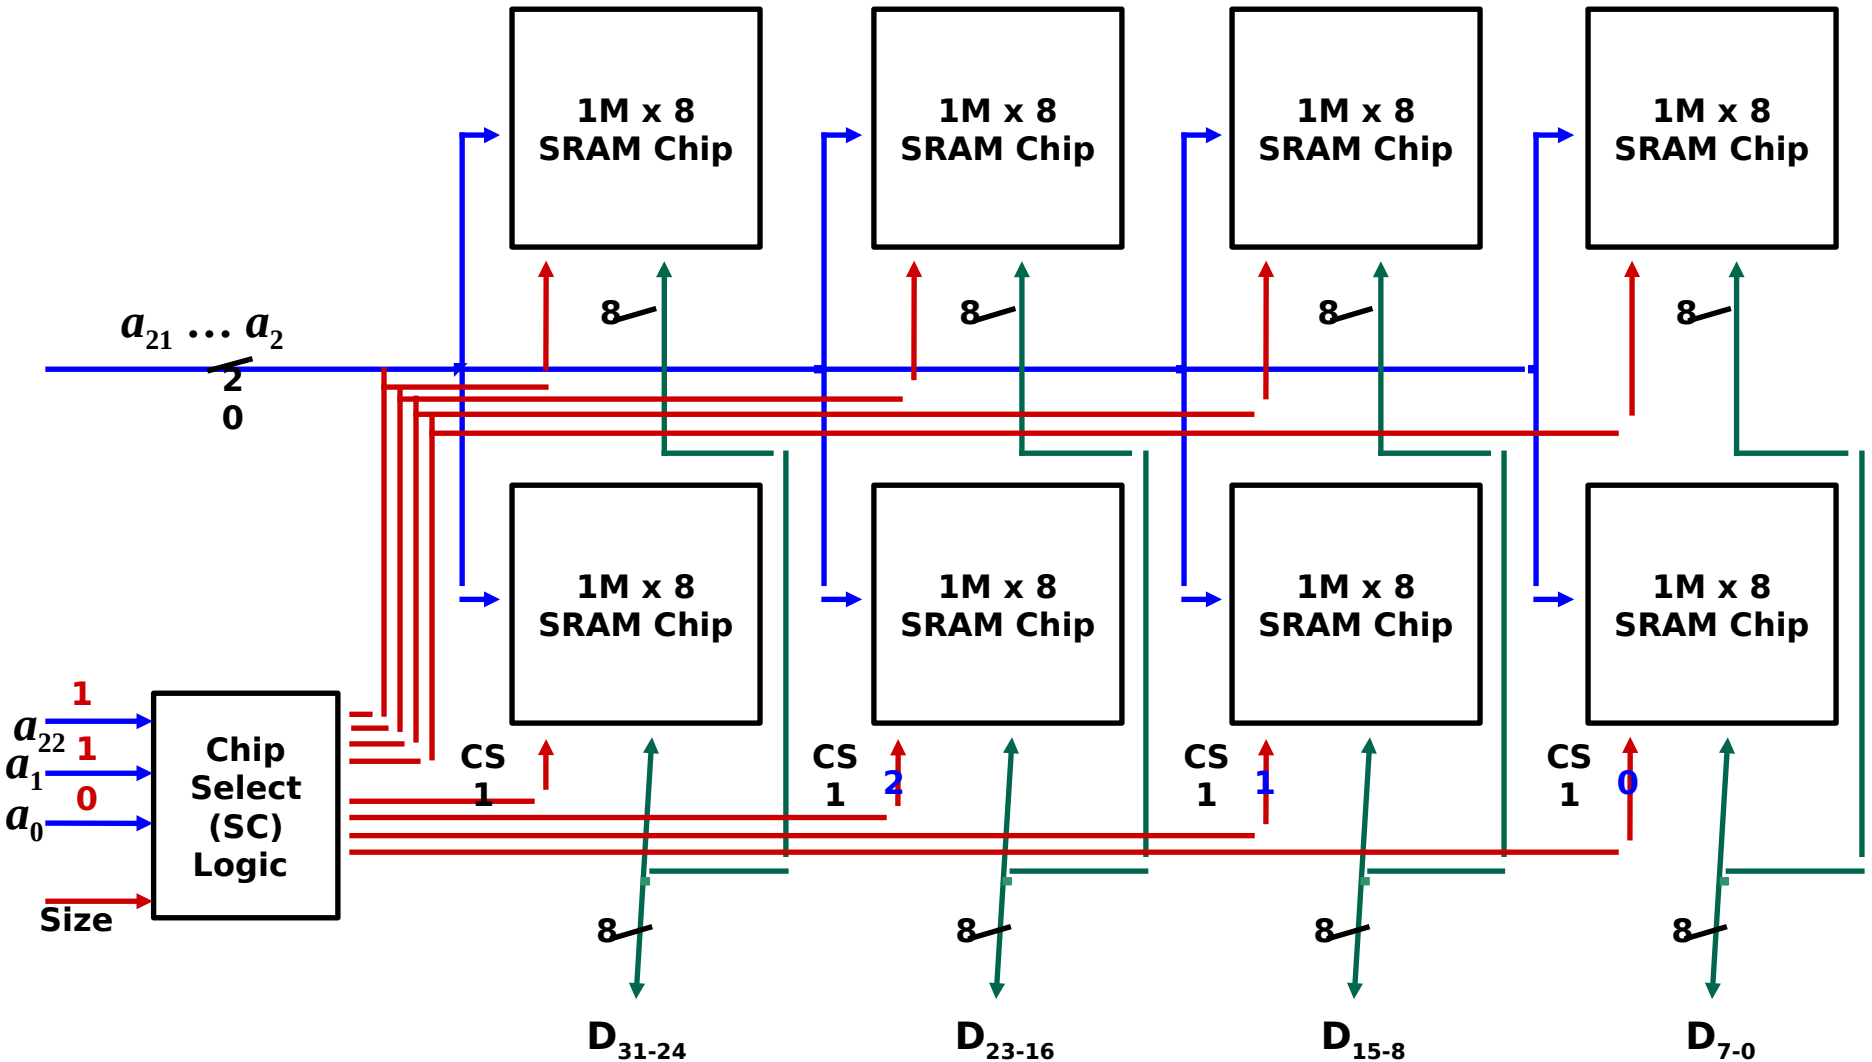
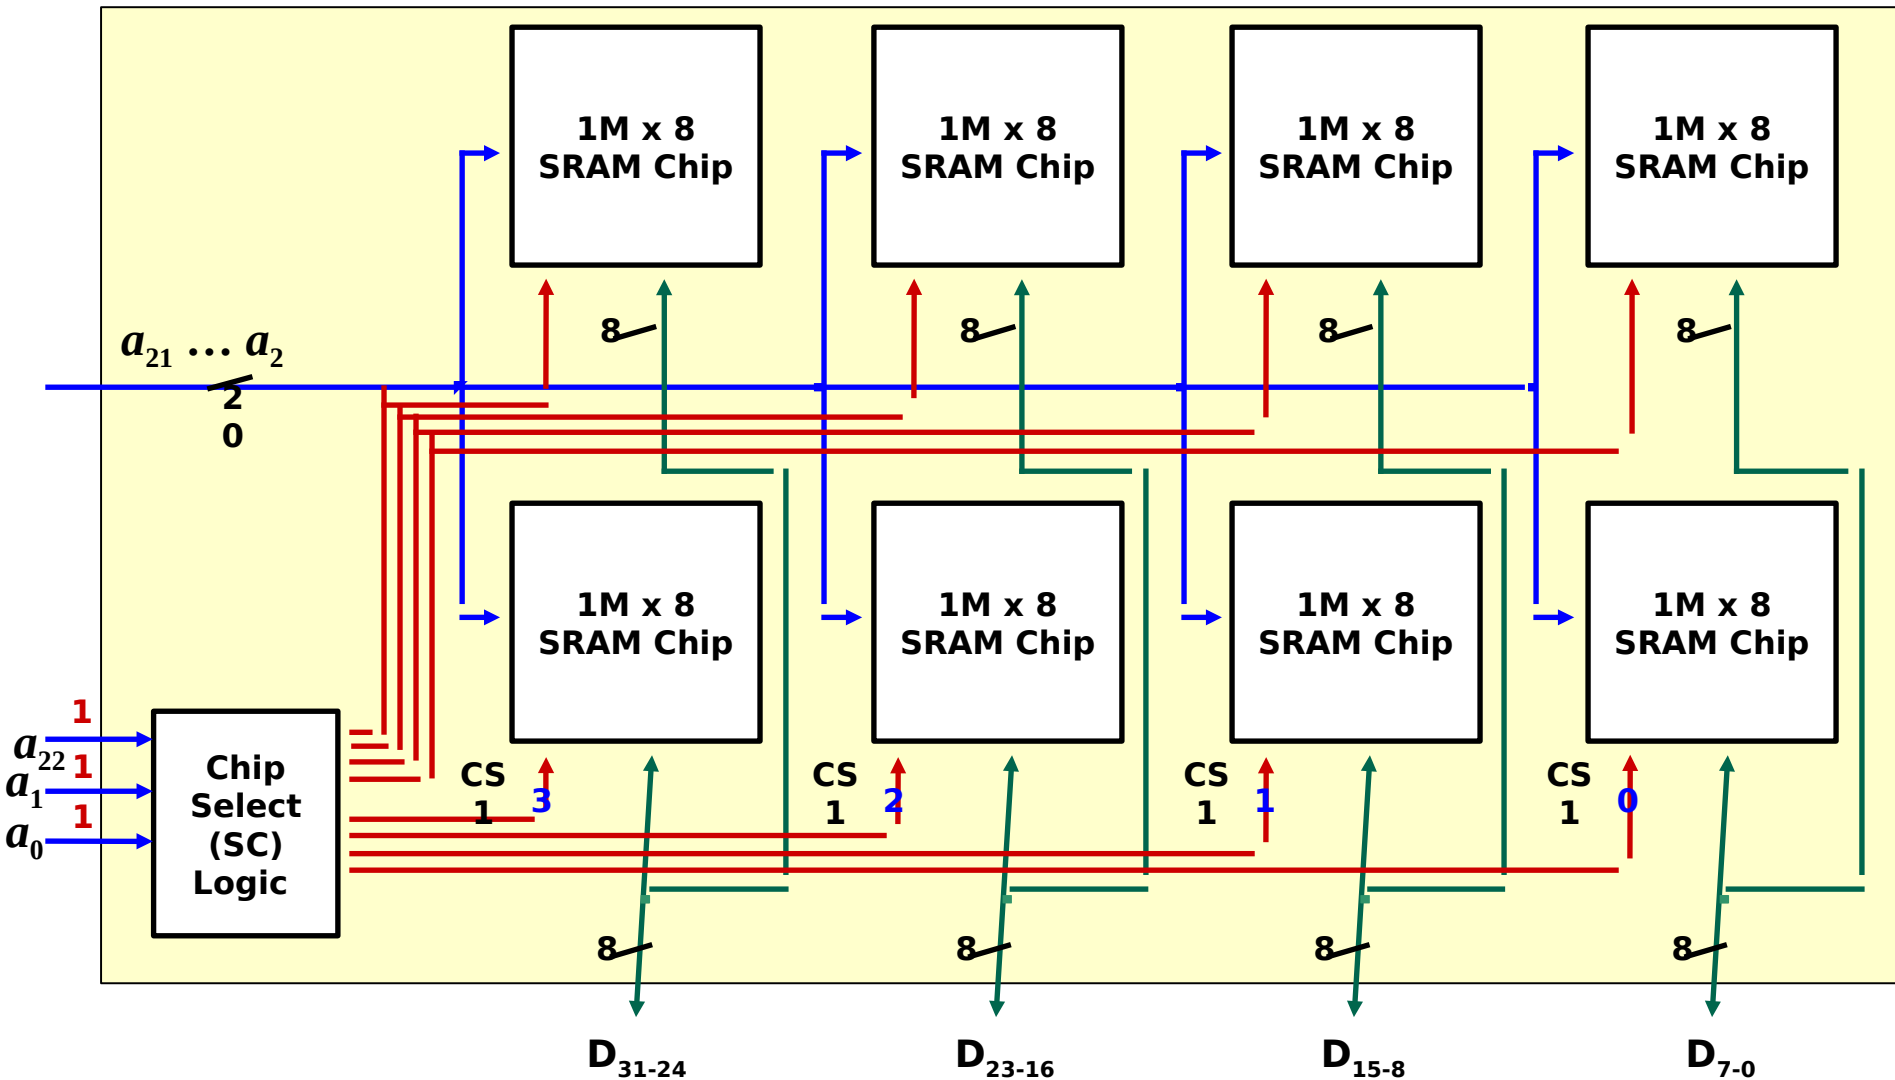


Illustration: 8MB Static RAM Module



Static RAM Module

- **4MB SRAM Module**

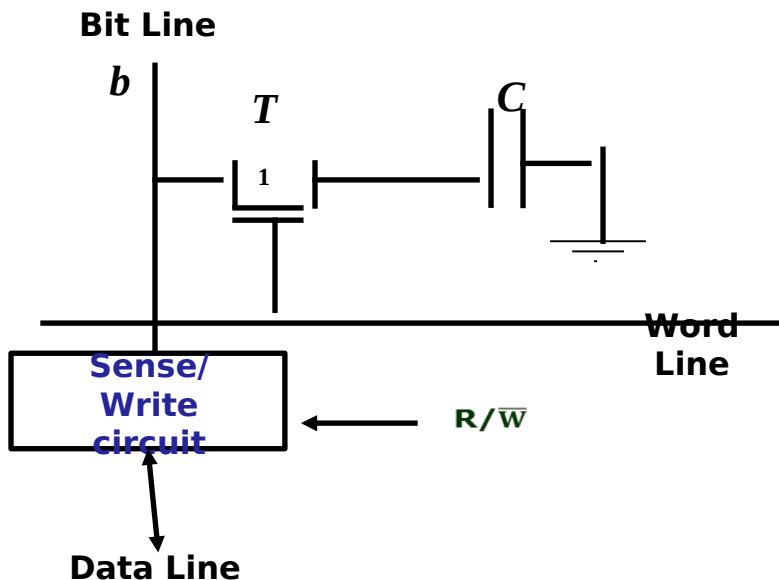


Static RAM Module

- N : Capacity of SRAM
- n bits: Width of SRAM
- M : Capacity of one SRAM chip
- m bits: Width of one SRAM chip
- Number of rows of memory chips: N/M
- Number of chips in a row: n/m
- According to the number of rows of memory chips and number of chips in a row, Chip Select (CS) logic is designed
- Higher order bits in address select a row of memory chips
- Lower order bits in address select a byte in a word
- Size line in CS logic indicates how many bytes in a word need to be selected

Dynamic RAM (DRAM) Cell

- Less expensive and simpler cell
- Information is stored in the form of a charge on a capacitor (C)
 - Charge in capacitor is stored only for short time
 - However, a cell is required to store information for a much longer time
- To retain information for longer time, content of capacitor must be periodically refreshed



- Low speed as refresh needed
- Only 1 transistor is used
- Used to build main memory

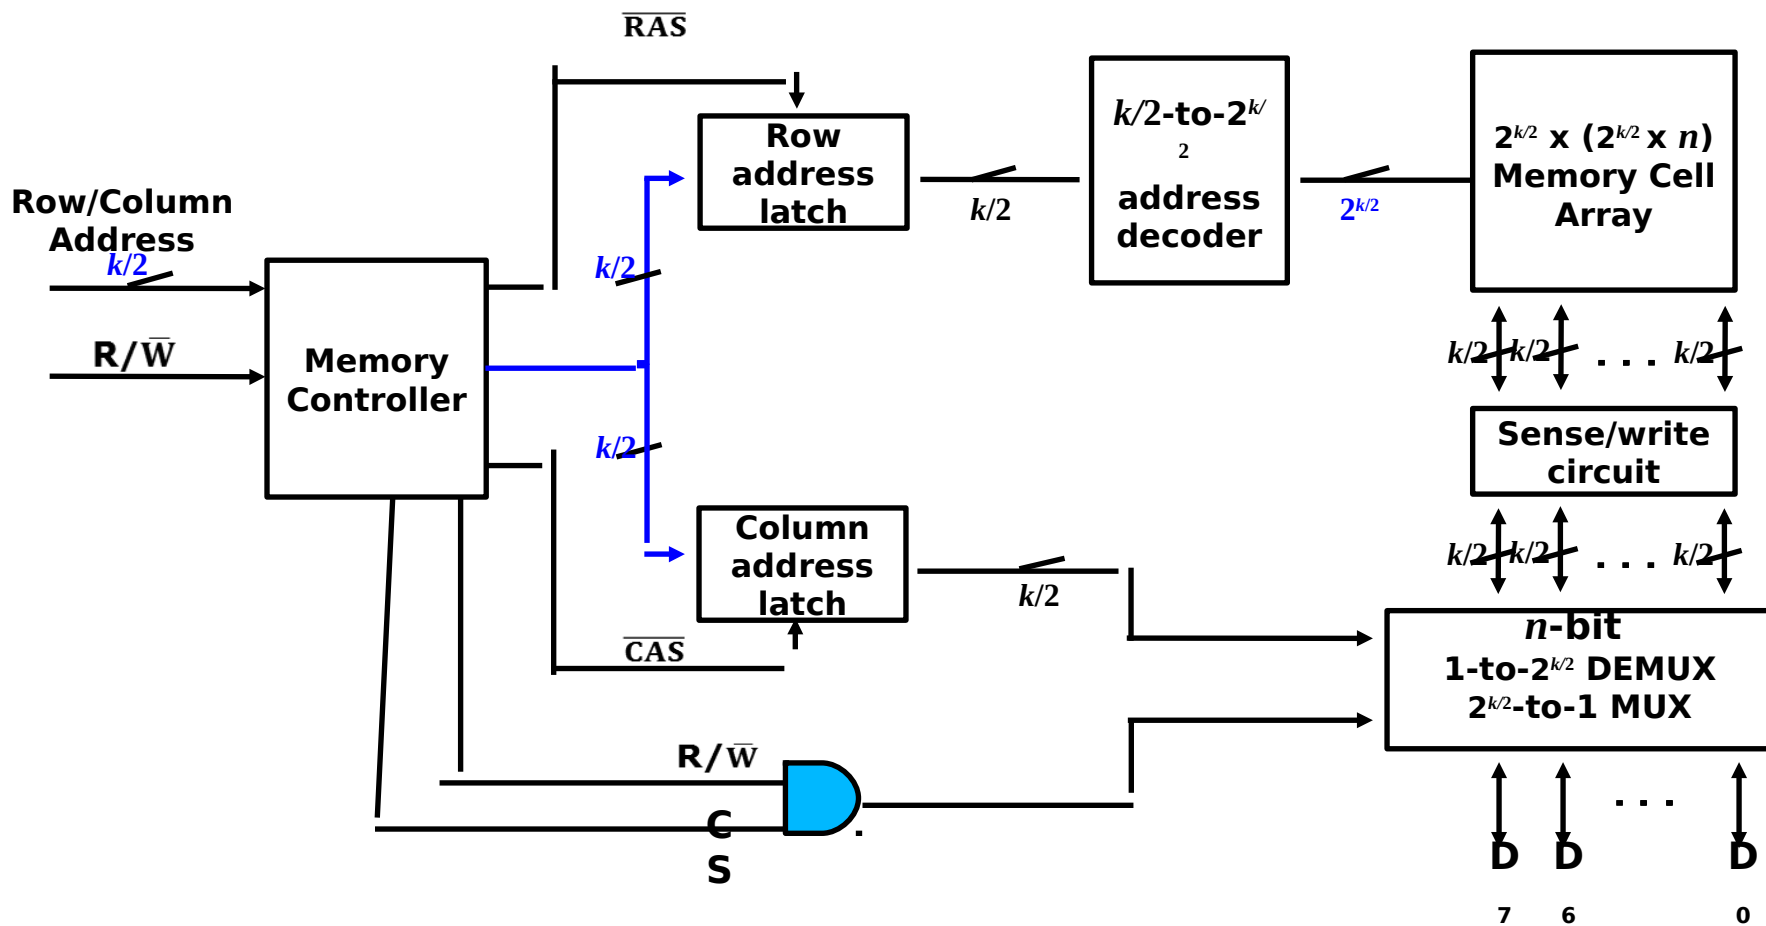
Dynamic RAM (DRAM) Chip Organization

- DRAM cells are also arranged in 2-dimensional array form
- Here also, row address lines are used to select row and column address lines are used to select column
- Two important factors influence the design of the DRAM chip are:
 - Number of input/output pins i.e. external pins
 - Need to refresh the cells
- Scheme for saving pins:
 - Row address and column address are transmitted over the same line one after the other
 - This is called time multiplexing
 - This is usually performed by a memory controller circuit
 - It generates the different control signals

Dynamic RAM (DRAM) Chip Organization

- Two additional control signals are needed to inform the chip when the row address and column address is valid on address line
 - Row address strobe ($\overline{\text{RAS}}$):
 - Inform the chip when the row address is valid on address lines
 - Column address strobe ($\overline{\text{CAS}}$):
 - Inform the chip when the row address is valid on address lines
 - Both $\overline{\text{RAS}}$ and $\overline{\text{CAS}}$ are usually active low

DRAM Chip Organization



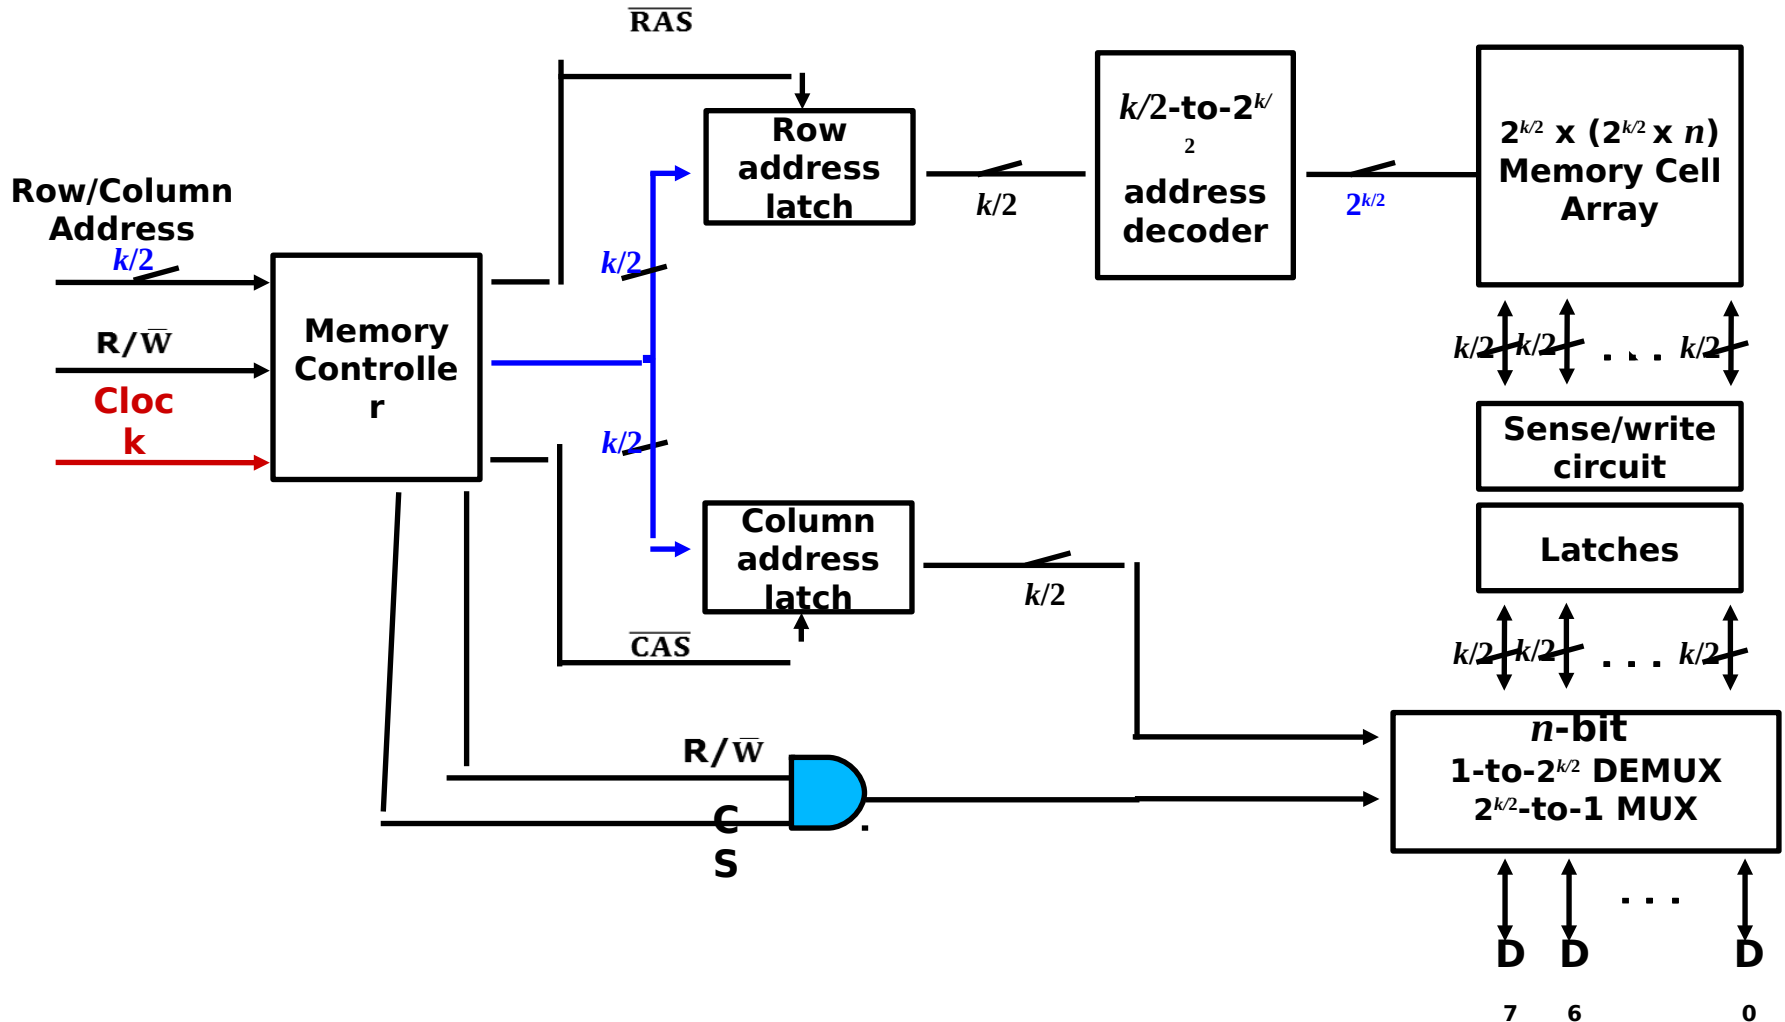
DRAM Chip Organization

- During read or write operation, the row address is applied first
 - It is loaded into row address latch in response to $\overline{\text{RAS}}$
 - Then read operation is initiated, in which all cells in a selected row are read and refreshed
- Shortly after row address is loaded, the column address is applied
 - It is loaded onto the column address latch under the control of $\overline{\text{CAS}}$
- The information in the column address latch is decoded and appropriate n sense/write circuits are selected
- To ensure that the contents of DRAM are maintained, each row of cells must be accessed and refreshed periodically

DRAM Chip Organization

- **Memory controller circuit** provide the necessary signals CAS and RAS that governs timing
- These operations are directly synchronised with clock signal
- Such a DRAM chip is called **Synchronous DRAM (SDRAM)**

SDRAM Chip Organization



Fast Page Mode Feature of SDRAM


- Transfer capability of SDRAM
- Contents of all $2^{k/2}$ cells in a selected row are sensed in each on the n cell arrays
- Only n bits (one from each of the cells arrays) are placed in the data lines, D_{7-0}
- To access other bytes in the same row, without having to reselect the row, a latch is used at the output of the sense/write circuits in each column
- The row address will load the latches corresponding to all bits in the selected row
- Then different column address are applied to place the different bytes on data lines
- Transfer bytes in sequential order
- This arrangement allows transferring a block of data at much faster rate

Refresh Overhead in SDRAM

- All dynamic memories need to be refreshed
- In SDRAM typical period of refreshing all rows is **64ms**
- Each row is refreshed at least in 64ms
- **Example:**
 - Suppose a SDRAM chip has 8K (8192) rows
 - Number of clock cycles to access each row: 4 clock cycles
 - Number of clock cycles to refresh all rows:
 $8192 \times 4 = 32768$ clock cycles
 - Suppose clock rate is 133 MHz
 - Times needed to refresh all rows:
 $32768 / 133 \times 10^6 = 246 \times 10^{-6} \text{ s} = 0.246 \text{ ms}$
 - Refreshing overhead is **0.246ms** out of **64ms**

Double-Data-Rate SDRAM (DDR SDRAM)

- Faster version of SDRAM
- The standard SDRAM performs all actions on the raising edge of the clock cycle
- **DDR SDRAM** access the cell array in the same way, but transfers the data on both edges of the clock
- Hence, their bandwidth is essentially doubled for long burst transfers
- **Bandwidth**: The number of bits/bytes that can be transferred in one second



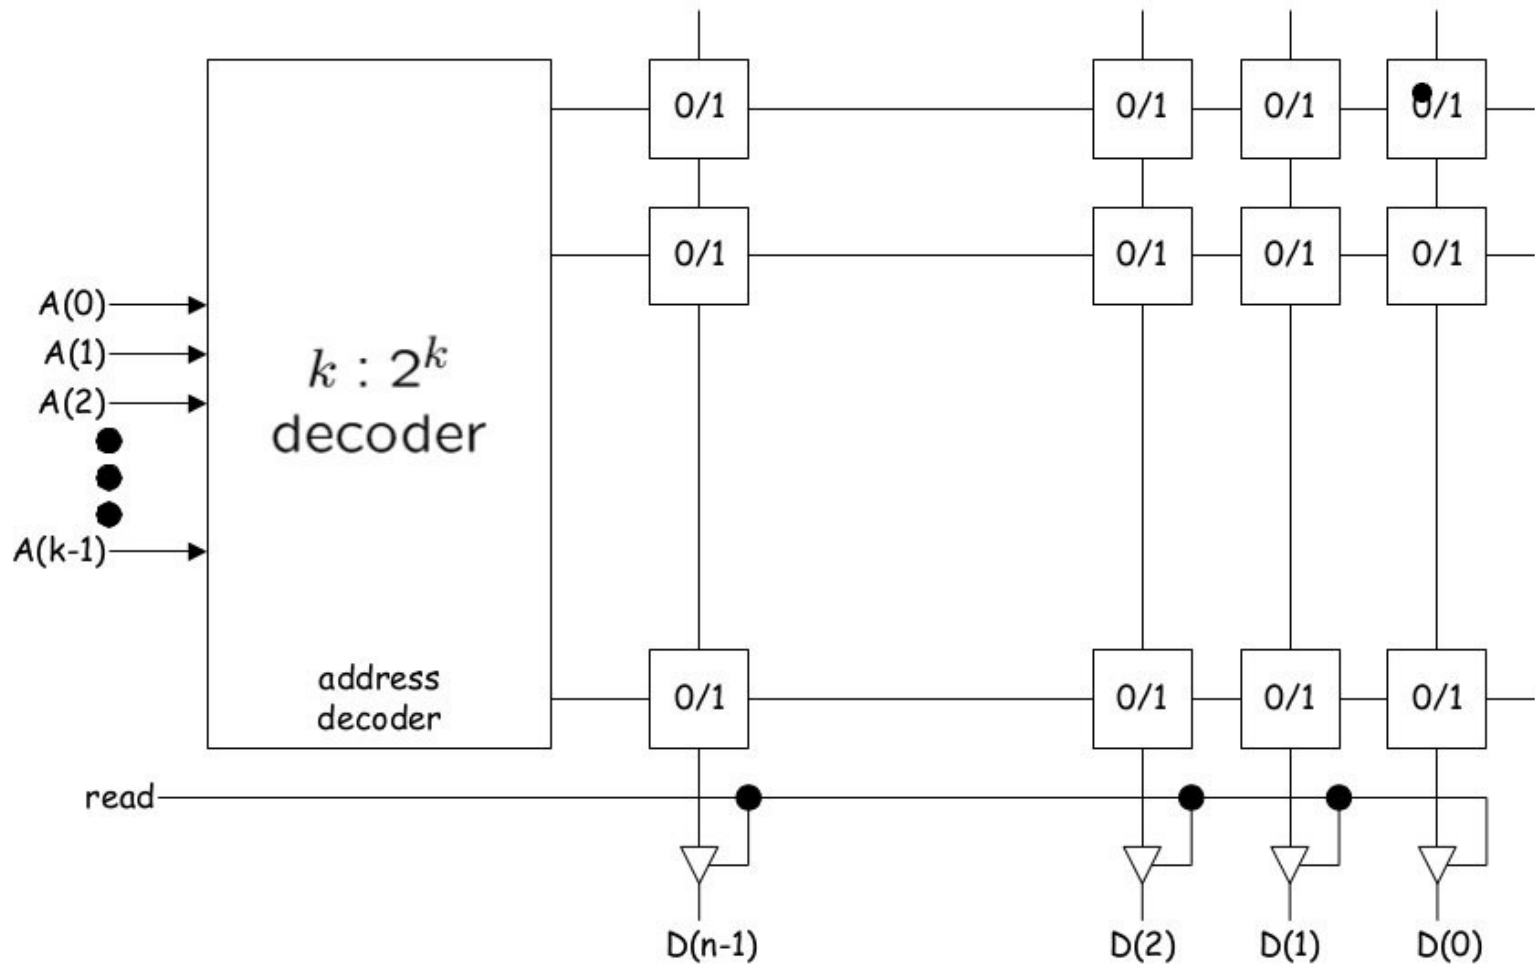
**BM 33L5039 RAM
Module**
**(1 GB, DDR2 RAM,
266 MHz, DIMM
184-pin)**



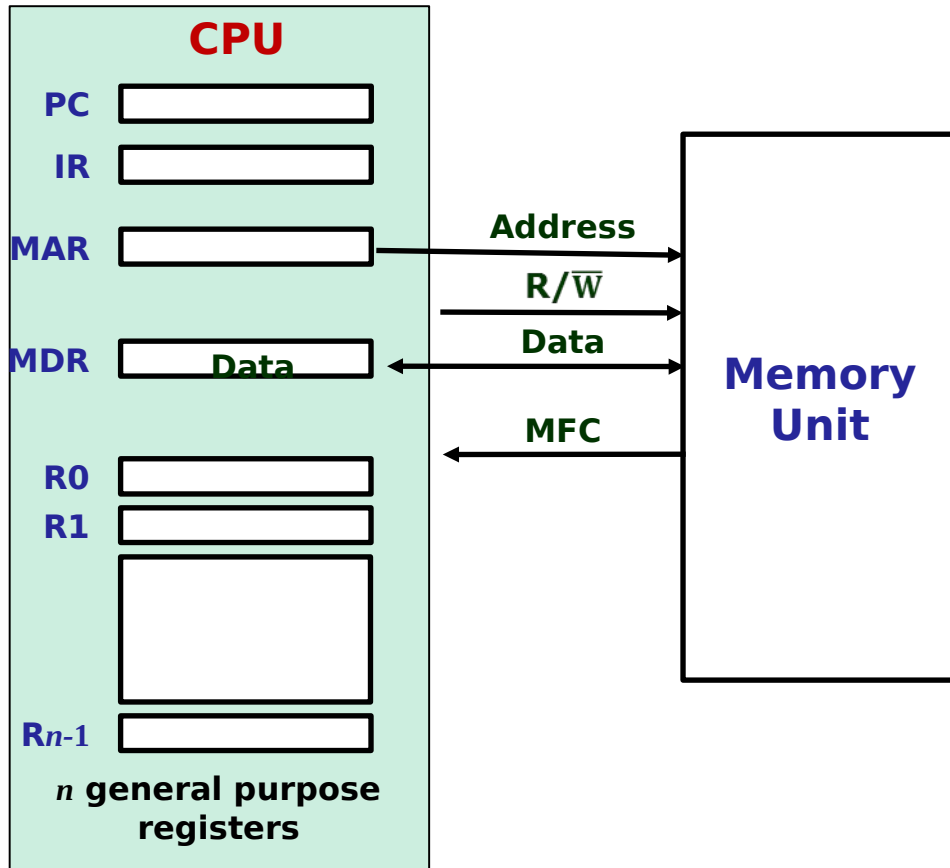
Read-Only Memories (ROMs)

- SRAM and DRAM are volatile i.e. they lose the stored information if power is turned off
- Read-only memories are semiconductor, non-volatile memories
- Their normal operation involves only reading the stored data
- They are extensively used in embedded systems
- Different types of ROMs
 - Read Only Memory (ROM)
 - Programmable ROM (PROM)
 - Erasable, reprogrammable ROM (EPROM)
 - Electrically erasable reprogrammable ROM (EEPROM)
 - Flash memory

Read-Only Memories (ROMs)



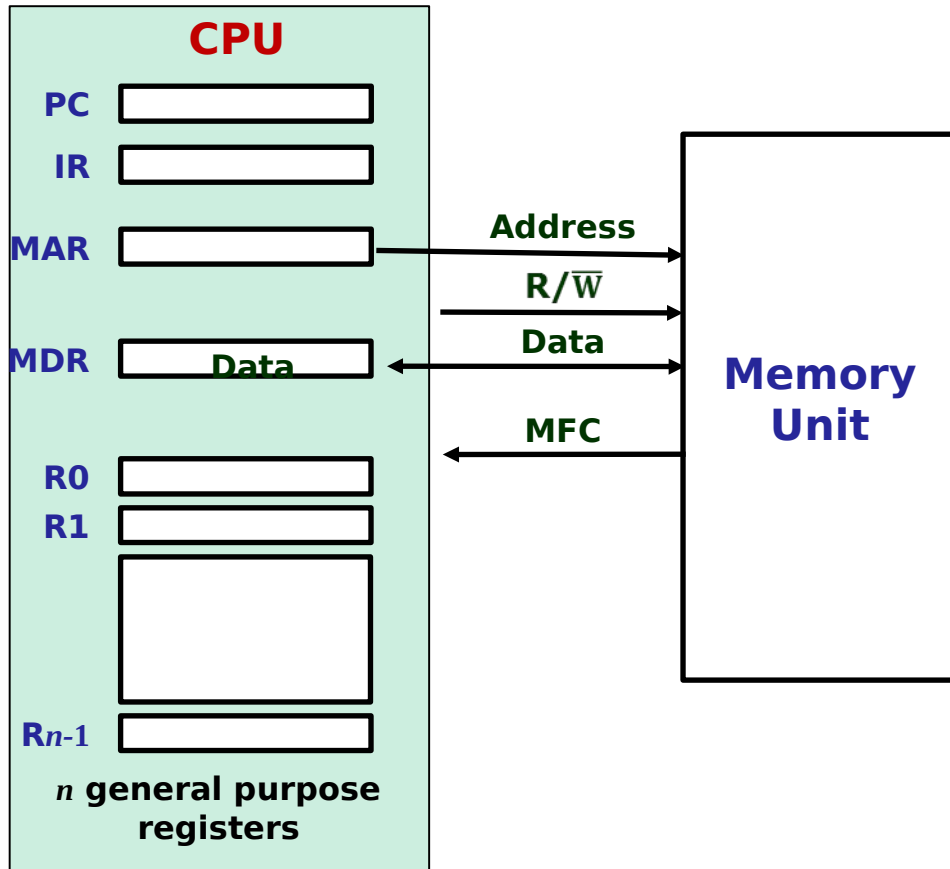
Memory Read and Write Operation



- **Read**

- Processor loads the address of memory location into MAR
- Set the R/\bar{W} line to 1
- Memory responds by placing the data from address location onto data line
- Confirm the action by asserting MFC (memory function complete) signal
- Upon receiving MFC signal, processor loads the data on data line into MDR

Memory Read and Write Operation




- **Read**

- Processor loads the address of memory location into MAR
- Set the R/\bar{W} line to 1
- Memory responds by placing the data from address location onto data line
- Confirm the action by asserting MFC (memory function complete) signal
- Upon receiving MFC signal, processor loads the data on data line into MDR

Cache Memory

- The cache memories are designed to exploit the **locality of reference** in the program
- **Locality of reference:**
 - Many instructions in **localized areas of the program** are **executed repeatedly** during some time period, and the remainder of the program is accessed relatively infrequently
- Different ways of locality of reference
 1. **Temporal locality:**
 - Recently executed instruction/data is likely to be used again very soon
 2. **Spatial locality:**
 - Instructions in close proximity to a recently executed instruction/data (with respect to instruction address) are likely to be executed very soon

Use of a Cache Memory



- Unit of transfer between main memory and cache is **block**
 - A block is a set of fixed number of words in contiguous address locations
 - Cache block is also called as **Cache line**
- **Mapping function**: Correspondence between main memory blocks and those in the cache

Replacement Algorithm

- Read hit/Write hit [Cache hit]:
 - When processor issues read or write request, **cache control circuit** determines whether the requested word exists in cache
 - If it exists in cache, the read or write operation is performed on the appropriate cache location
 - This means, **read hit or write hit is said to have occurred**
- Cache miss:
 - If the desired word is not there in cache during read/write operation, then **cache miss is said to have occurred**
 - During that time new block need to be brought into cache
- **Cache control hardware** decide which block to removed to create space for new block that contain referenced word when cache is full
- The collection of rules for making this decision constitutes the **replacement algorithms**

Read and Write Operations on Cache


- Read operation:
 - Handling read miss:
 - Approach 1:
 - The block of words that contains the requested word is copied into the cache
 - After entire block is loaded into cache, particular requested word is forwarded to processor
 - Approach 2: Load through or early restart
 - A word is sent to processor as soon as it is read from the main memory
 - At the same time, the block of words that contains the requested word is also copied into the cache
 - This approach reduces the processor waiting period, but with the expense of complex circuitry

Read and Write Operations on Cache

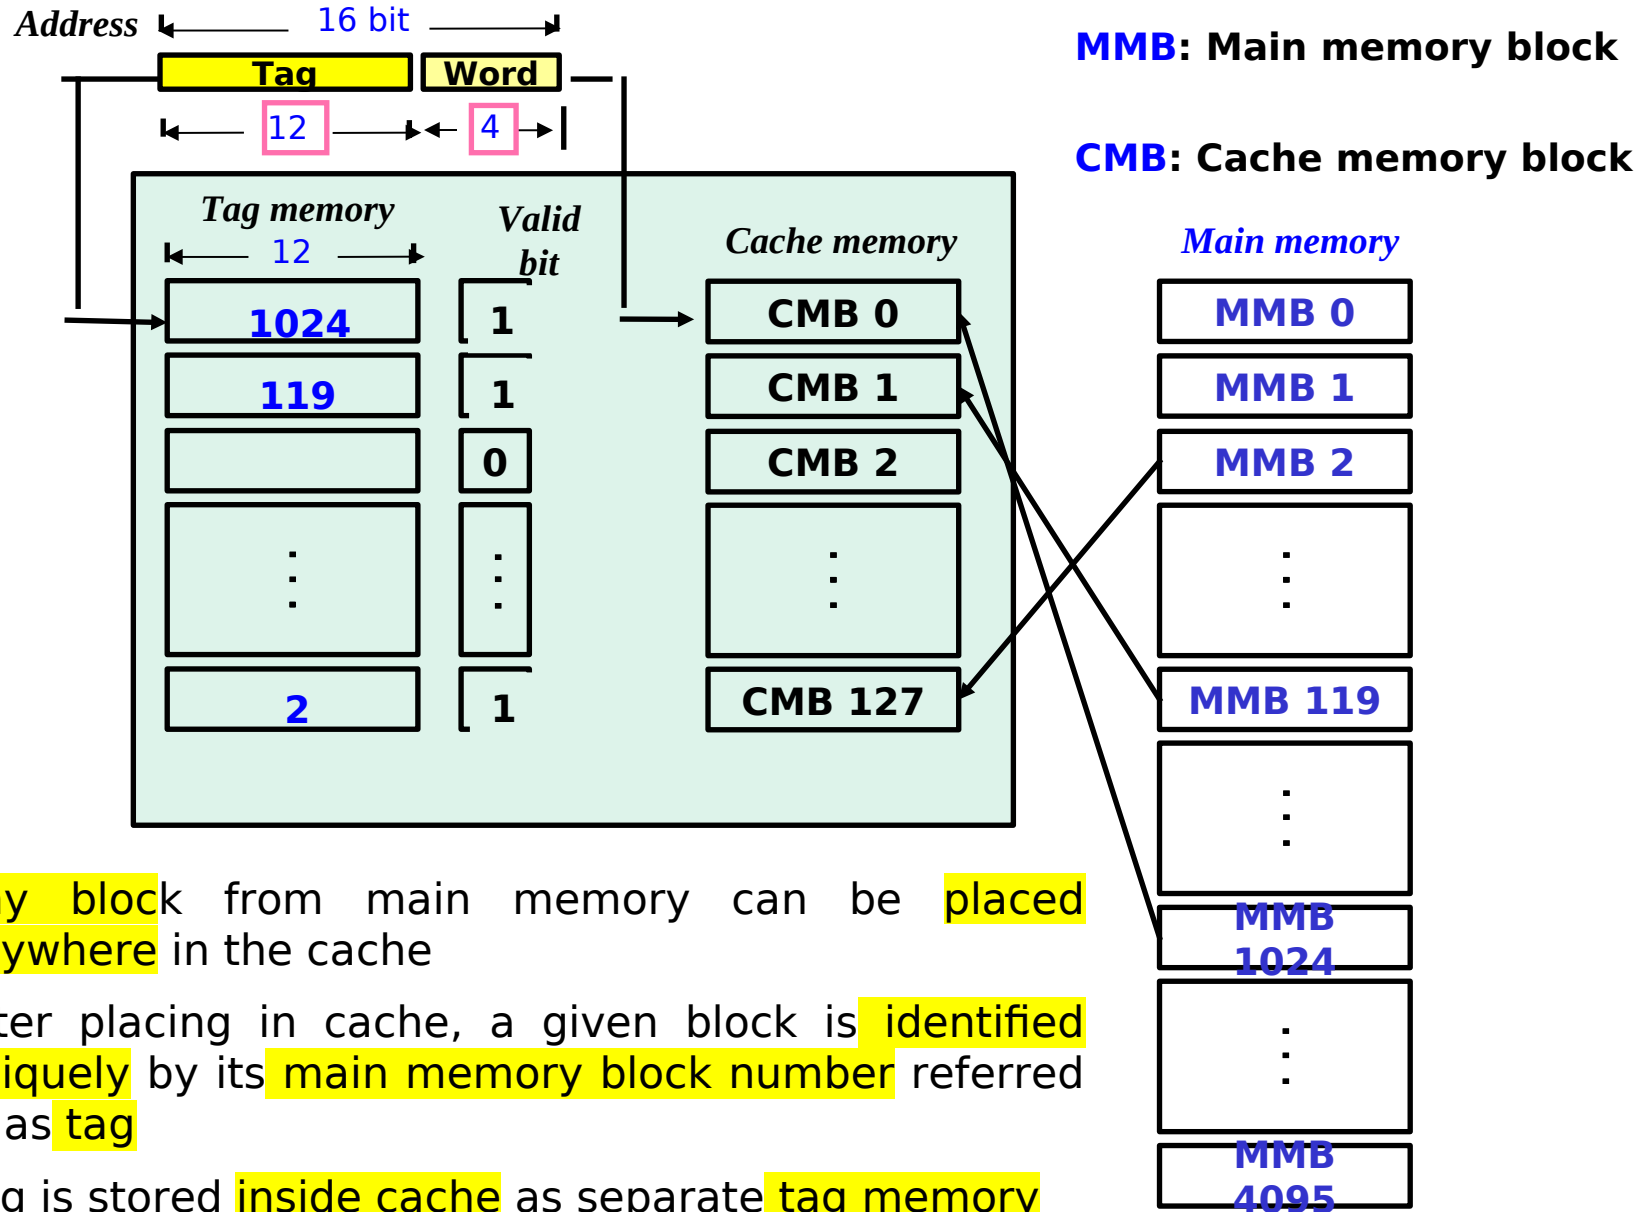
- Write operation:
 - Two techniques for write operation:
 - Write-through protocol
 - The cache location and main memory location are updated simultaneously
 - Write-back (copy-back) protocol
 - Update only the cache location and mark it as updated in a flag bit called dirty bit or modified bit
 - The main memory location is updated later when the block containing that modified word need to be removed from the cache
 - Handling write miss:
 - Write-through protocol:
 - The information is written directly into the main memory
 - Write-back (copy-back) protocol:
 - The block containing the addressed word is first brought into the cache
 - Then the desired word in the cache is overwritten with new information

Mapping Function

- Specifies where memory blocks are placed in the cache
- Cache and main memory are viewed as collection of fixed number of blocks
- **Example:** Consider a cache and main memory with 2K words and 64K words respectively and each blocks are of size 16 words
 - Block size: 16 words
 - Number of blocks in a cache, $N = 2K/16 = 128$
 - Number of blocks in main memory, $M = 64K/16 = 4K = 4096$
 - Addressable location in main memory: 2^{16}



Mapping Function: Associative Mapping (Associative Mapped Cache)



Mapping Function:

Associative Mapping (Associative Mapped Cache)

- Any block from main memory can be placed anywhere in the cache
- After placing in cache, a given block is identified uniquely by its main memory block number referred to as **tag**
- Tag is stored inside cache as separate **tag memory**
- Cache maintains a **control bit called valid bit** for each block **to indicate whether the block contains valid data**
- Main memory address reference is partitioned into two parts: **Tag and word**
- **Advantage:**
 - **Simple and most flexible**
 - **Complete use of its capacity**
- **Disadvantage:**
 - Tag memory must be **searched entirely** for each memory reference: **Associative search**
 - **Expensive**

Mapping Function: Direct Mapping (Direct Mapped Cache)

- Main memory block is placed in one and only one place in the cache
- Simplest way to determine the cache location
- Let N be the number of blocks in cache
- The j th main memory block is placed (mapped) onto $(j \bmod N)$ th block of the cache

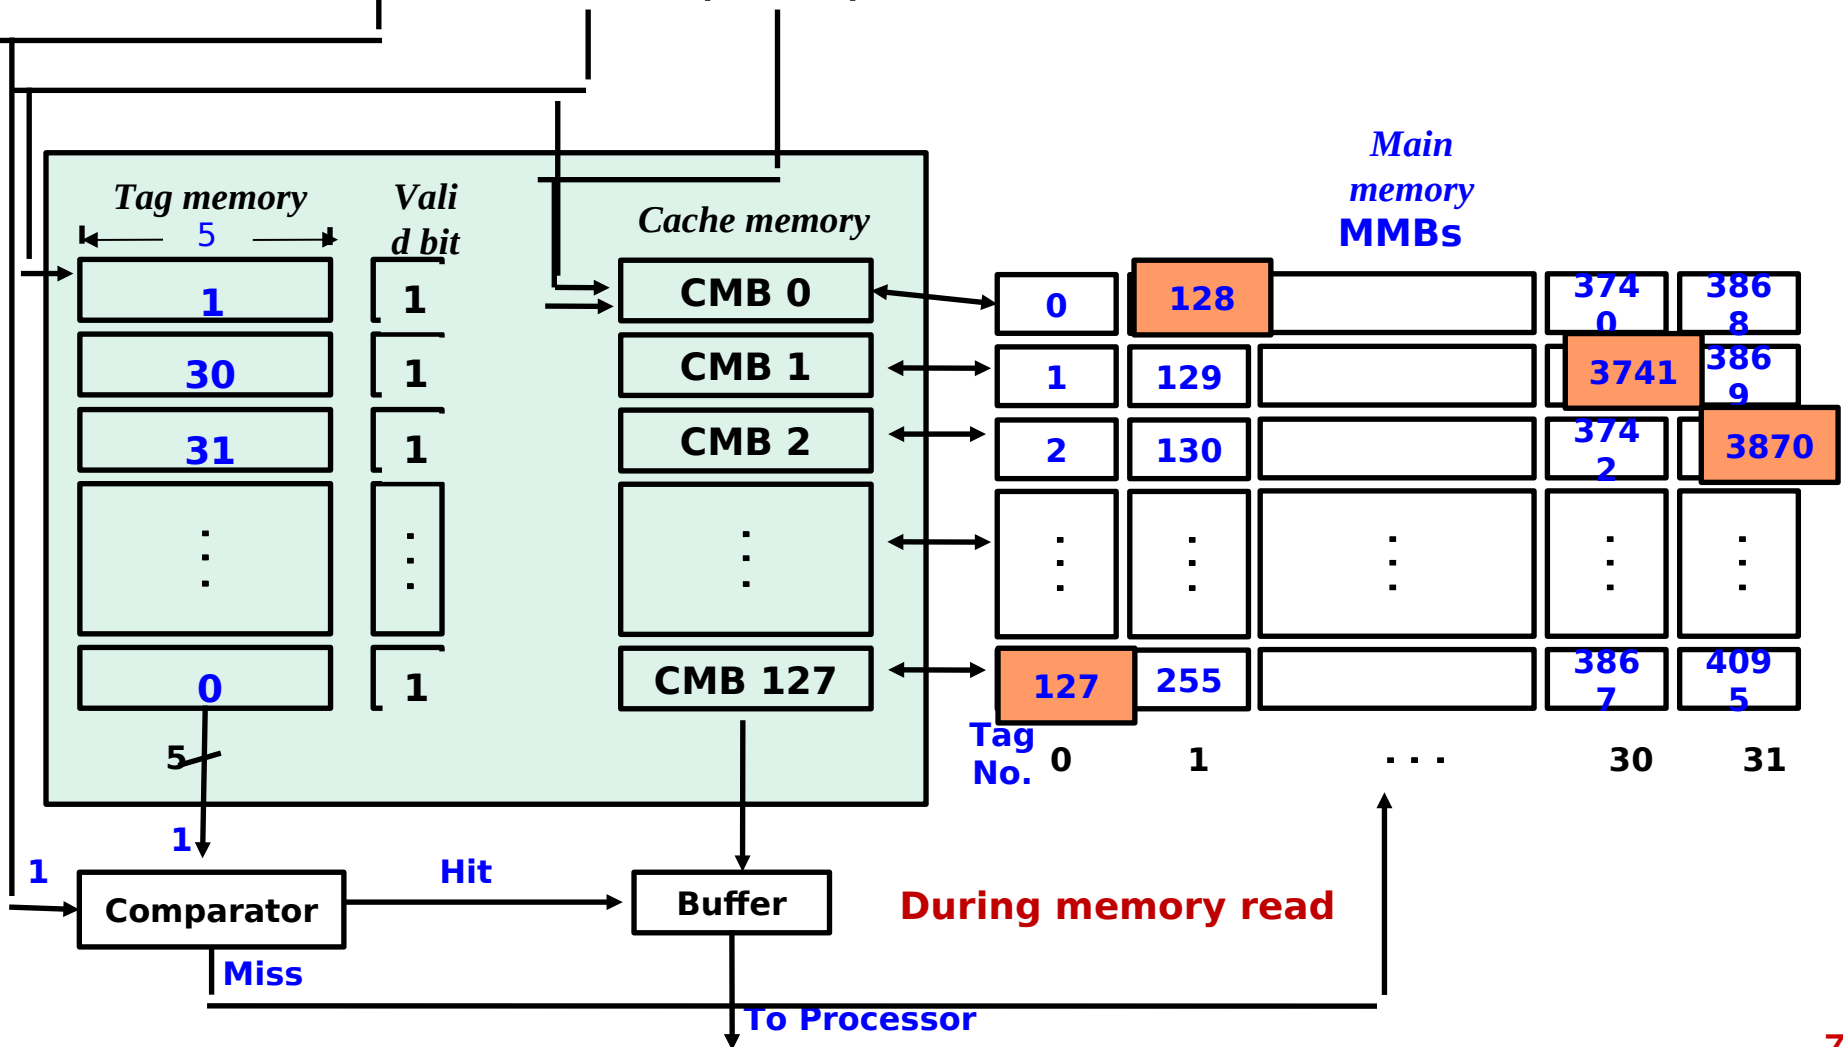

$$\text{CMB}_i = \text{MMB}_j \bmod N$$

- **Example:** For number of blocks in a cache, $N = 128$ and the number of blocks in main memory, $M = 4096$
 - The main memory blocks 0 or 128 or 256 etc. are mapped onto cache block 0
 - The main memory blocks 1 or 129 or 257 etc. are mapped onto cache block 1
- For the understanding sake, let's consider the main memory as rectangular array of blocks

Mapping Function: Direct Mapping (Direct Mapped Cache)

$$\text{CMB}_i = \text{MMB}_j \bmod 128$$

$$3828 \bmod 128 = 127$$




Mapping Function: Direct Mapping (Direct Mapped Cache)

- Advantage:
 - Simplicity
 - Cost is less compared to fully associative cache
- Disadvantage:
 - Only a single block from a given group can be present in the cache
 - Not flexible
 - Imposes considerable amount of rigidity on the cache organization
- It relies on the principle of locality of reference for its success
- If two blocks of same group are frequently referenced, it leads to trashing

Mapping Function: Set-Associative Mapping

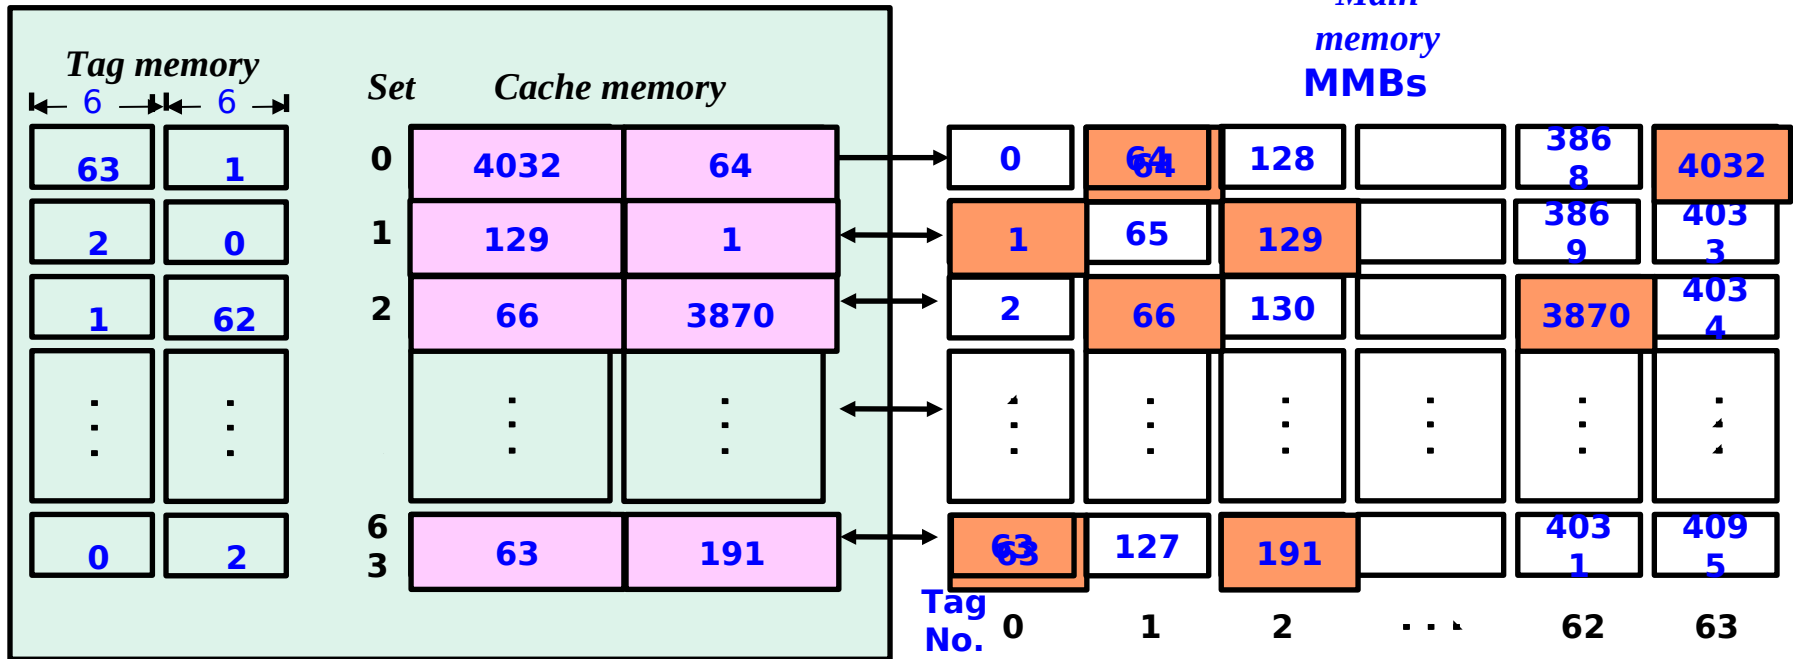
- Combination of the direct and fully associative mapping
- Blocks of the cache are grouped into sets
- Mapping allows block of the main memory to reside in any block of a specific set
- Within a set, it is fully associative
- **K-way set associative:**
 - A set may hold K number of blocks
 - Number of sets = N/K
 - $CMS_i = MMB_j \bmod (N/K)$ where CMS = Cache memory set
- Intel P-III and P-IV processors contain 8-way set associative cache
- **Example:** Let number of blocks in a cache, $N = 128$ and the number of blocks in main memory, $M = 4096$

Mapping Function: 2-way Set-Associative Mapping



$$CMS_i = MMB_j \bmod 128/2$$

$$3870 \bmod 64 = 62$$



Replacement Algorithm

- Direct mapped cache:
 - The position of each block is predetermined
 - No replacement strategy exists
- In associative and set-associative cache, there exists flexibility
 - Replacement strategy needed
- Replacement algorithms helps in deciding which of the old blocks in the cache to be replaced when the cache is full and a new block is brought into the cache
- Objective: To keep the blocks in cache that are likely to be referenced
- Locality of reference in program gives a clue to a reasonable strategy
- It is sensible to replace the block that has gone longest time without being referenced

Least Recently Used (LRU) Algorithm

- The cache controller need to track references to all blocks
- LRU algorithm is implemented by using counters for each block
- LRU algorithm:
 - **During cache hit:**
 - Counter of the block referenced is set to 0
 - Counter of the empty block locations are unchanged
 - Counter of other occupied block locations are incremented by 1
 - **During cache miss and cache is not full:**
 - Load the main memory block to empty space and set the counter of that block to 0
 - Increment the counter of other block location by 1
 - **During cache miss and cache is full:**
 - Block with largest counter value (LRU) is removed
 - New block is loaded in that emptied location
 - Counter of that location is set to 0
 - Increment the counter of other locations by 1

Illustration: LRU Algorithm

- Consider **fully associative cache**
 - Let number of blocks in a cache, $N = 4$ and the number of blocks in main memory, $M = 8$
 - Let each **block is 4 words**

Read LOC#3

Read LOC#7

Read LOC#11

Read LOC#2

Read LOC#8

Read LOC#14

Read LOC#16

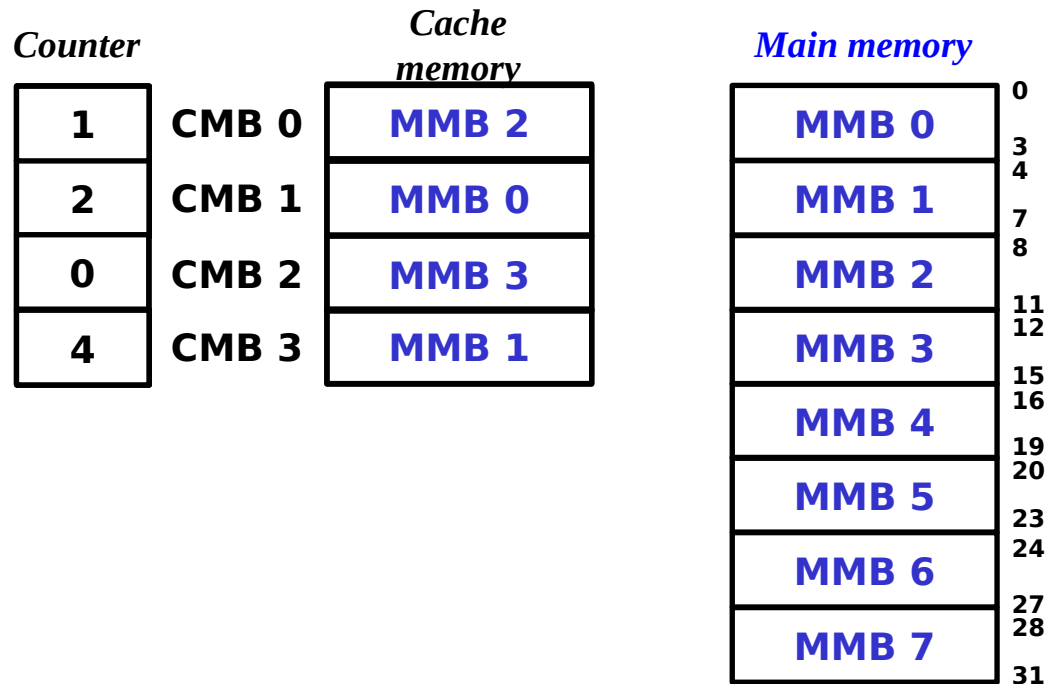


Illustration: LRU Algorithm

- Consider **fully associative cache**
 - Let number of blocks in a cache, $N = 4$ and the number of blocks in main memory, $M = 8$
 - Let each block is 4 words

Read LOC#3
Read LOC#7
Read LOC#11
Read LOC#2
Read LOC#8
Read LOC#14
Read LOC#16
Read LOC#23

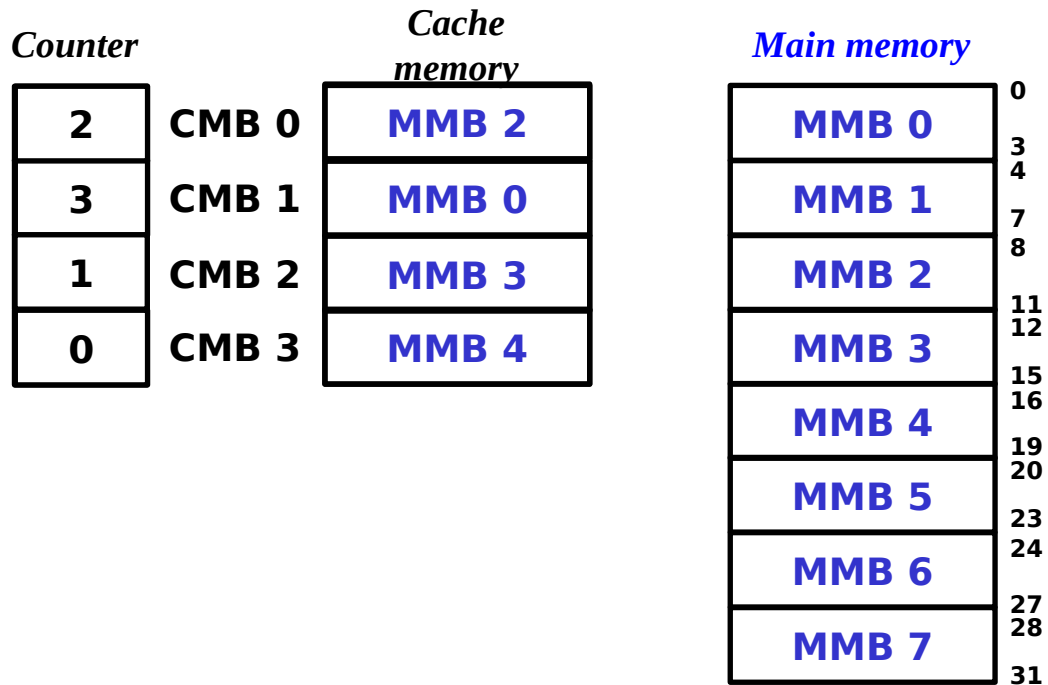


Illustration: LRU Algorithm

- Consider **fully associative cache**
 - Let number of blocks in a cache, $N = 4$ and the number of blocks in main memory, $M = 8$
 - Let each block is 4 words

Read LOC#3
Read LOC#7
Read LOC#11
Read LOC#2
Read LOC#8
Read LOC#14
Read LOC#16
Read LOC#23
Read LOC#25

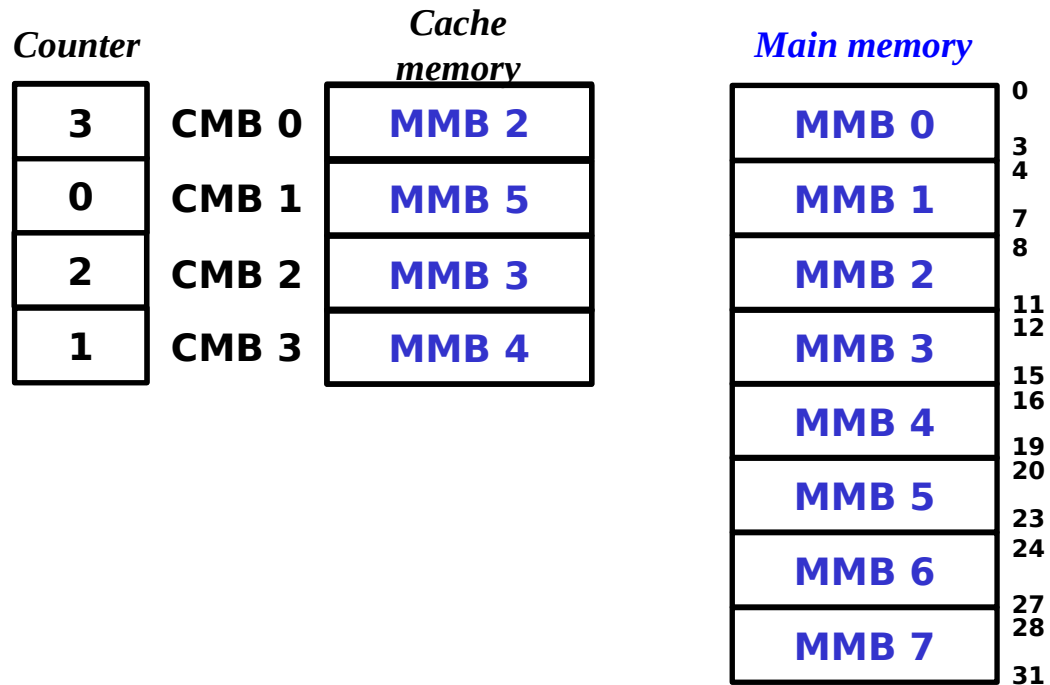


Illustration: LRU Algorithm

- Consider **fully associative cache**
 - Let number of blocks in a cache, $N = 4$ and the number of blocks in main memory, $M = 8$
 - Let each block is 4 words

Read LOC#3
Read LOC#7
Read LOC#11
Read LOC#2
Read LOC#8
Read LOC#14
Read LOC#16
Read LOC#23
Read LOC#25
Read LOC#27

Counter		Cache memory	Main memory	
0	CMB 0	MMB 6	MMB 0	0
2	CMB 1	MMB 5	MMB 1	3
4	CMB 2	MMB 3	MMB 2	4
3	CMB 3	MMB 4	MMB 3	7
			MMB 4	8
			MMB 5	11
			MMB 6	12
			MMB 7	15
				16
				19
				20
				23
				24
				27
				28
				31

Cache Memory Architectures

- Cache memory hierarchy: L1 cache, L2 cache and L3 cache
- Two architectures of cache memories:
 - Harvard architecture cache
 - Separate data and instruction cache
 - Advantages:
 - Prevents conflicts between blocks of instruction and data that might map onto same location
 - Program generally do not modify instructions
 - Instructions take less memory than program data
 - Generally used in L1 cache
 - Unified cache (Princeton architecture cache)
 - Cache contain both instruction and data

Categories of Cache Misses

- **Goal:** To reduce number of cache misses
- This requires to know the reasons for cache misses

1. Compulsory miss:

- Occurs when the **cache is first referenced**
- It causes the block to be brought into the cache

2. Capacity miss:

- Occurs when the amount of **data referenced** by the program **exceeds the capacity** of the cache
- It causes some blocks to be **evicted to make room** to new data
- If the **evicted data** is **referenced again** by the program, **capacity miss** occurs

3. Conflict miss:

- Occurs in associative/set-associative mapping
- Occurs when program references **more blocks of data** mapped on to the **same set**
- It causes **one of the block to be removed**
- If the removed block is **referenced again**, **conflict miss** occurs

Cache Performance Considerations

- Ideally, **entire memory** hierarchy would appear to the processor as a **single memory unit**, that has the **access time of a cache** on processor (L1) and the **size of memory disk**
- **Performance is adversely affected** by the actions that must be taken after a miss
- **Hit**: Successful access to data in a cache
- **Hit rate (h)**: **Ratio of number of hits** over **all attempted access**
 - **Example**: $h=0.9$ means 90% of the time required block is in cache
- **Miss rate**: Ratio of number of misses over all attempted access

Cache Performance Considerations

- Miss penalty:
 - The extra time needed to bring the desired information into the cache
 - It is the time that the processor is stalled during waiting
 - It is also the time needed to bring a block of data from a slower unit in the hierarchy to a faster unit
- Suppose we have a cache and a main memory in the hierarchy
 - Let h be the hit rate
 - t_M be the miss penalty i.e. time to access information in the main memory
 - t_C be the cache hit latency i.e. time to access information in cache
 - The average access time experienced by the processor:
$$t_{\text{avg}} = h t_C + (1-h) t_M \quad \text{i.e.} \quad t_{\text{avg}} = \text{hit rate} * t_C + \text{miss rate} * t_M$$

Cache Performance Considerations

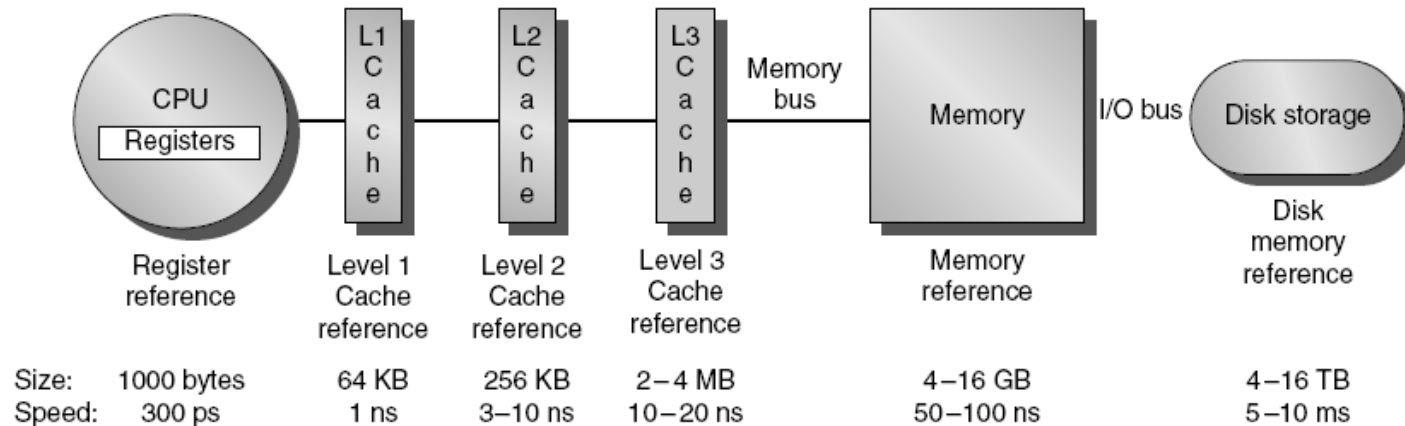
- Suppose we have L1 & L2 caches and a main memory in the hierarchy
 - h_1 : hit rate of L1 cache
 - h_2 : hit rate of L2 cache
 - t_M be the miss penalty i.e time to access information in the main memory
 - t_{C1} : Time to access L1 cache (L1 cache latency)
 - t_{C2} : Time to access L2 cache (L2 cache latency)
 - The **average access time** experienced by the processor in the two level cache:

$$t_{\text{avg}} = h_1 t_{C1} + (1-h_1) h_2 t_{C2} + (1-h_1) (1-h_2) t_M$$


Cache Performance Considerations

- A cache has a hit rate of 95%, and a block capacity of 128-byte and a cache hit latency of 5 ns. Each word in a block is 32 bits. The main memory takes 100 ns to return a block.
 - What is the cache block size?
 - What is the average memory access time?
- Ans: 1. Cache block size: 32 words
2. Average memory access time: 9.75 ns

Memory Hierarchy




(a) Memory hierarchy for server



(b) Memory hierarchy for a personal mobile device

Memory Performance Gap



Memory Hierarchy Design


- Memory hierarchy design becomes more crucial with recent multi-core processors:
 - Aggregate peak bandwidth grows with # cores:
 - Intel Core i7 can generate two references per core per clock
 - Four cores and 3.2 GHz clock
 - 25.6 billion 64-bit data references/second +
 - 12.8 billion 128-bit instruction references
 - = 409.6 GB/s!
 - DRAM bandwidth is only 6% of this (25 GB/s)
 - Requires:
 - Multi-port, pipelined caches
 - Two levels of cache per core
 - Shared third-level cache on chip

Memory Hierarchy

- Average memory access time= Hit time+ Miss rate x miss penalty
- basic cache optimizations:
 - Reducing miss rate
 - Larger block size
 - Larger cache
 - Higher associativity
 - Reducing miss penalty
 - Multi level cache
 - Reducing the time to hit in the cache
 - Avoid address translation when indexing the cache

Memory Hierarchy

- basic cache optimizations:
 - Larger block size
 - Reduces compulsory misses
 - Increases capacity and conflict misses, increases miss penalty
 - Average access time is used to decide the size of cache and block



Memory Hierarchy

- Basic cache optimizations:
 - Higher associativity
 - Reduces conflict misses
 - Eight-way set associativity is as effective in reducing miss rate for these sized caches as fully associativity
 - 2:1 cache rule of thumb: Direct map of cache of size N has about the same miss rate as two-way set associativity of size $N/2$
 - Increases hit time, increases power consumption

Miss rate components (relative percent) (sum = 100% of total miss rate)								
Cache size (KB)	Degree associative	Total miss rate	Compulsory		Capacity		Conflict	
4	1-way	0.098	0.0001	0.1%	0.070	72%	0.027	28%
4	2-way	0.076	0.0001	0.1%	0.070	93%	0.005	7%
4	4-way	0.071	0.0001	0.1%	0.070	99%	0.001	1%
4	8-way	0.071	0.0001	0.1%	0.070	100%	0.000	0%
8	1-way	0.068	0.0001	0.1%	0.044	65%	0.024	35%
8	2-way	0.049	0.0001	0.1%	0.044	90%	0.005	10%
8	4-way	0.044	0.0001	0.1%	0.044	99%	0.000	1%
8	8-way	0.044	0.0001	0.1%	0.044	100%	0.000	0%
16	1-way	0.049	0.0001	0.1%	0.040	82%	0.009	17%
16	2-way	0.041	0.0001	0.2%	0.040	98%	0.001	2%
16	4-way	0.041	0.0001	0.2%	0.040	99%	0.000	0%
16	8-way	0.041	0.0001	0.2%	0.040	100%	0.000	0%
32	1-way	0.042	0.0001	0.2%	0.037	89%	0.005	11%
32	2-way	0.038	0.0001	0.2%	0.037	99%	0.000	0%
32	4-way	0.037	0.0001	0.2%	0.037	100%	0.000	0%
32	8-way	0.037	0.0001	0.2%	0.037	100%	0.000	0%
64	1-way	0.037	0.0001	0.2%	0.028	77%	0.008	23%
64	2-way	0.031	0.0001	0.2%	0.028	91%	0.003	9%
64	4-way	0.030	0.0001	0.2%	0.028	95%	0.001	4%
64	8-way	0.029	0.0001	0.2%	0.028	97%	0.001	2%
128	1-way	0.021	0.0001	0.3%	0.019	91%	0.002	8%
128	2-way	0.019	0.0001	0.3%	0.019	100%	0.000	0%
128	4-way	0.019	0.0001	0.3%	0.019	100%	0.000	0%
128	8-way	0.019	0.0001	0.3%	0.019	100%	0.000	0%
256	1-way	0.013	0.0001	0.5%	0.012	94%	0.001	6%
256	2-way	0.012	0.0001	0.5%	0.012	99%	0.000	0%
256	4-way	0.012	0.0001	0.5%	0.012	99%	0.000	0%
256	8-way	0.012	0.0001	0.5%	0.012	99%	0.000	0%
512	1-way	0.008	0.0001	0.8%	0.005	66%	0.003	33%
512	2-way	0.007	0.0001	0.9%	0.005	71%	0.002	28%
512	4-way	0.006	0.0001	1.1%	0.005	91%	0.000	8%
512	8-way	0.006	0.0001	1.1%	0.005	95%	0.000	4%

2:1 cache rule of thumb: Direct mapped cache of size N has about the same miss rate as a two-way set associative cache of size N/2

Virtual Memory

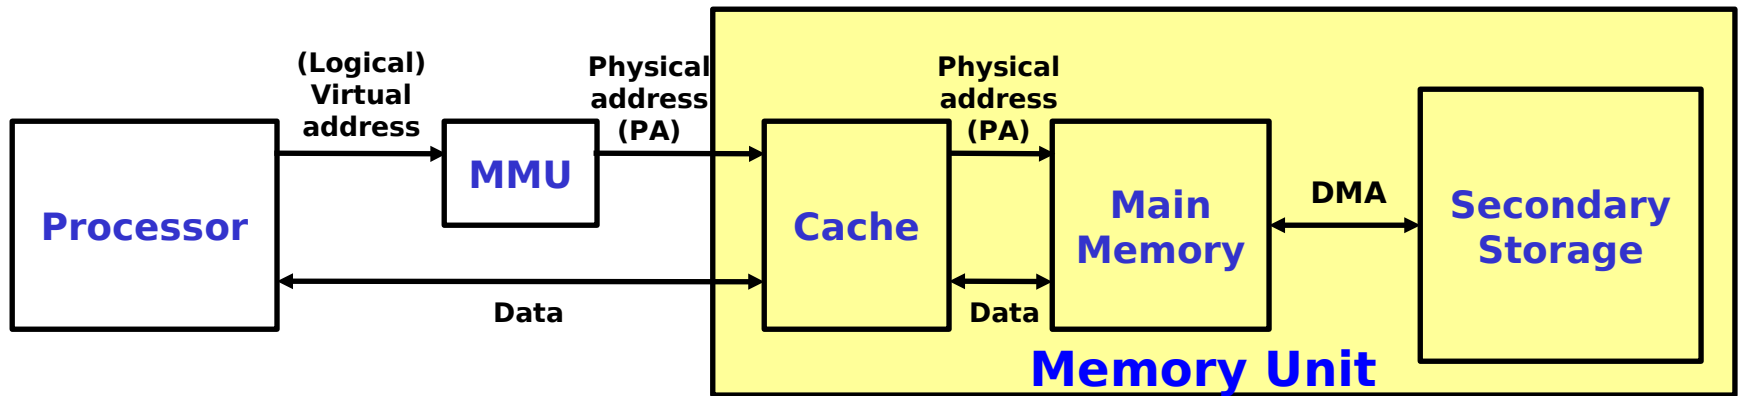
- Ideally, entire memory hierarchy would appear to the processor as a single memory unit
- In modern computer system, the physical main memory is not as large as the address space spanned by the address issued by the processor
- When a program (or process) does not completely fits into the main memory, parts of it will be there in secondary memory
- In modern computers, operating system moves the data automatically between main memory and secondary storage
- Programmer does not need to aware of the limitations imposed by the main memory

Virtual Memory Technique

- Technique that automatically move program and data blocks into the physical main memory when they are required for execution
- Using virtual program concept, each program may use entire CPU local address space, at least up to secondary storage
- The address issued by the processor either for instruction or data are called virtual address or logical address
- These addresses are translated into physical memory addresses by a combination of hardware and software

Memory Management Unit (MMU)


- MMU translate the logical address into physical main memory address
- It is a part of the processor



- If the data is not in main memory, MMU causes the operating system to bring data into memory from the disk
- Transfer of data between disk and main memory is performed using **direct memory access (DMA)** scheme

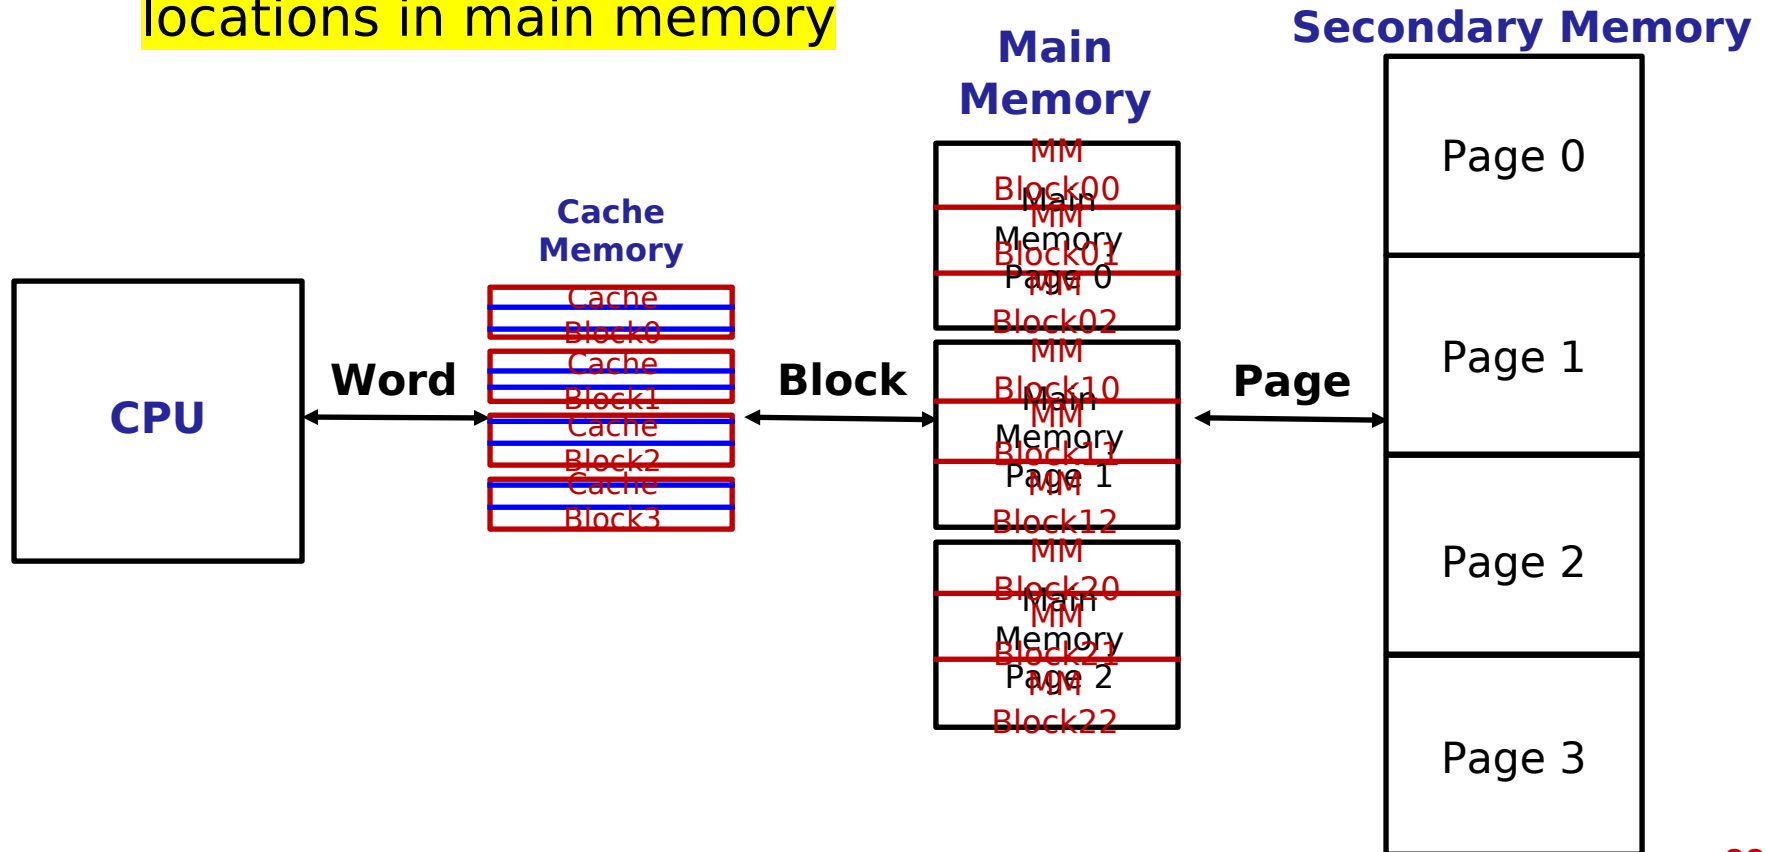
Address Translation

- The **virtual memory address translation** method based is based on the concept of **fixed length pages**
- The address translations assumes that programs and data are composed of **fixed size units called pages**
- **Unit of transfer** between secondary memory and main memory is **page**
 - A page is a block of words that occupy contiguous locations in main memory



Address Translation

- The address translations assumes that programs and data are composed of fixed size units called pages
- Unit of transfer between secondary memory and main memory is page
 - A page is a block of words that occupy contiguous locations in main memory



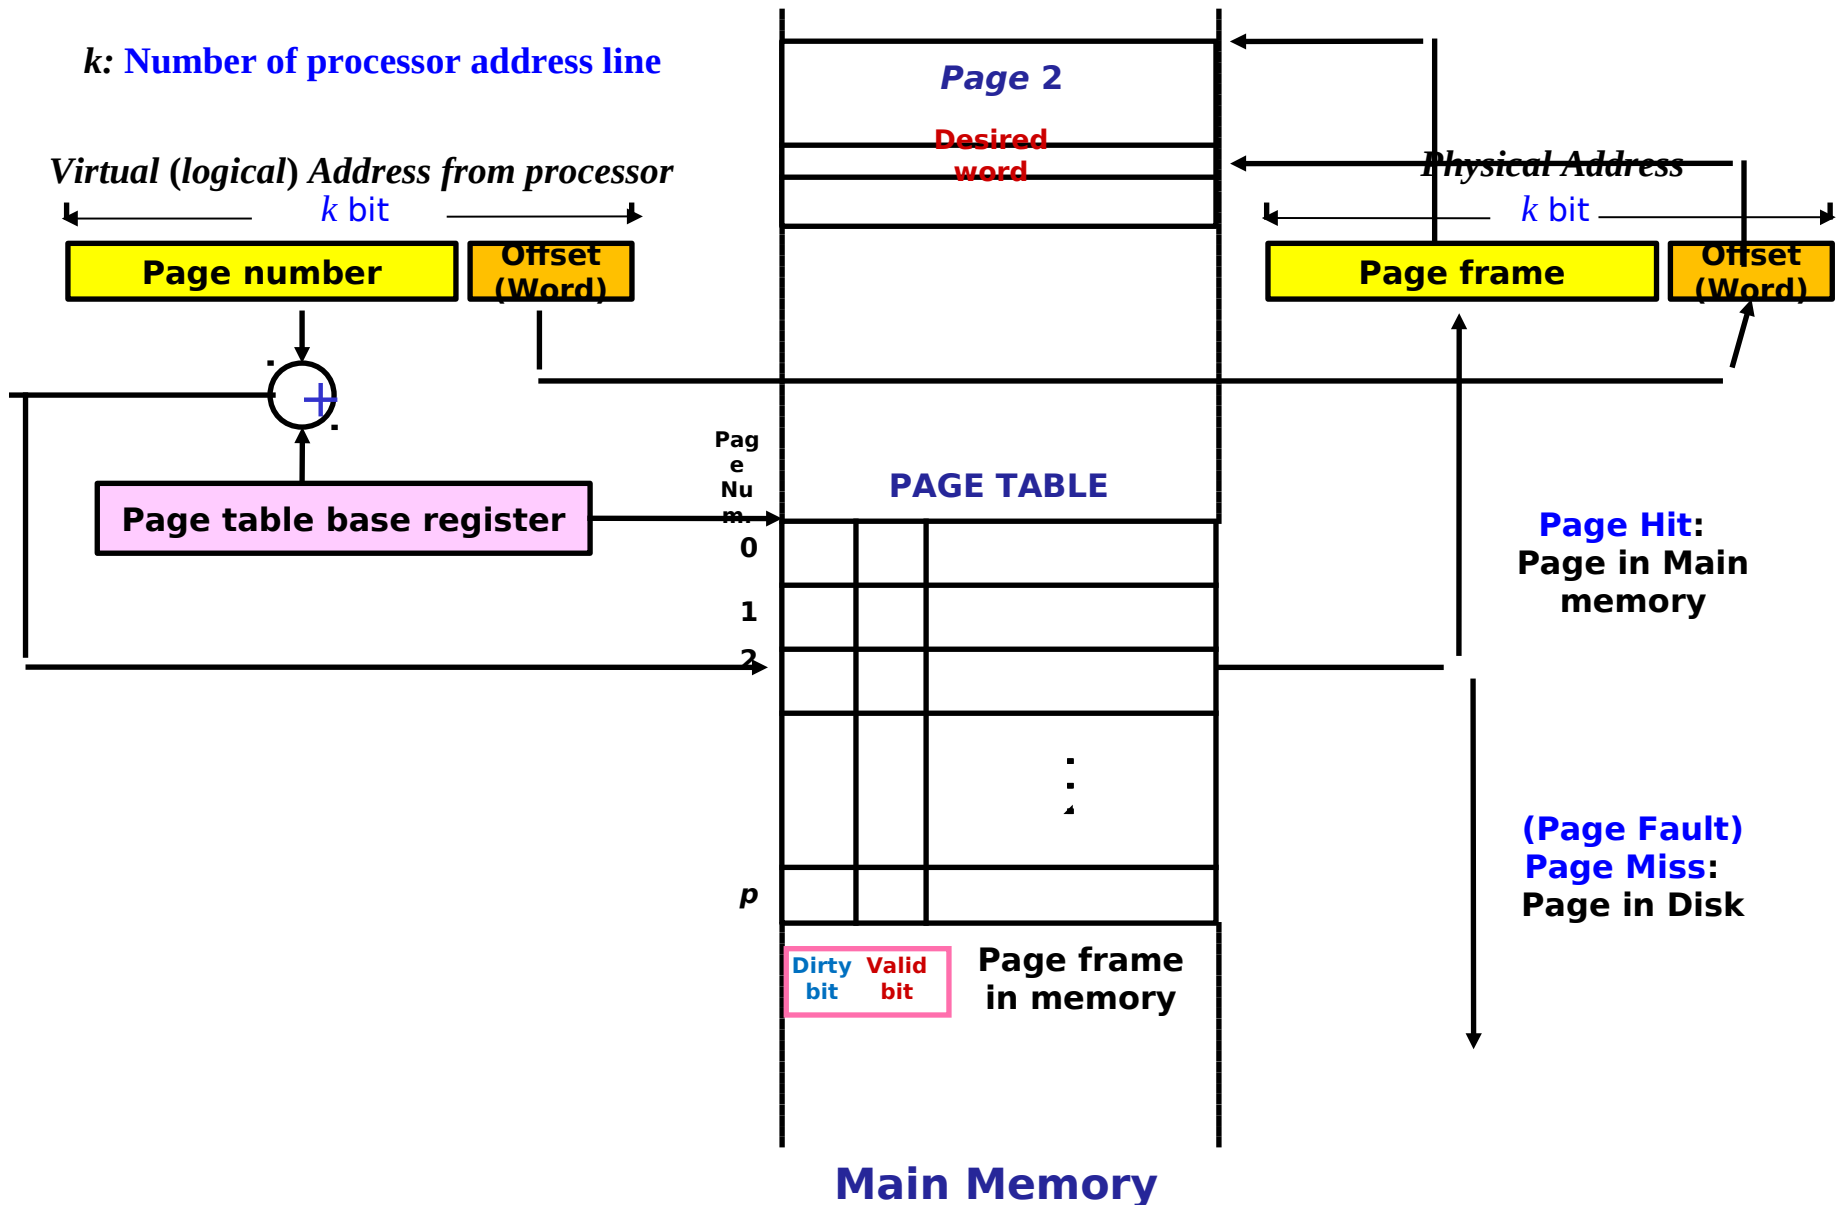
Page

- The programs or data in the disk are seen by the virtual memory as a collection of pages
- This page is the basic unit of information that is moved between the main memory and the secondary memory
- Each page is of the size 2 KB to 16 KB
- Page should not be too small
 - Disk access time is much longer
 - It take considerable time to locate data in the disk
- Page should not be too large
 - Substantial portion of a page may not be used
- Demand paging: Pages are copied to main memory when requested

Parallels Between the Concepts of Cache and Virtual Memory

- Cache:
 - Bridges the speed gap between the processor and the main memory
 - It is implemented in hardware
- Virtual memory mechanism:
 - Bridges the size and speed gap between the main memory and secondary storage
 - It is usually implemented in part by software techniques
- Conceptually, cache techniques and main memory techniques are very similar
- They differ mainly in the details of their implementation

Virtual Memory Address Translation



Virtual Memory Address Translation

- The **virtual address** generated by the processor contain **page number** and **offset (word)** in the page
- To make sure that required page is in main memory, **operating system** create **page table** for each **process**
- This **page table** is kept in main memory
- **Page table base register**: Operating system keep the **starting address of page table** in it
- Page table hold the **main memory location for each page**
 - The area in main memory that can **hold one page** is called **page frame**
- Every entry in page table also include **valid bit** and **dirty bit** to describe the status of the page while it is in main memory

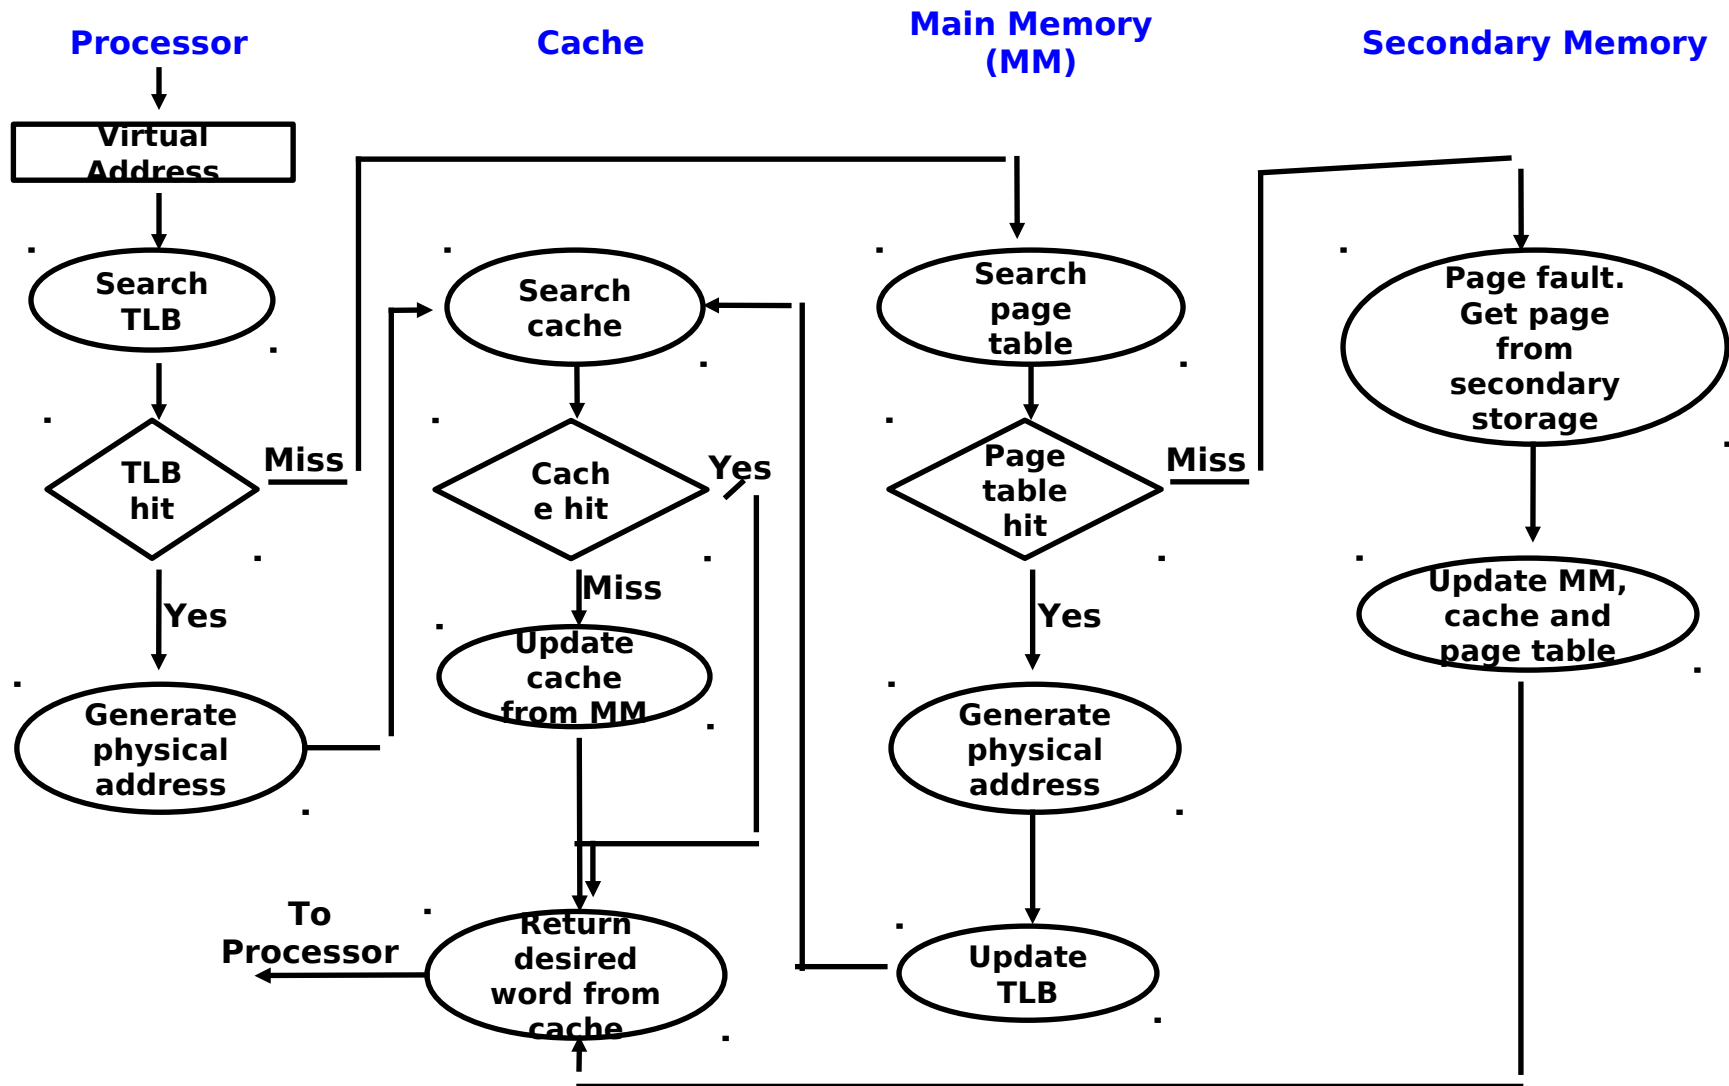
Virtual Memory Address Translation using Translation Lookaside Buffer

- Page table information is used by the MMU for every read and write access
- In order to speed up the address translation procedure, a small cache called Translation Lookaside Buffer (TLB) is incorporated in MMU
- It uses associative/set-associative mapping technique
- It holds a portion of page table corresponding to most recently accessed pages
- TLB holds only the page number and page frame number

Virtual Memory Address Translation

- Page table information is used by the MMU for every read and write access
- Page fault: Whenever a requested page is not present in the main memory, page fault is said to have occurred
- When a page fault occurs, MMU asks operating system to intervene and raise an exception (interrupt)
 - Process in active get interrupted and control goes to operating system
 - Operating system then copies the requested page from disk to main memory
 - Then returns the control to the interrupted task
- During write operation pages get modified are indicated by dirty bit
- Modified page has to be written back to disk before removed from main memory
- Uses write back policy only

Operation of Memory Hierarchy and Virtual Memory Technique



Operation of Memory Hierarchy and Virtual Memory Technique

